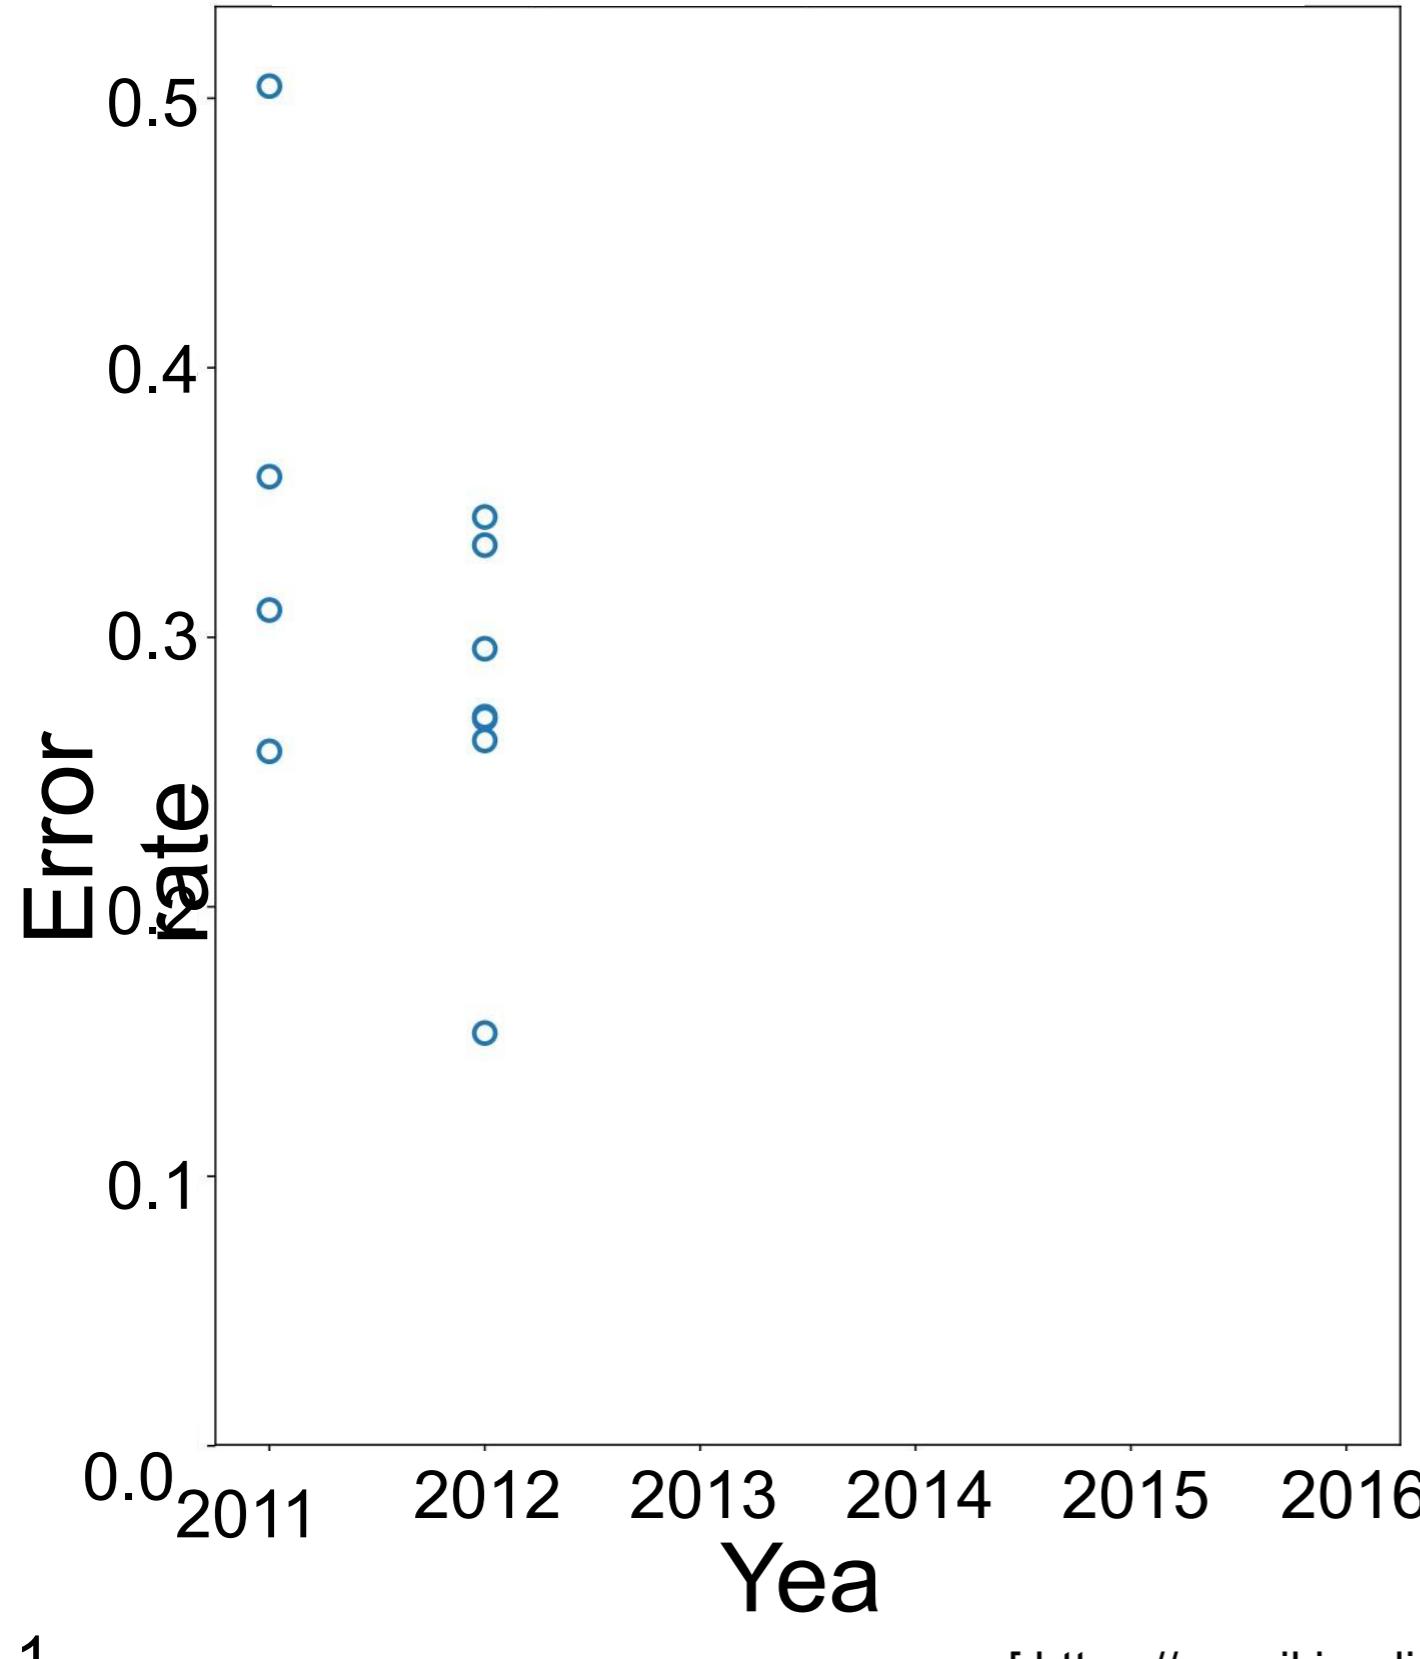


Convolutional Neural Networks (CNNs, ConvNets)

Somak Aditya
IIT KHARAGPUR

Impact of CNNs

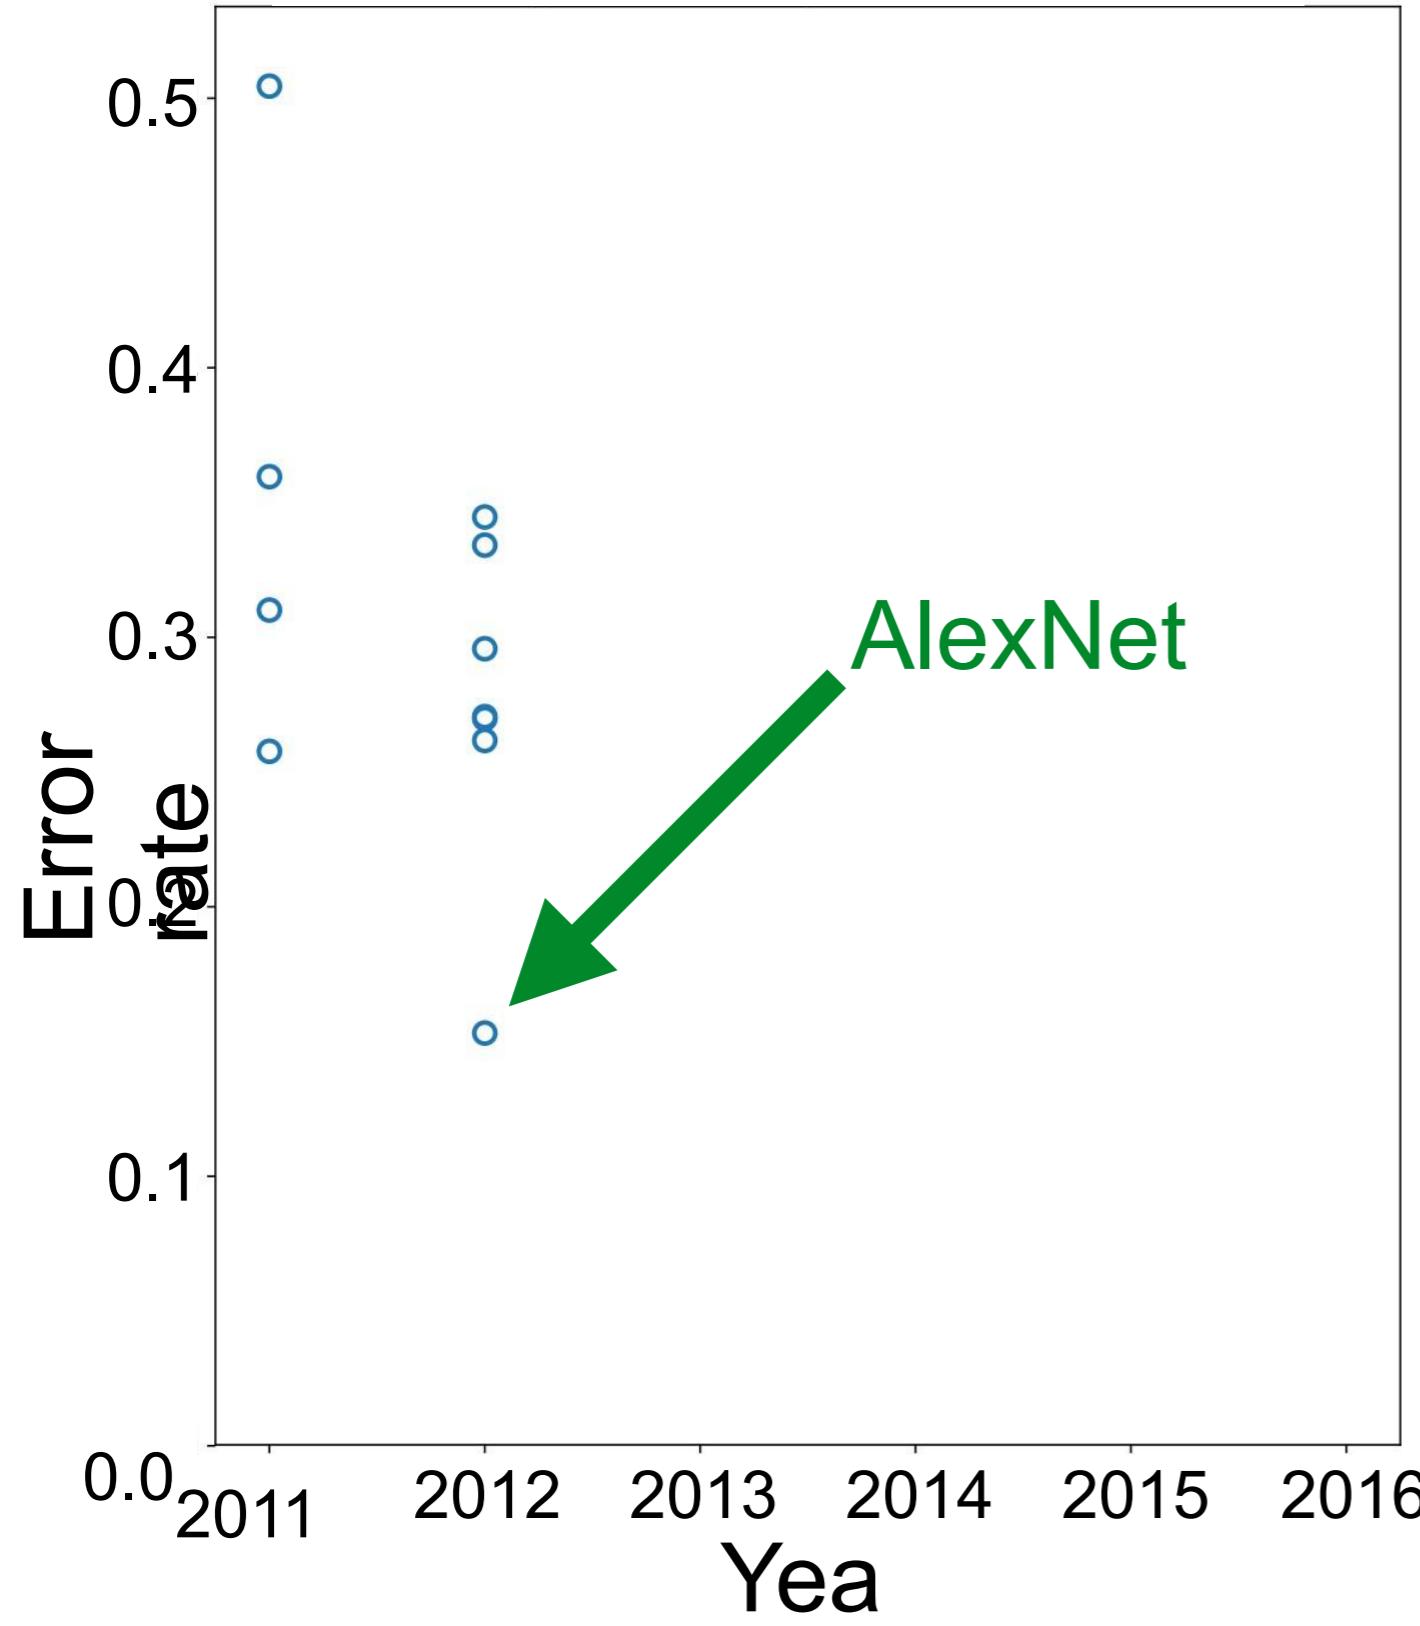
ImageNet results



- Since 2010: large-scale image classification challenge

Impact of CNNs

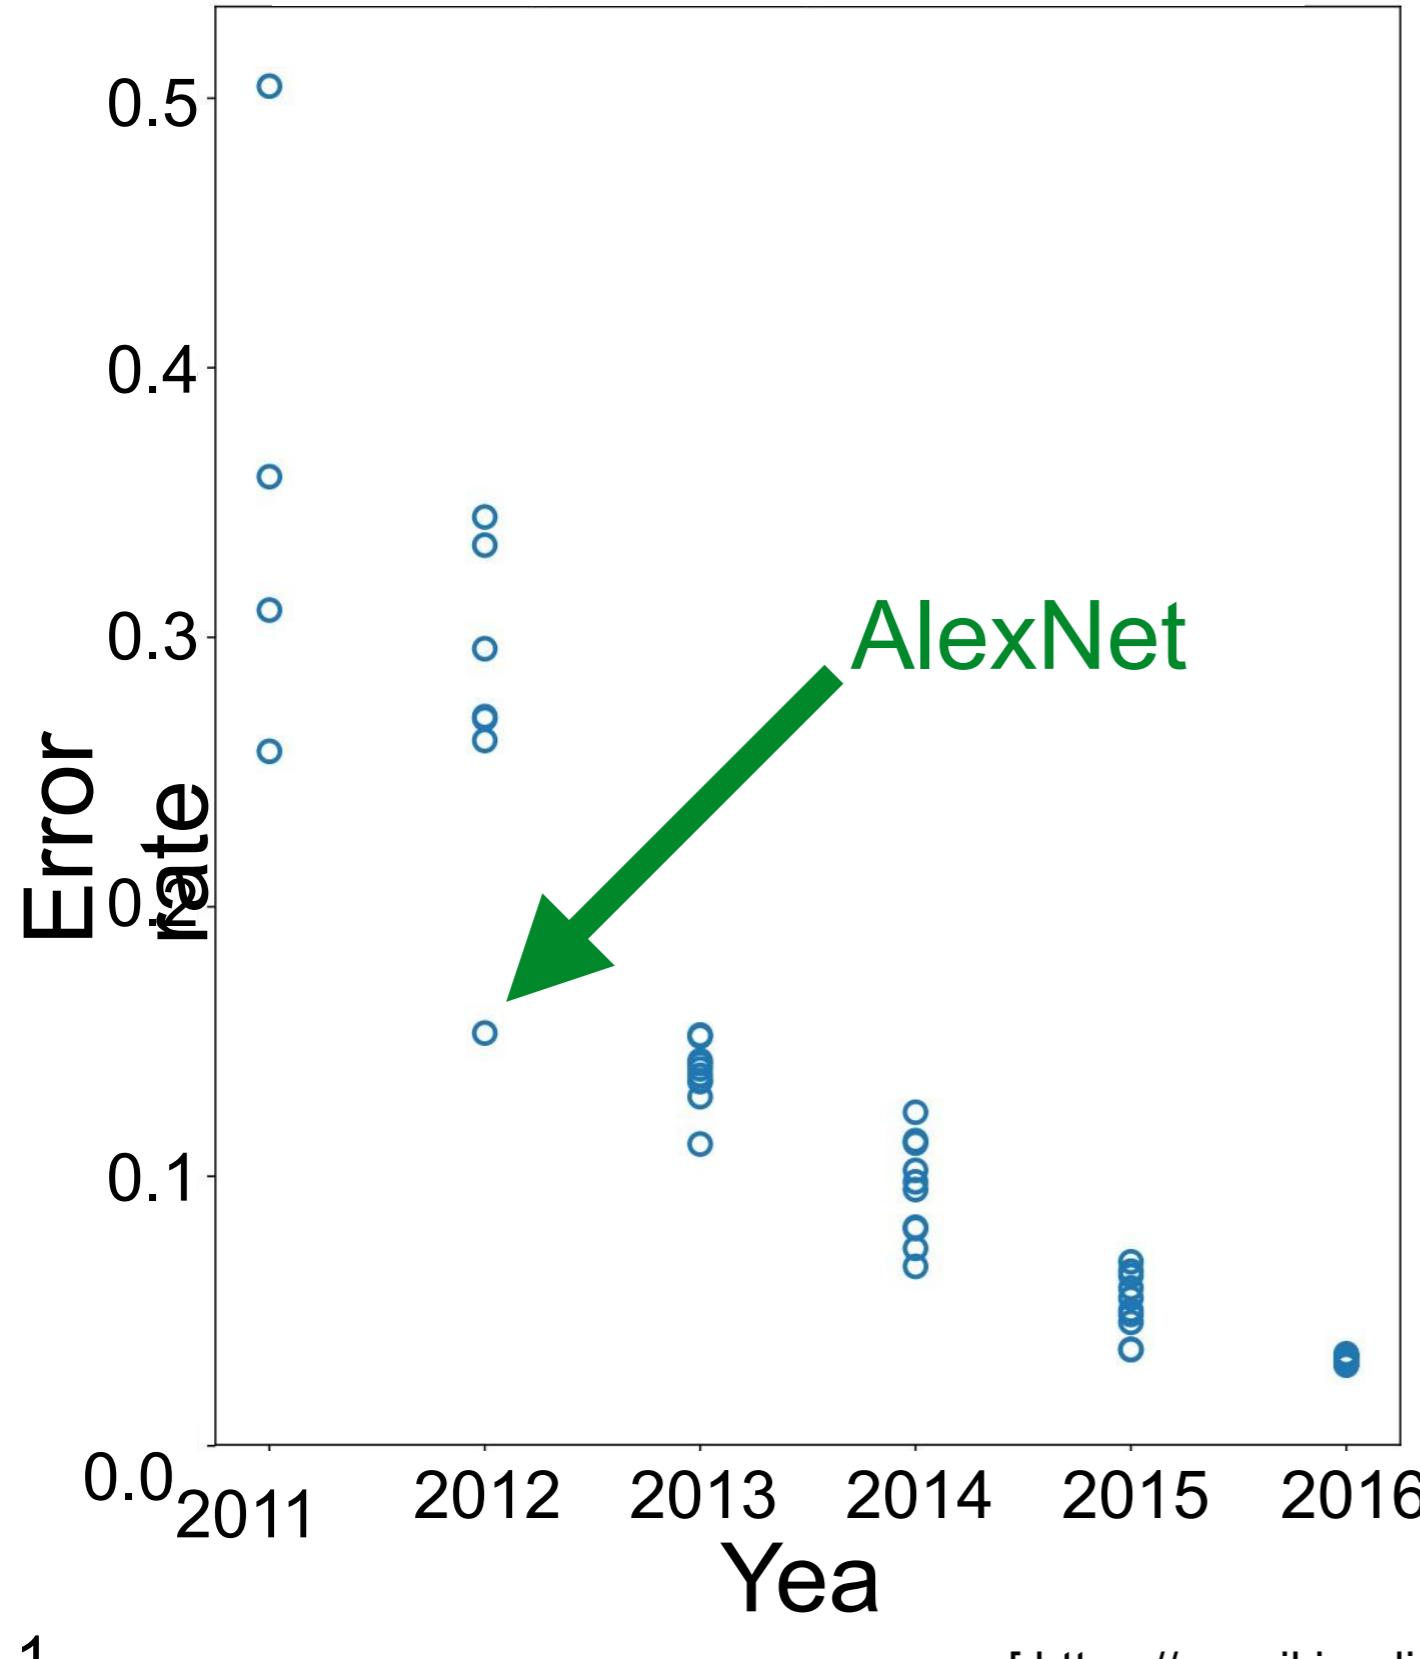
ImageNet results



- Since 2010: large-scale image classification challenge

Impact of CNNs

ImageNet results



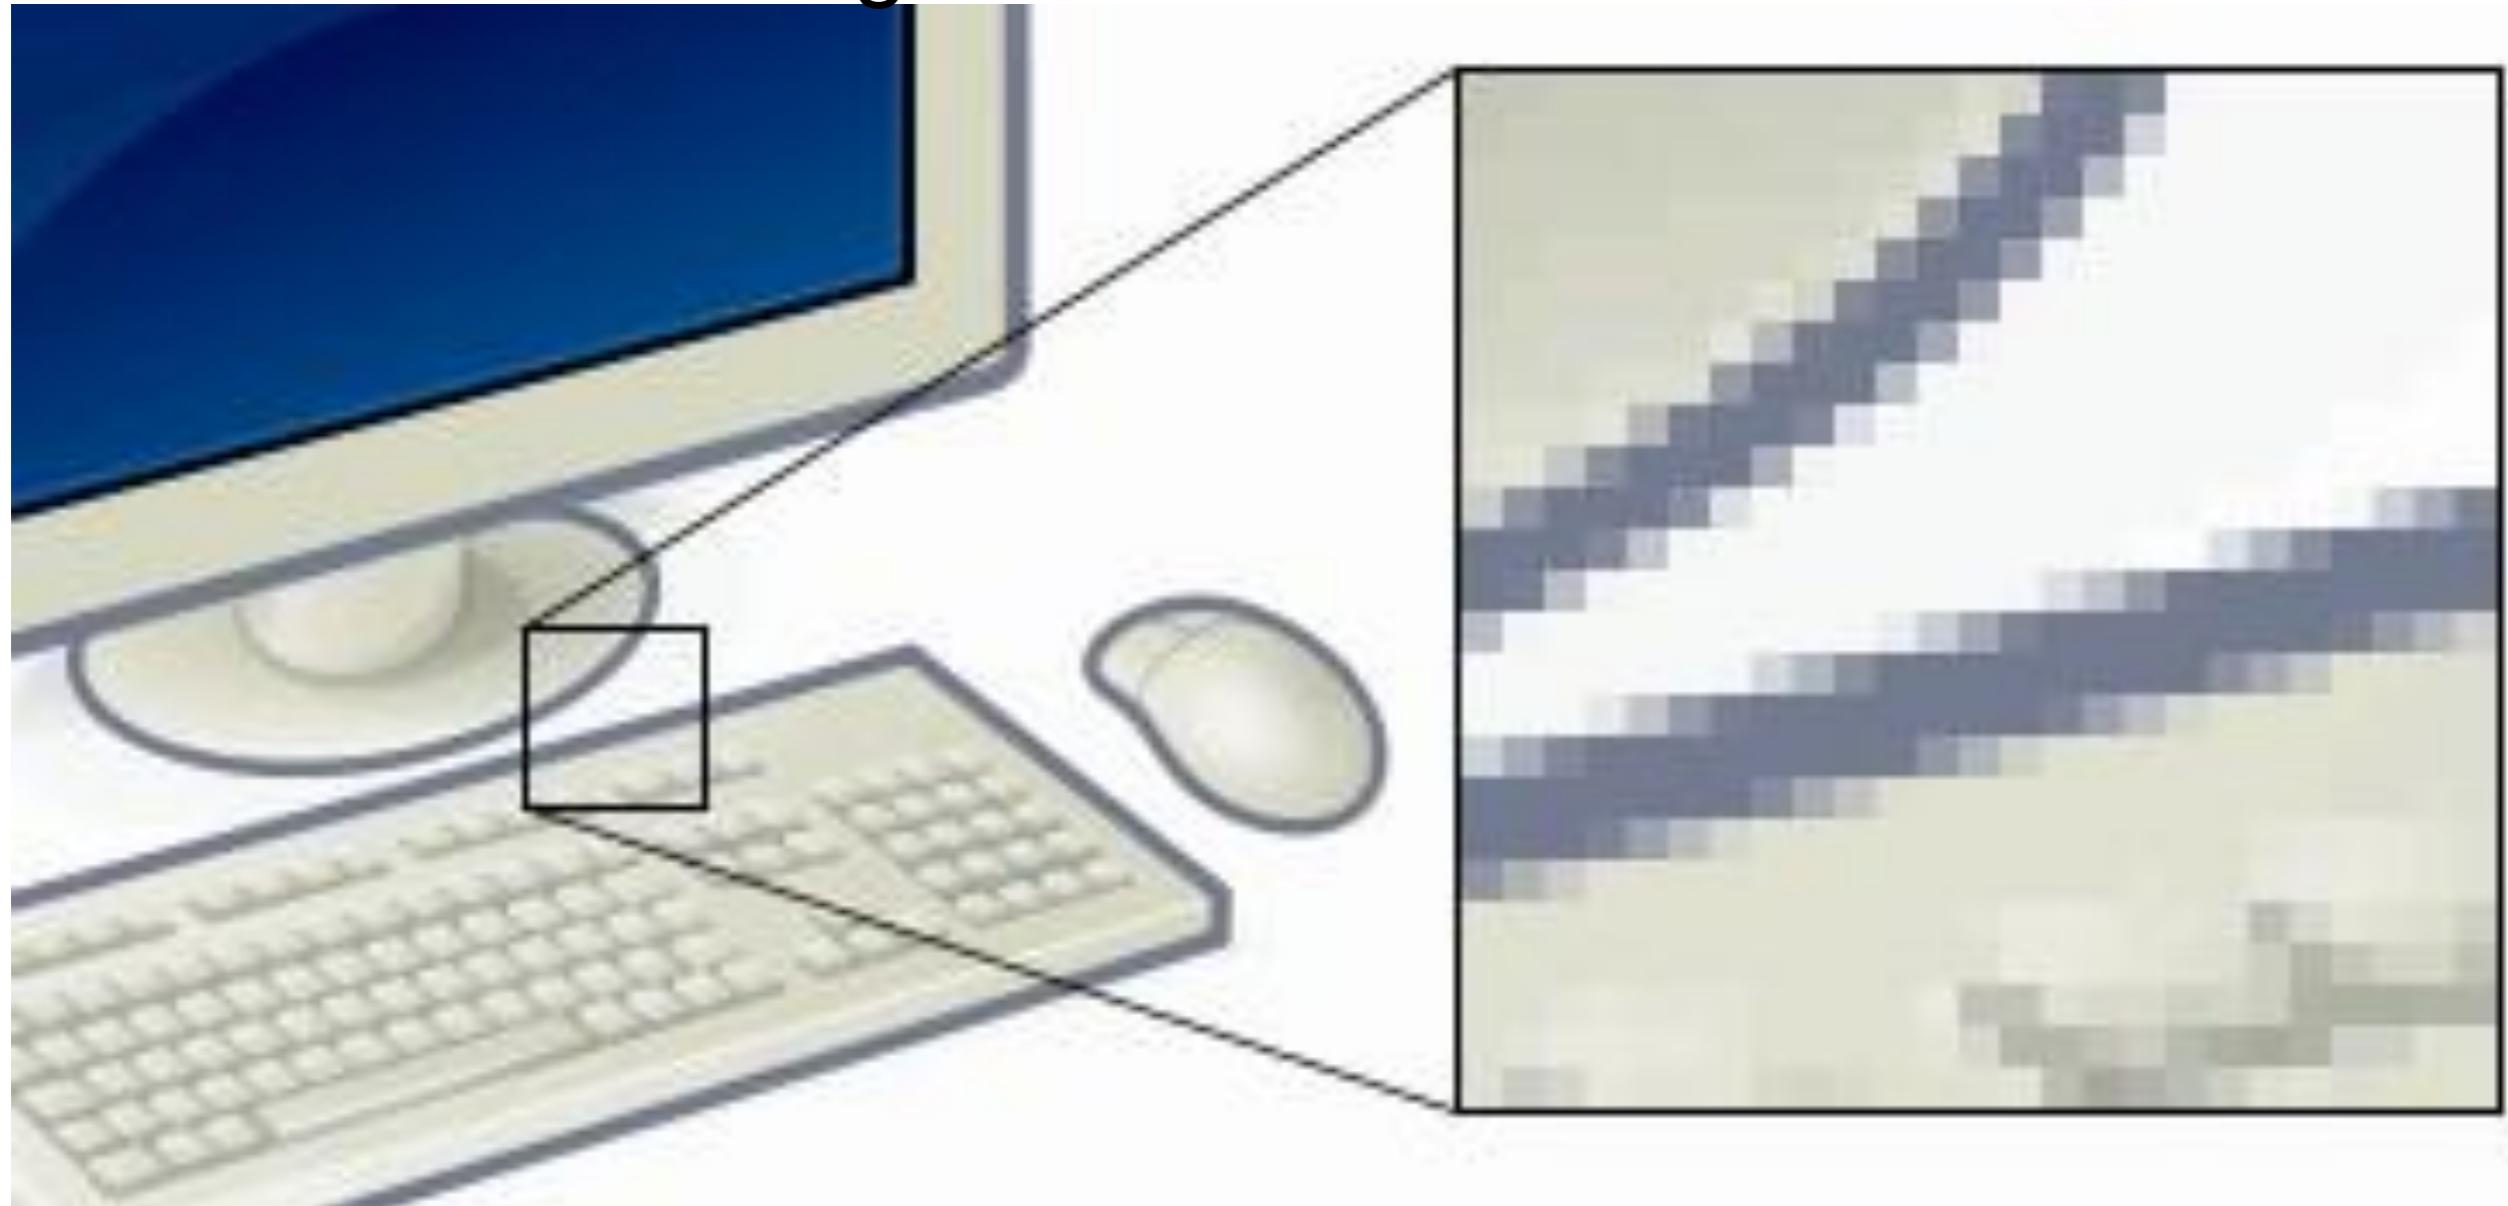
- Since 2010: large-scale image classification challenge
- Recent AI boom
- 1960s, 1980s, today: neural networks
- Since 1980s: CNNs

Images

- Potential uses of image classification: Detect tumor (type) from medical scans, image search online, autonomous driving

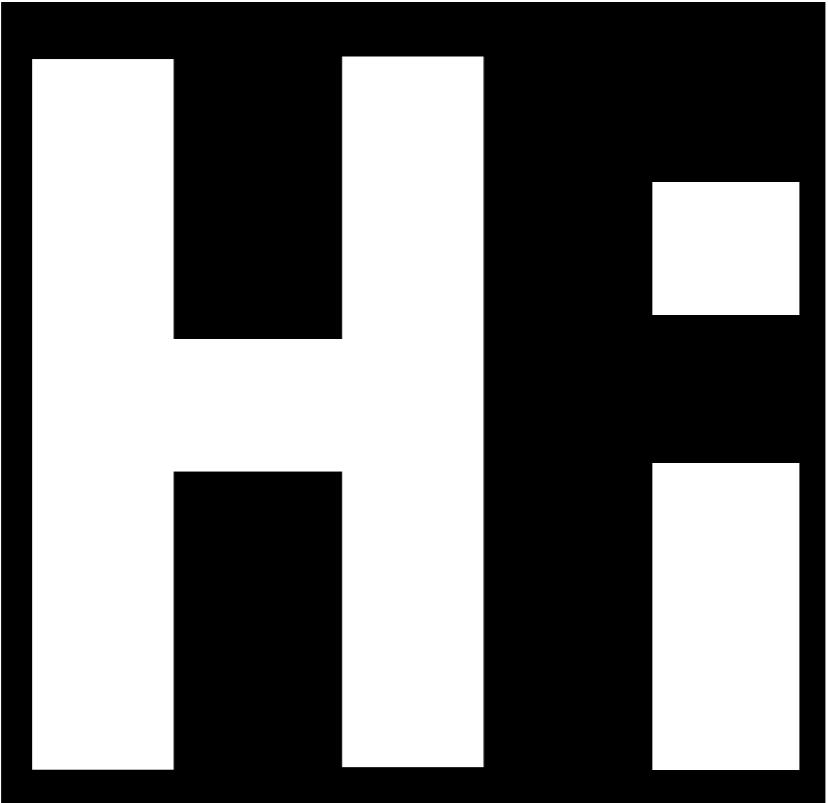
Images

- Potential uses of image classification: Detect tumor (type) from medical scans, image search online, autonomous driving



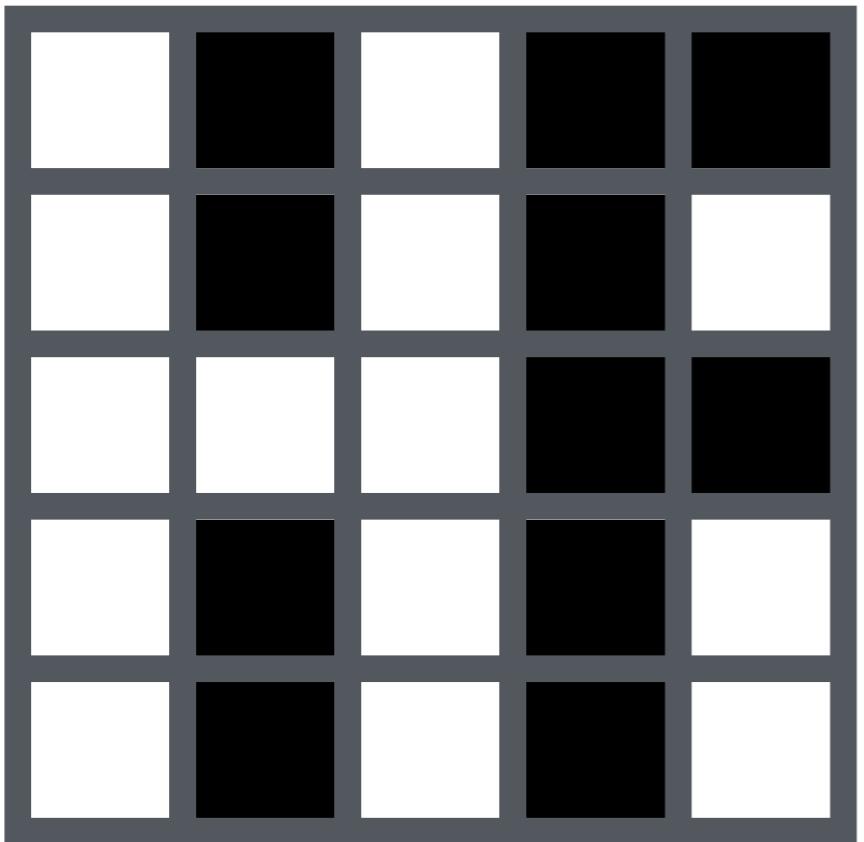
- Recall: images are made of pixels

Images



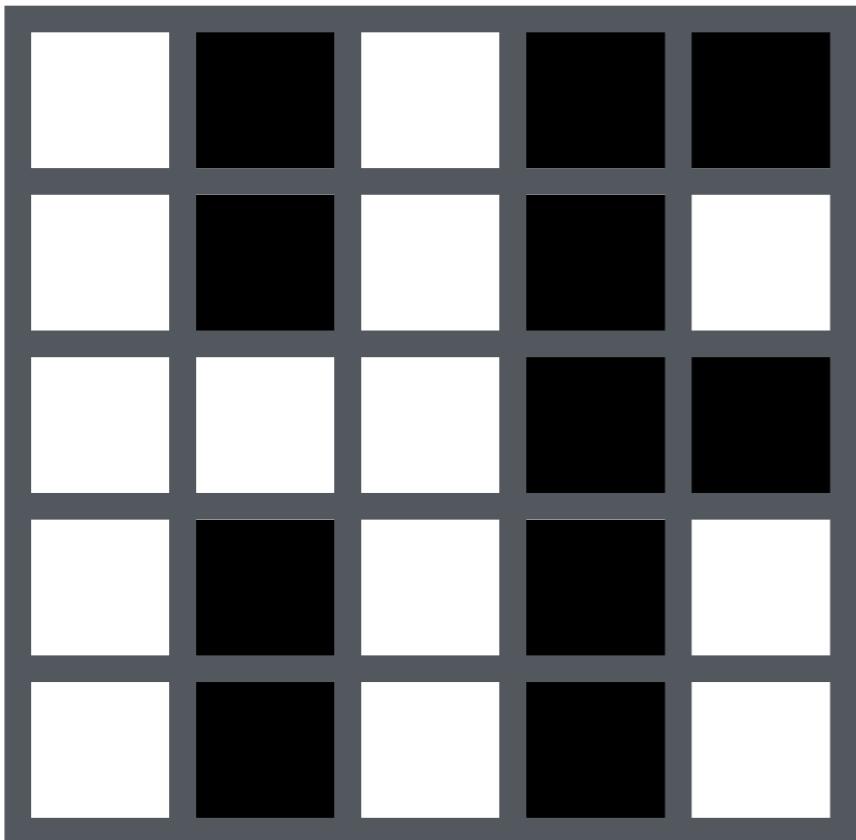
- We'll focus on grayscale images

Images



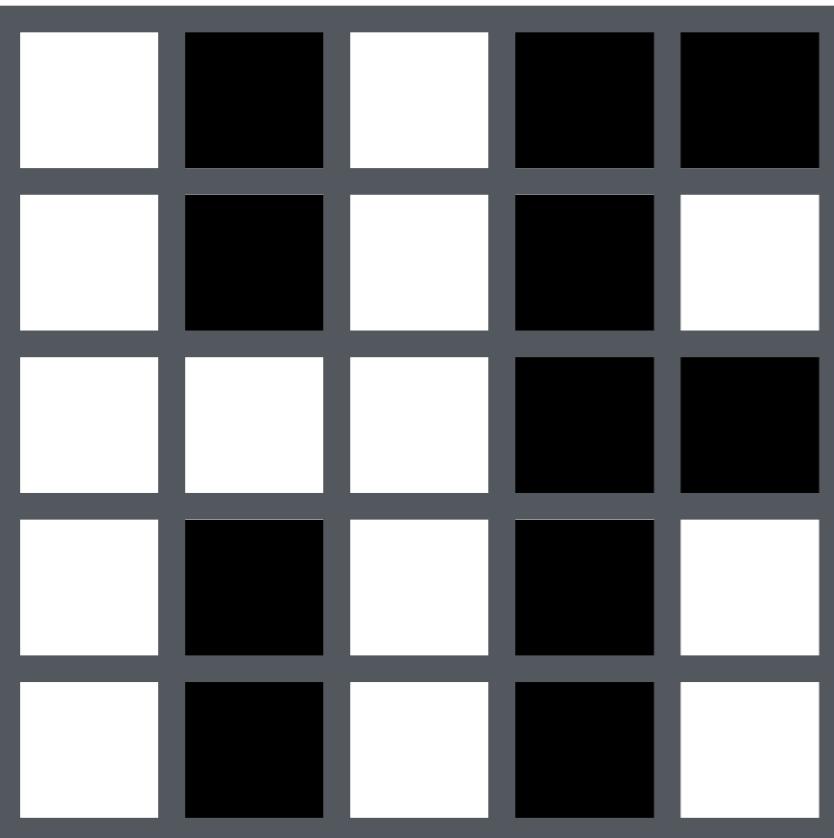
- We'll focus on grayscale images

Images



- We'll focus on grayscale images
 - Each pixel takes a value between 0 and P
 - Here, 0: black, 1: white
 - Larger P in Lab Week 08

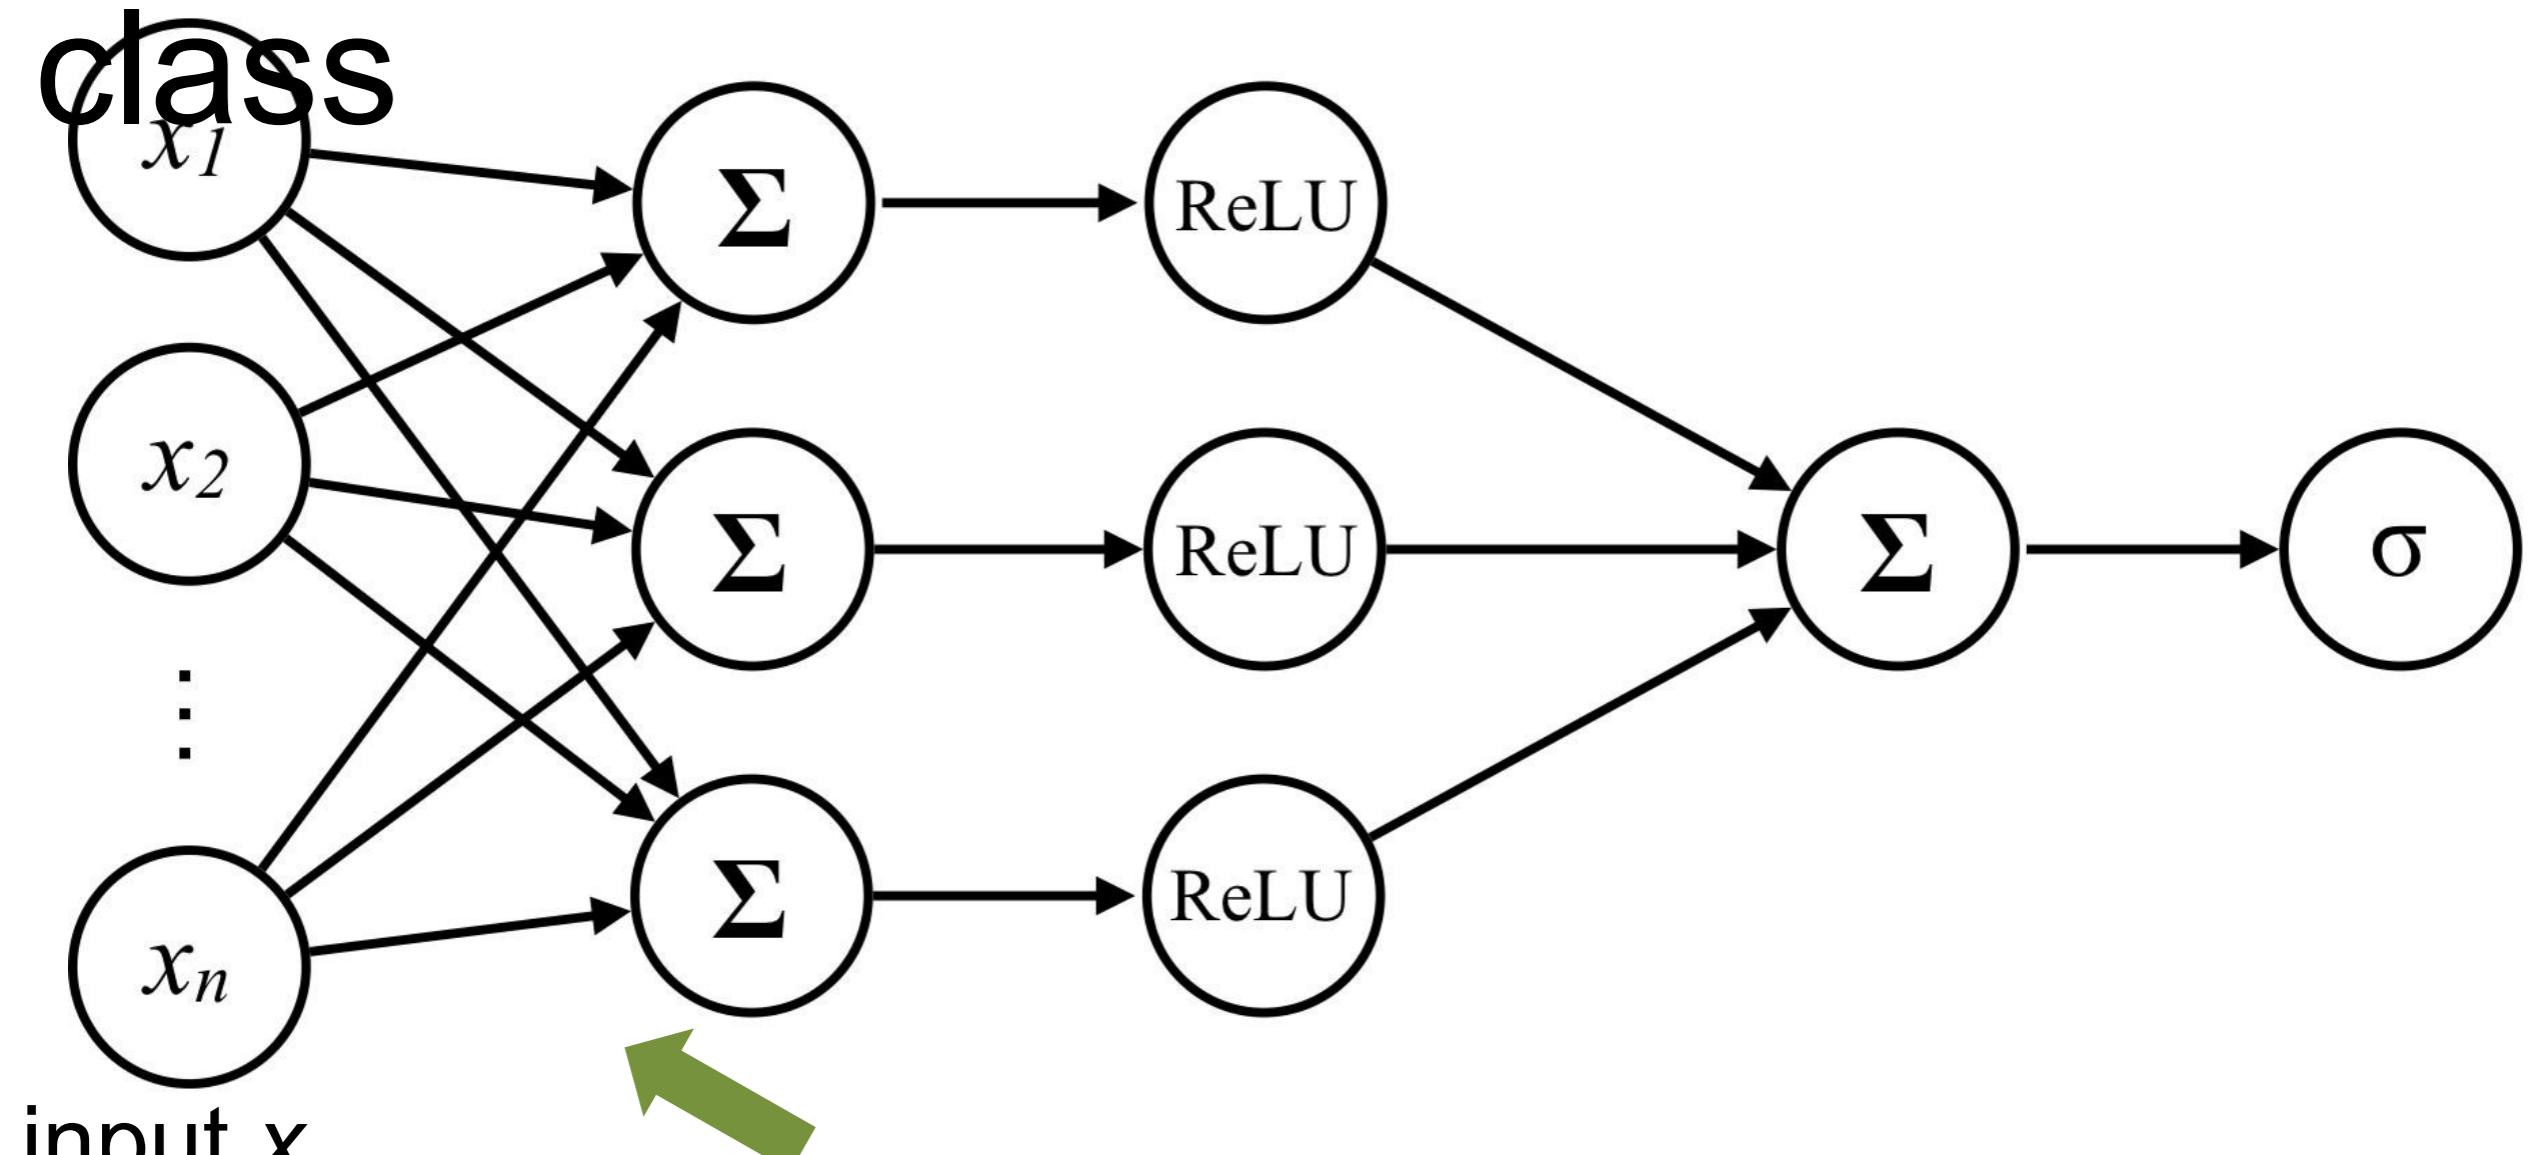
Images



1	0	1	0	0
1	0	1	0	1
1	1	1	0	0
1	0	1	0	1
1	0	1	0	1

- We'll focus on grayscale images
 - Each pixel takes a value between 0 and P
 - Here, 0: black, 1: white
 - Larger P in Lab Week 08
- How do we use an image as an input for a neural net?

Previous neural nets in this class



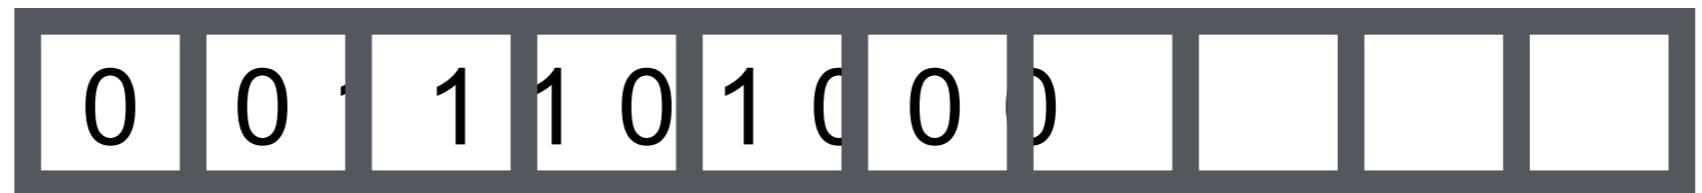
Fully connected layer: every input is connected to every output by a weight

But we know more about images:

- Spatial locality
- Translation invariance

Convolutional Layer: 1D example

A 1D image:



Convolutional Layer: 1D example

A 1D image:

0	0	1	1	10	11	0	0	1	0	0	0
---	---	---	---	----	----	---	---	---	---	---	---

A filter:

-1	1	-1
----	---	----

Convolutional Layer: 1D example

A 1D image:

0	0	1	1	10	11	0	0	1	0	0	0
---	---	---	---	----	----	---	---	---	---	---	---

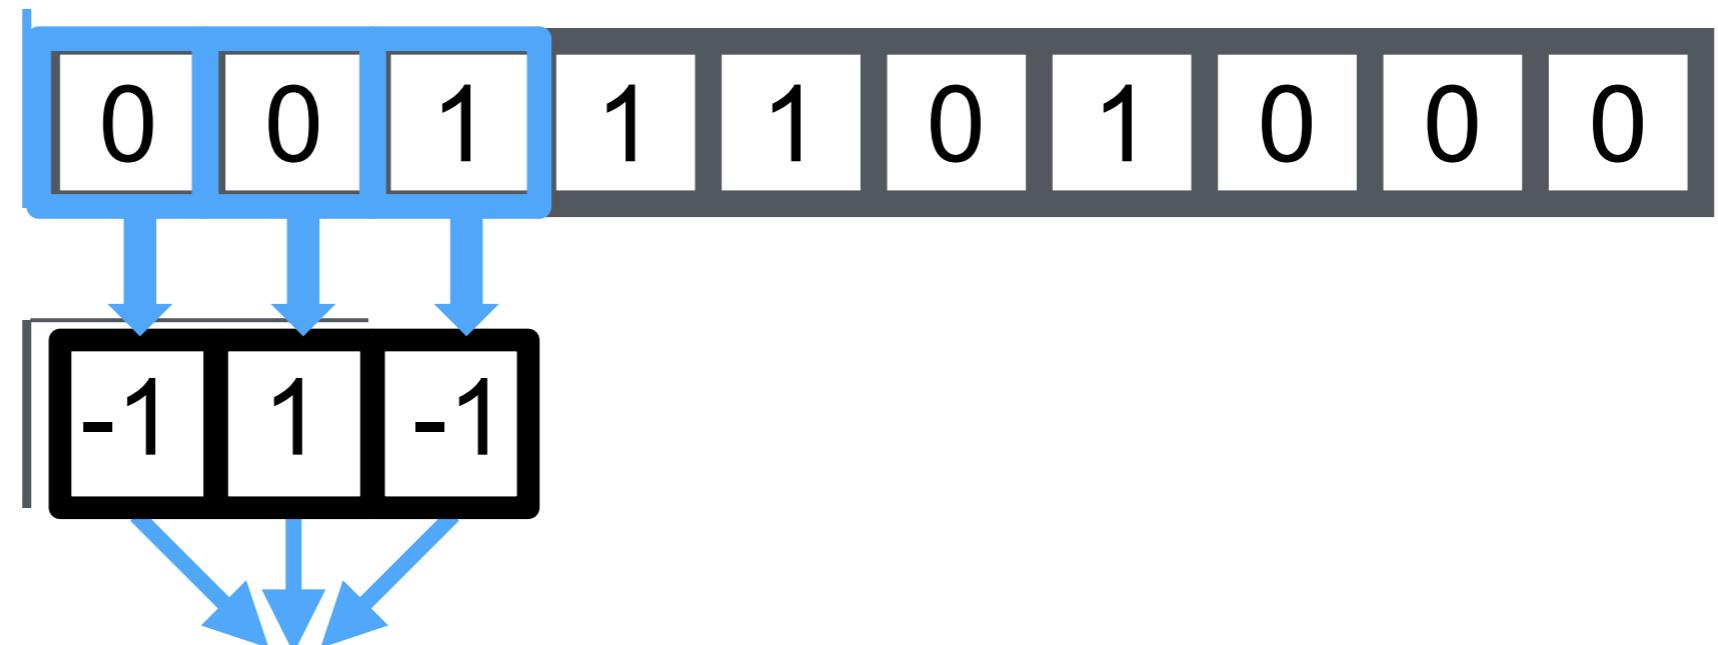
A filter:

-1	1	-1
----	---	----

After
convolution*:

Convolutional Layer: 1D example

A 1D image:



A filter:

After
convolution*:

Convolutional Layer: 1D example

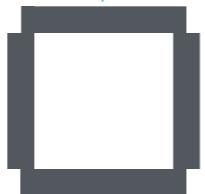
A 1D image:



A filter:

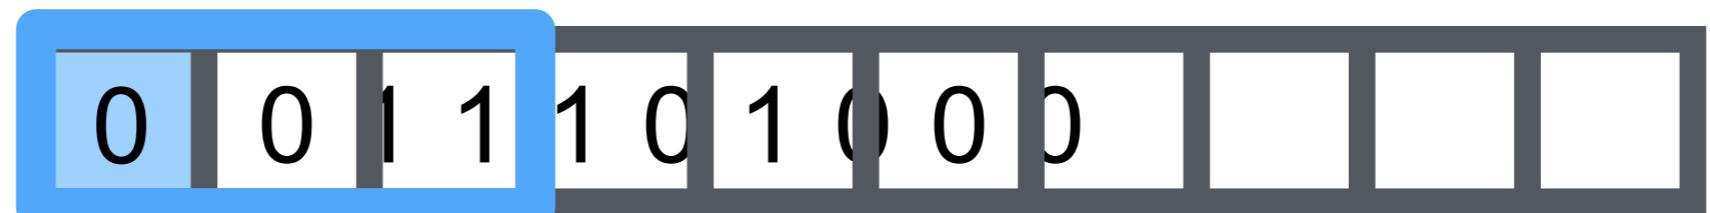


After
convolution*:

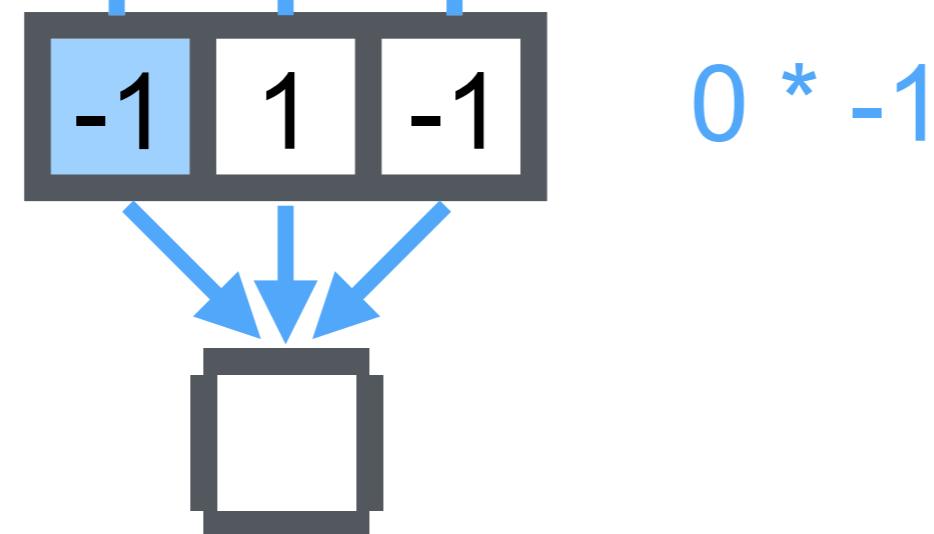


Convolutional Layer: 1D example

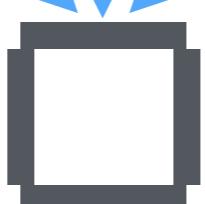
A 1D image:



A filter:



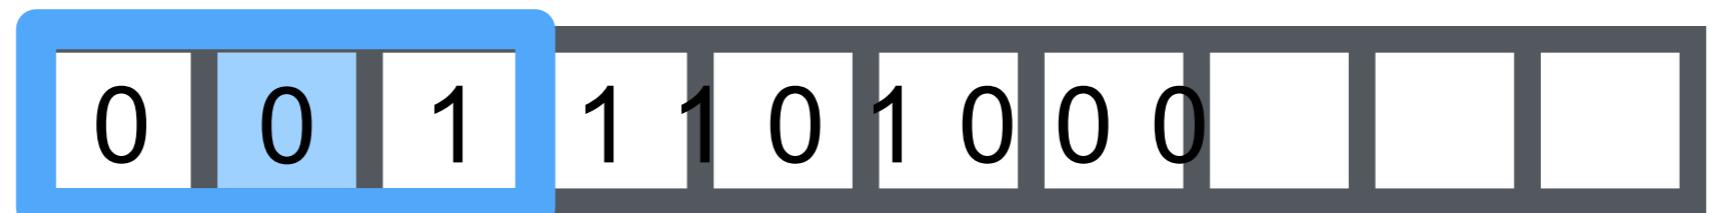
After
convolution*:



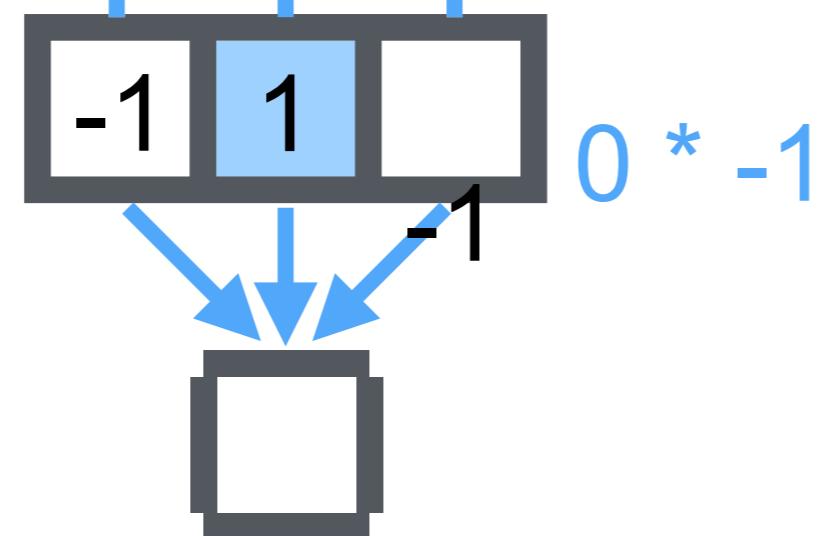
$$0 * -1$$

Convolutional Layer: 1D example

A 1D image:



A filter:



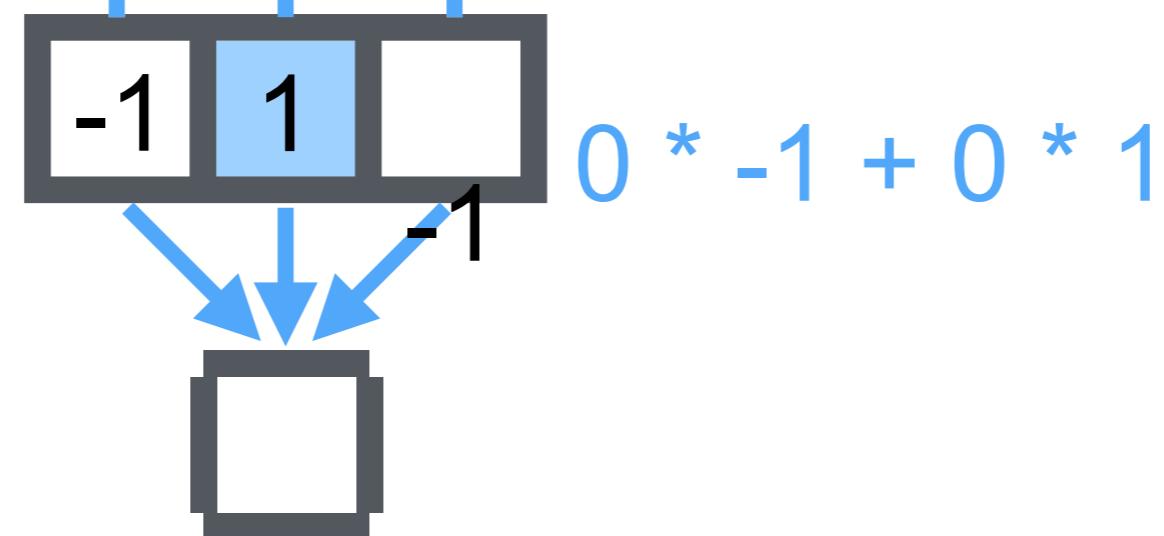
After
convolution*:

Convolutional Layer: 1D example

A 1D image:



A filter:

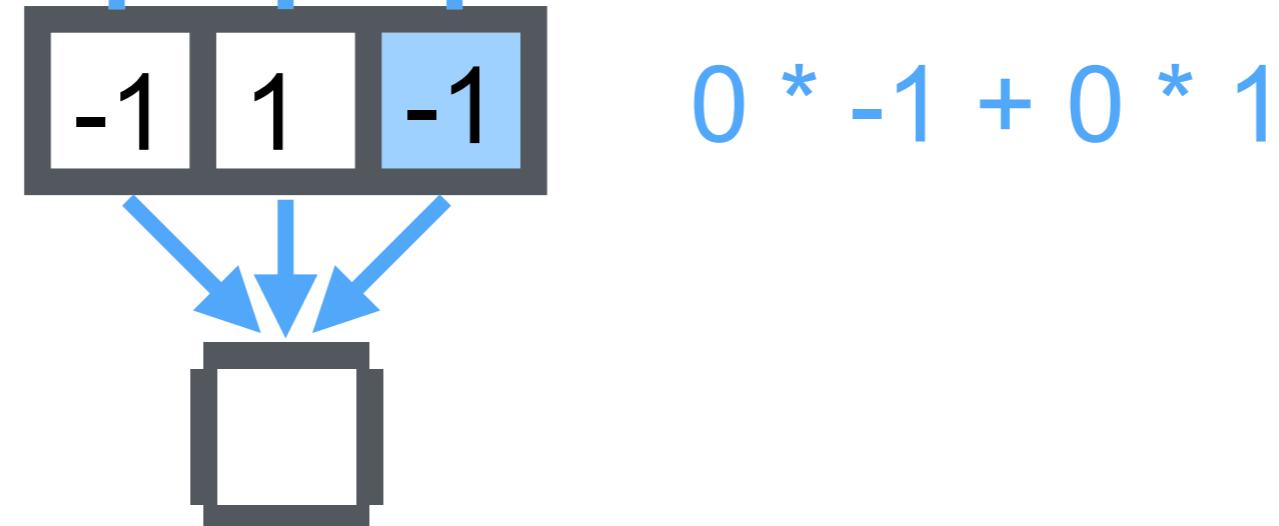


Convolutional Layer: 1D example

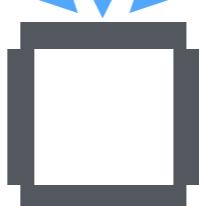
A 1D image:



A filter:

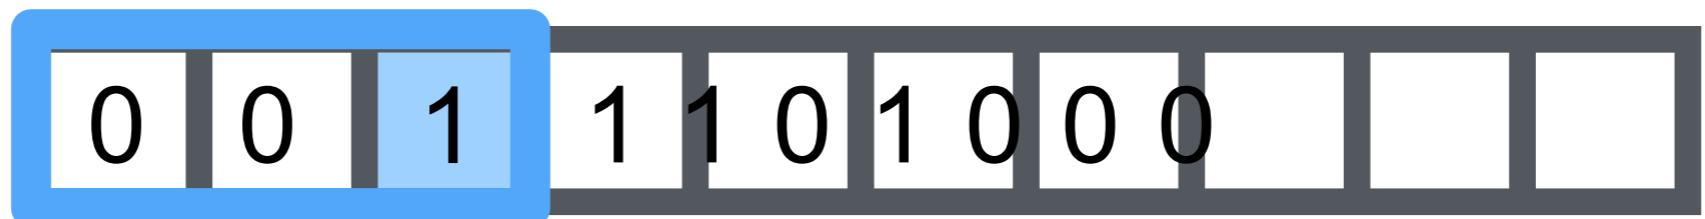


After
convolution*:



Convolutional Layer: 1D example

A 1D image:

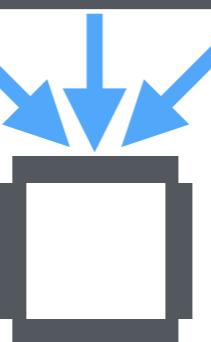


A filter:



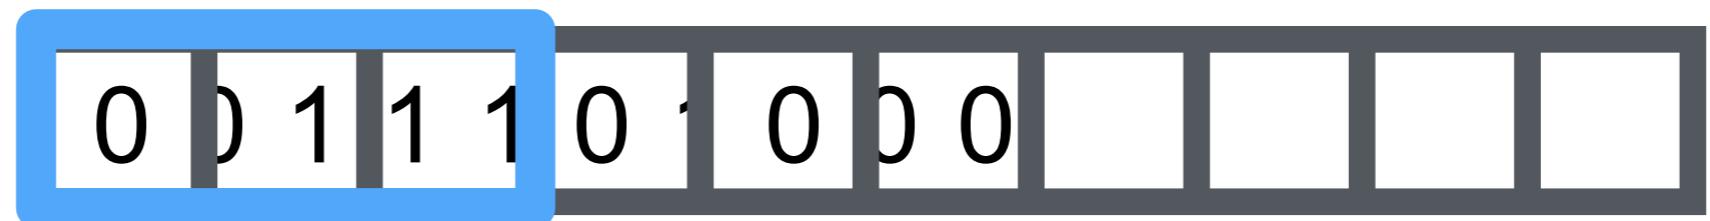
$$0 * -1 + 0 * 1 + 1 * -1$$

After
convolution*:



Convolutional Layer: 1D example

A 1D image:

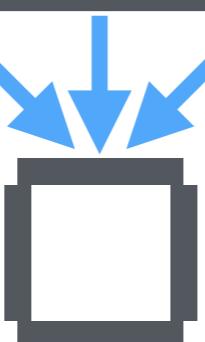


A filter:



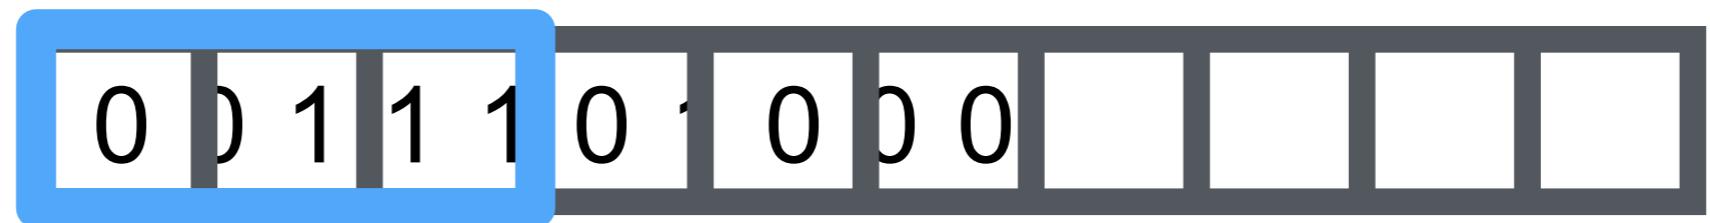
$$0 * -1 + 0 * 1 + 1 * -1$$

After
convolution*:

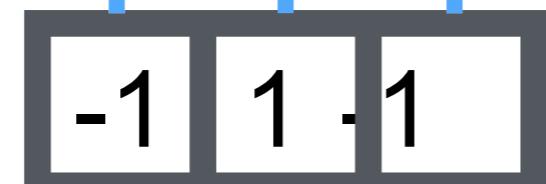


Convolutional Layer: 1D example

A 1D image:

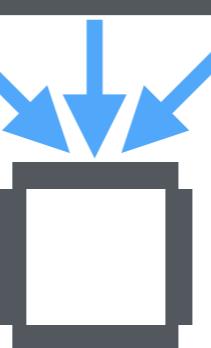


A filter:



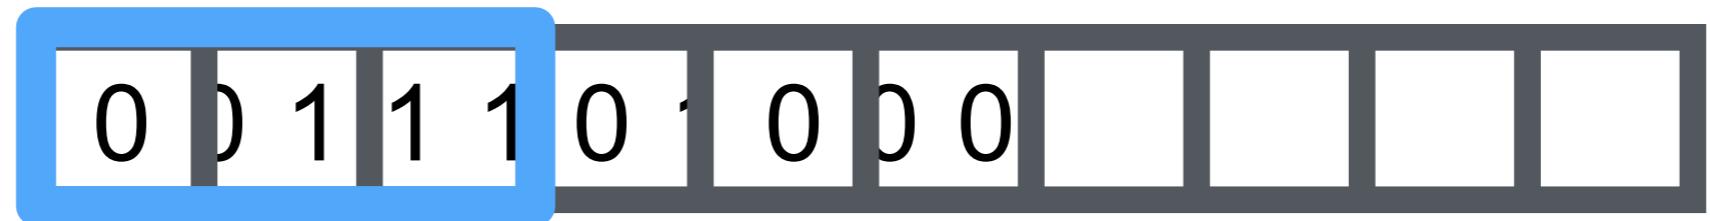
$$0 * -1 + 0 * 1 + 1 * -1 = -1$$

After
convolution*:



Convolutional Layer: 1D example

A 1D image:

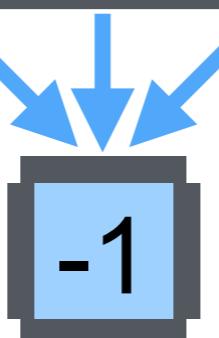


A filter:



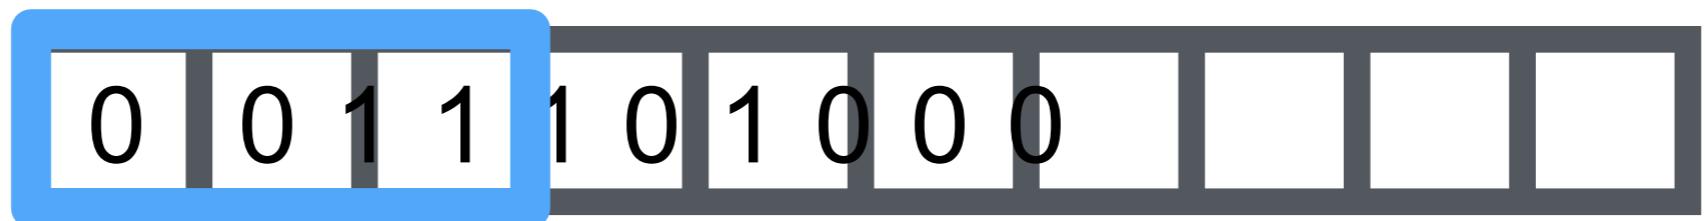
$$0 * -1 + 0 * 1 + 1 * -1 = -1$$

After
convolution*:



Convolutional Layer: 1D example

A 1D image:



A filter:

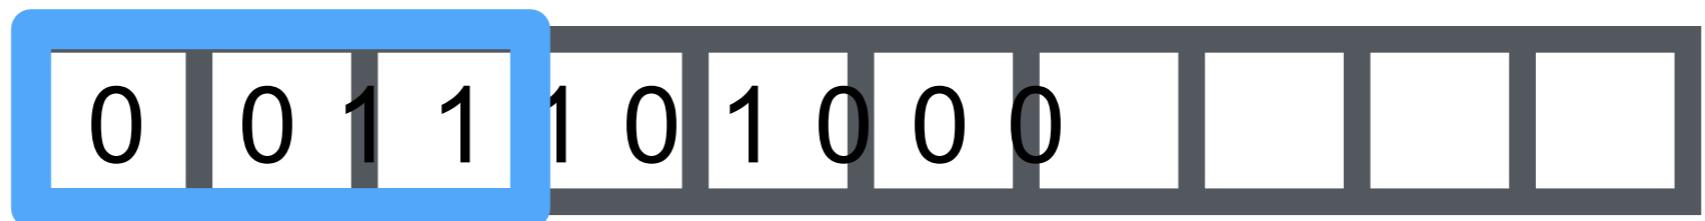


After
convolution*:



Convolutional Layer: 1D example

A 1D image:



A filter:



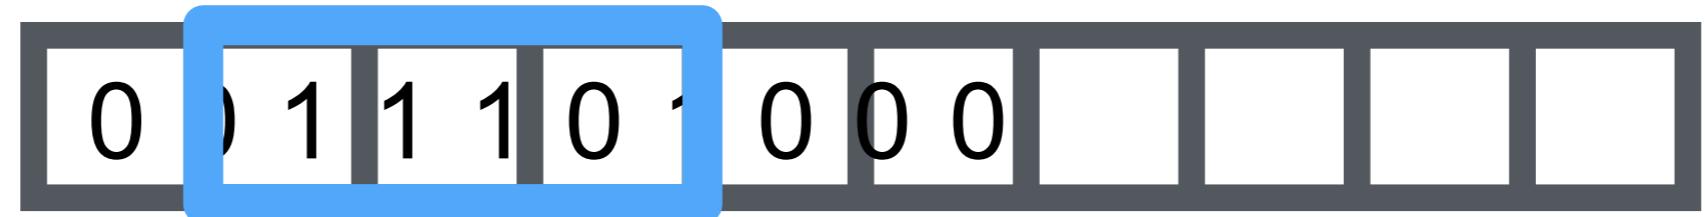
After
convolution*:



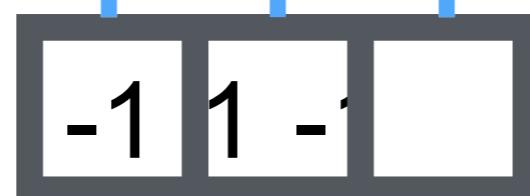
*correlation

Convolutional Layer: 1D example

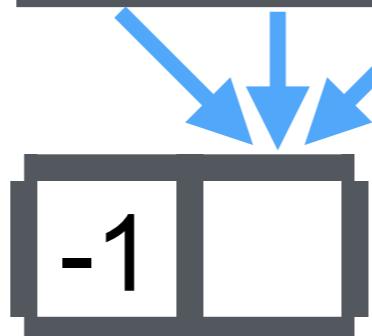
A 1D image:



A filter:



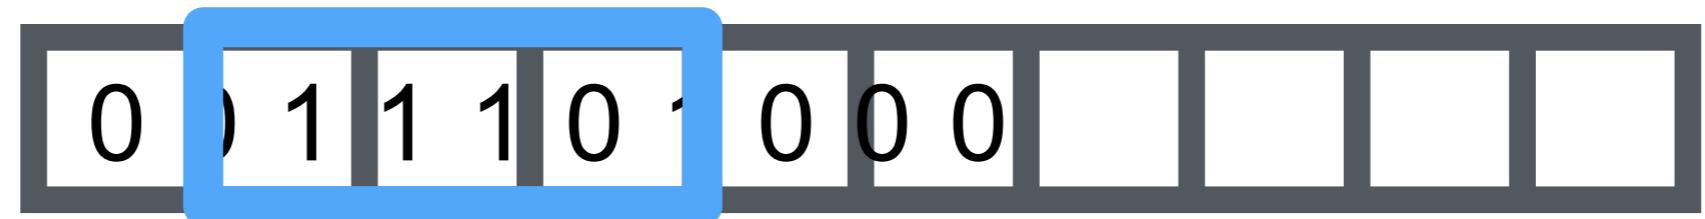
After
convolution*:



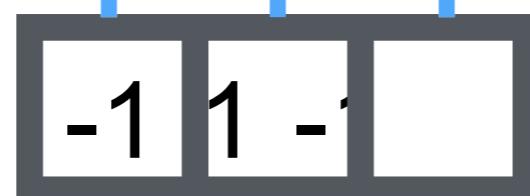
*correlation

Convolutional Layer: 1D example

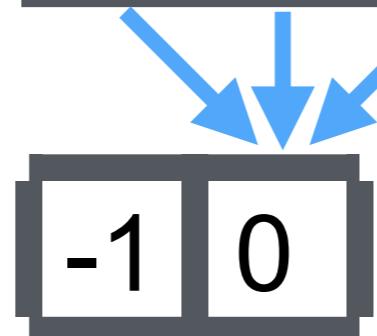
A 1D image:



A filter:



After
convolution*:



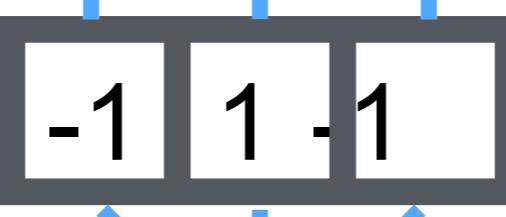
*correlation

Convolutional Layer: 1D example

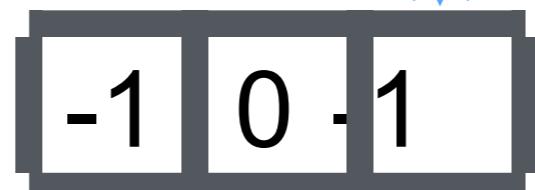
A 1D image:



A filter:



After
convolution*:



*correlation

Convolutional Layer: 1D example

A 1D image:



A filter:



After
convolution*:



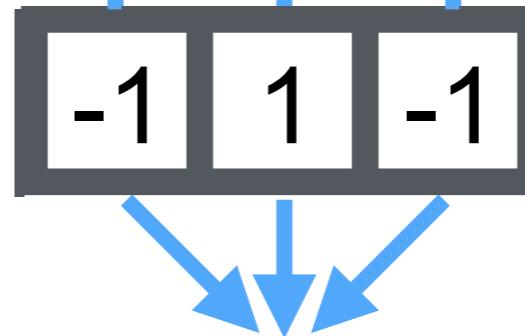
*correlation

Convolutional Layer: 1D example

A 1D image:



A filter:



After convolution*:



*correlation

Convolutional Layer: 1D example

A 1D image:



A filter:



After
convolution*:

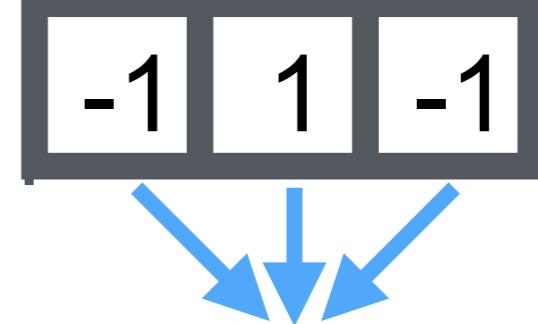


Convolutional Layer: 1D example

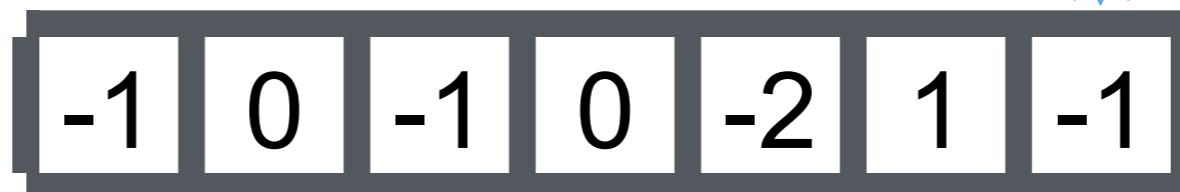
A 1D image:



A filter:



After convolution*:



*correlation

Convolutional Layer: 1D example

A 1D image:

0	0	1	1	1	0	1	0	0	0
---	---	---	---	---	---	---	---	---	---

A filter:

-1	1	-1
----	---	----

After
convolution*:

-1	0	-1	0	-2	1	-1	0
----	---	----	---	----	---	----	---

Convolutional Layer: 1D example

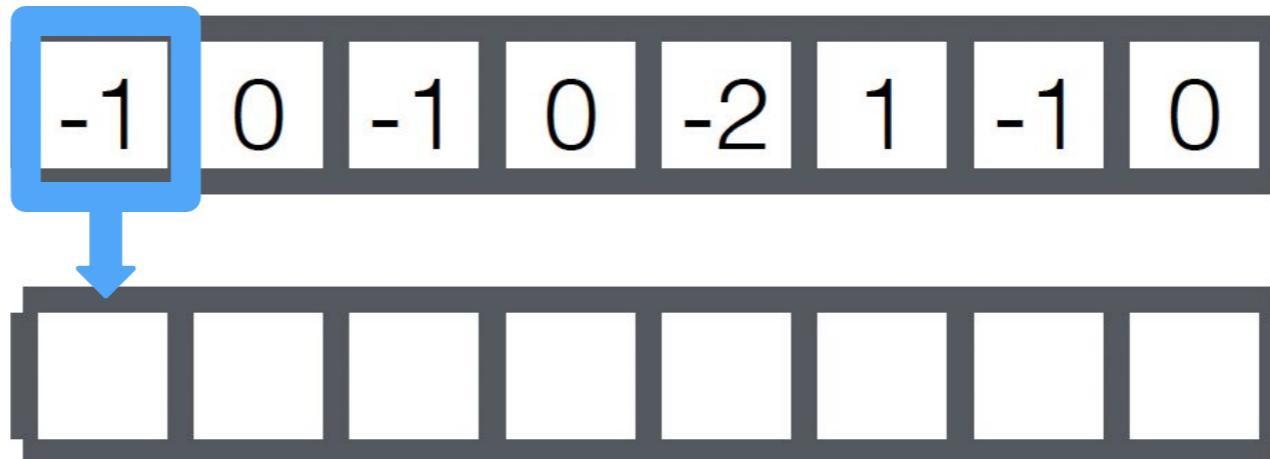
A 1D image:



A filter:



After convolution*:



After ReLU:

*correlation

Convolutional Layer: 1D example

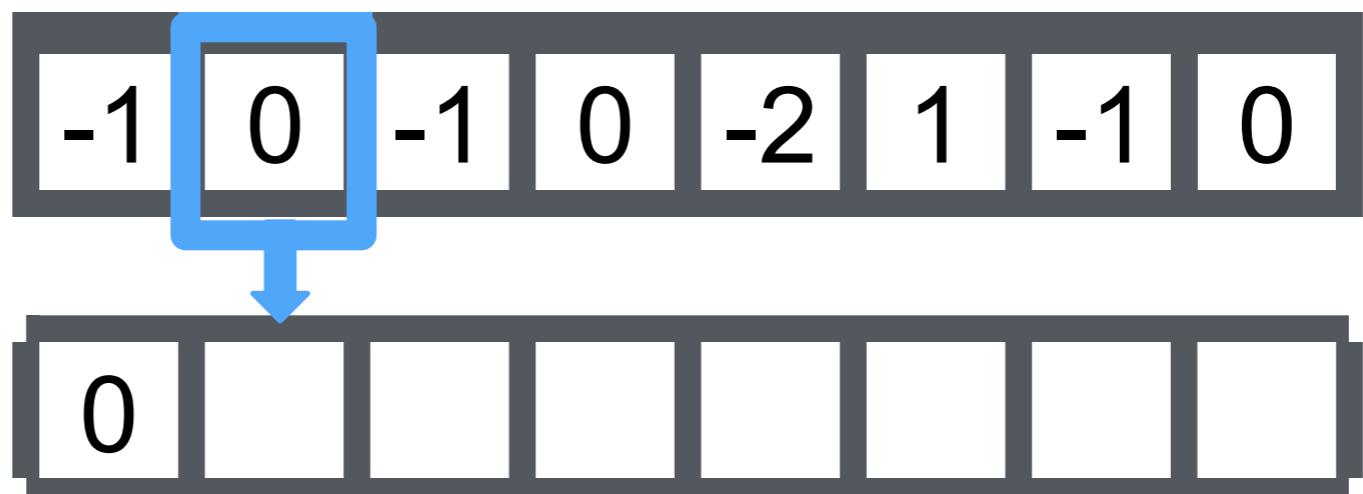
A 1D image:



A filter:



After convolution*:



After ReLU:

*correlation

Convolutional Layer: 1D example

A 1D image:

0	0	1	1	1	0	1	0	0	0
---	---	---	---	---	---	---	---	---	---

A filter:

-1	1	-1
----	---	----

After convolution*:

-1	0	-1	0	-2	1	-1	0
----	---	----	---	----	---	----	---

After ReLU:

0	0	0	0	0	1	0	0
---	---	---	---	---	---	---	---

**What does
the filter do?**

Convolutional Layer: 1D example

A 1D image:



A filter:



After convolution*: :



After ReLU:



*correlation

Convolutional Layer: 1D example

A 1D image:



A filter:



After convolution:

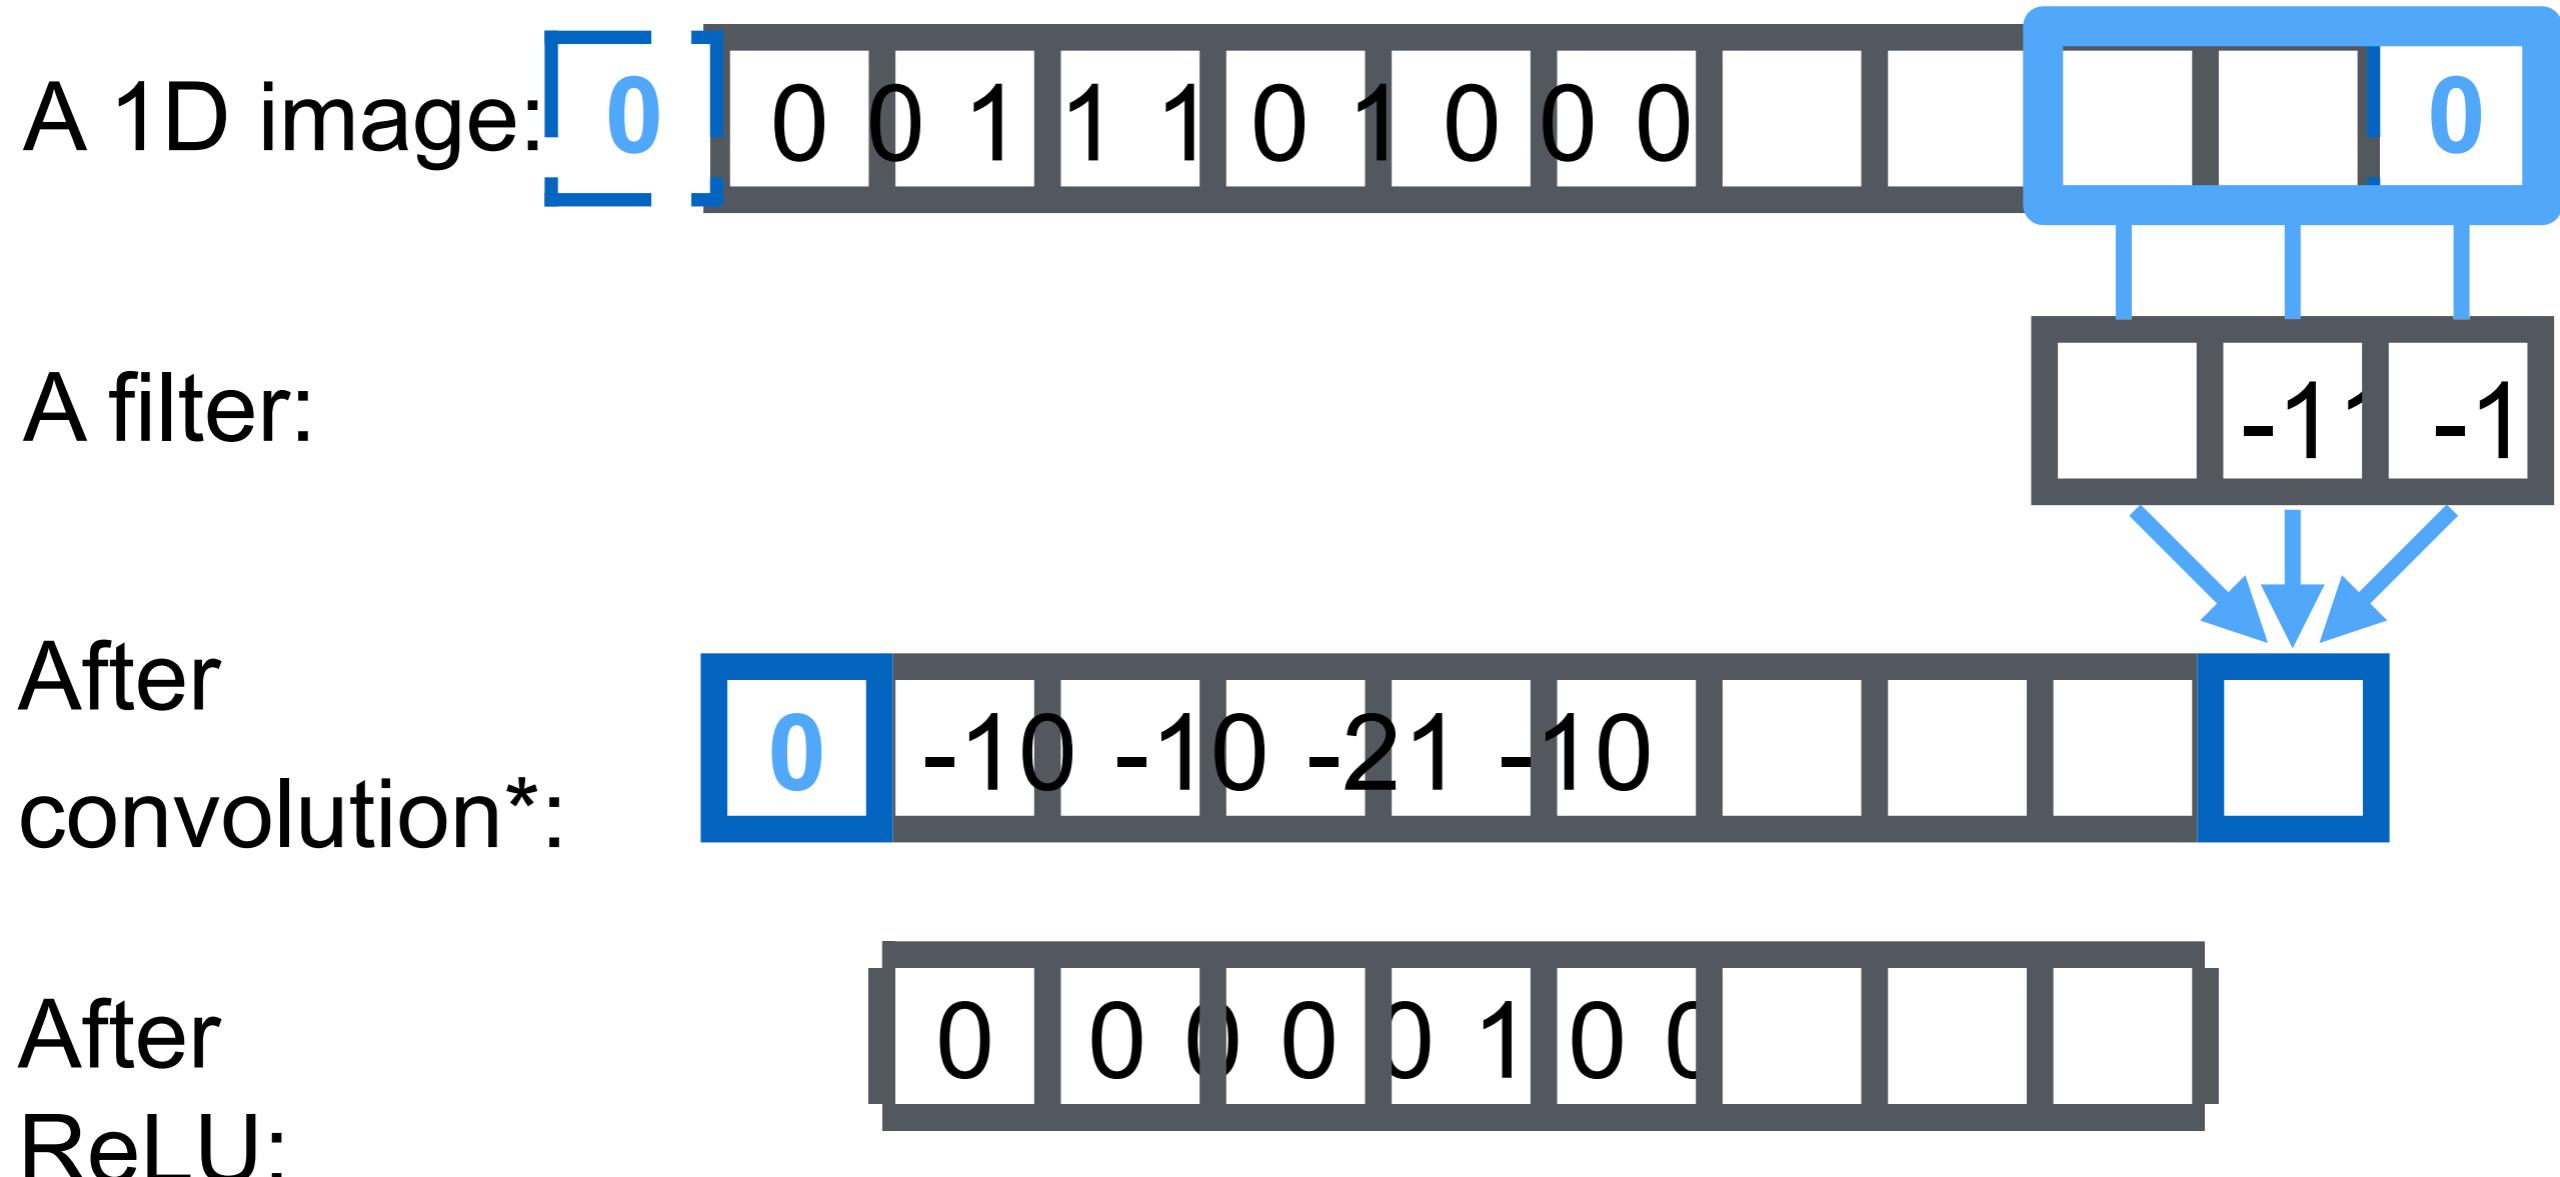


After ReLU:



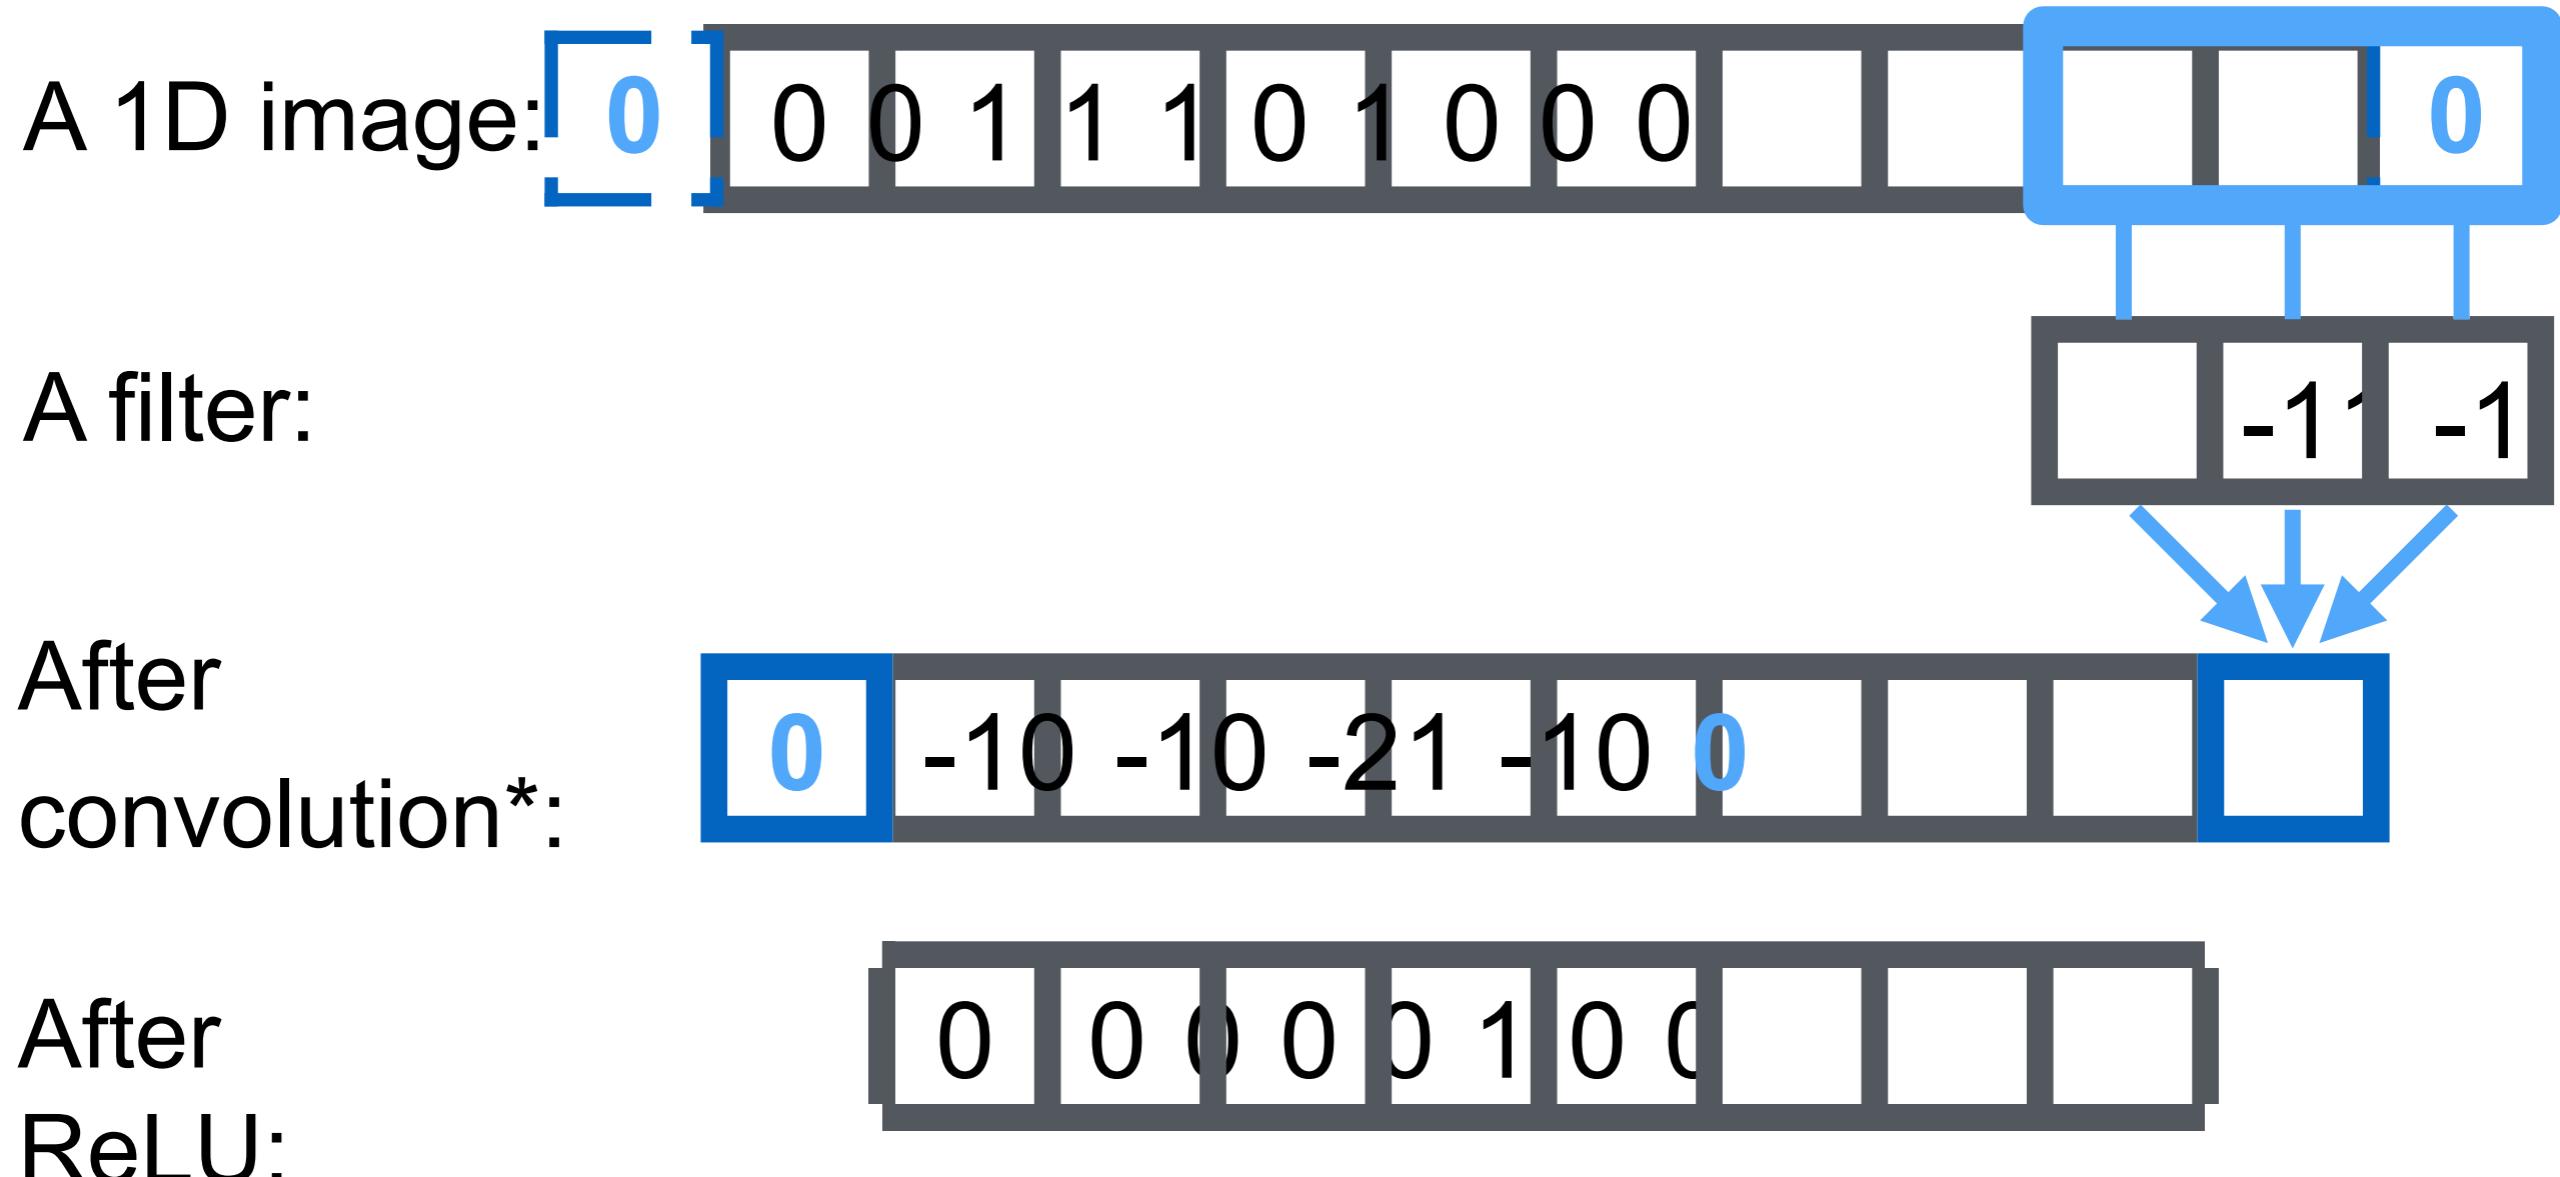
*correlation

Convolutional Layer: 1D example



*correlation

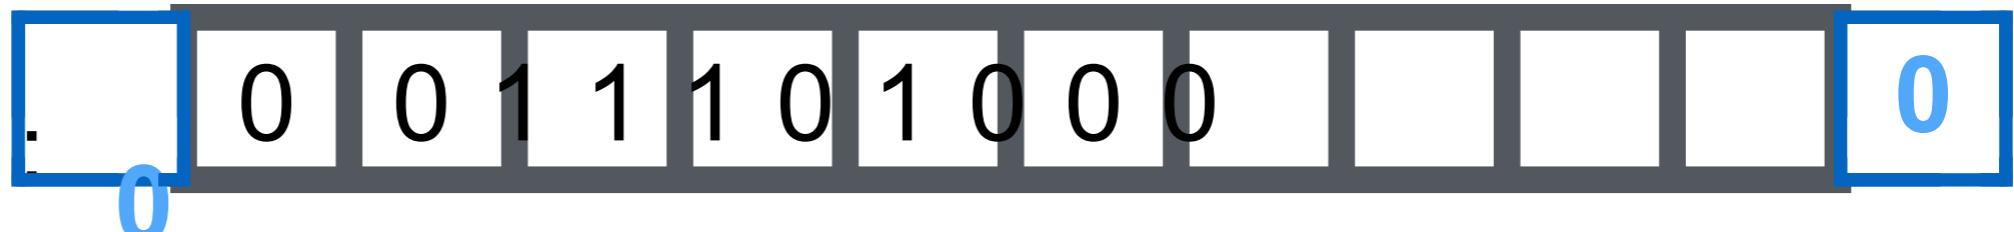
Convolutional Layer: 1D example



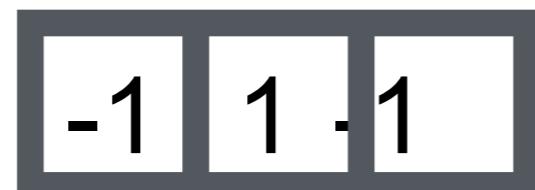
*correlation

Convolutional Layer: 1D example

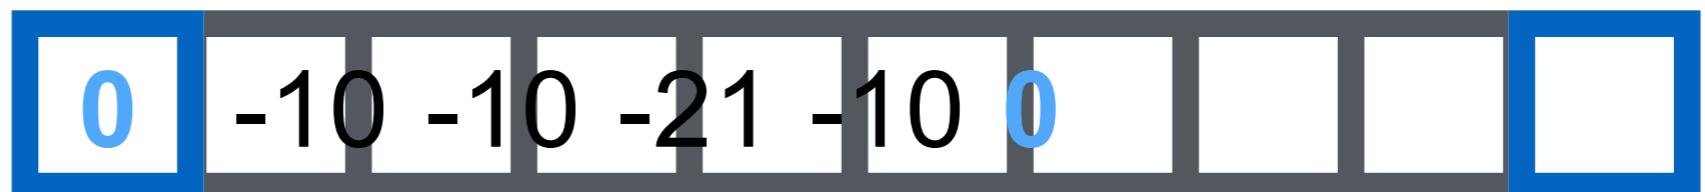
A 1D image



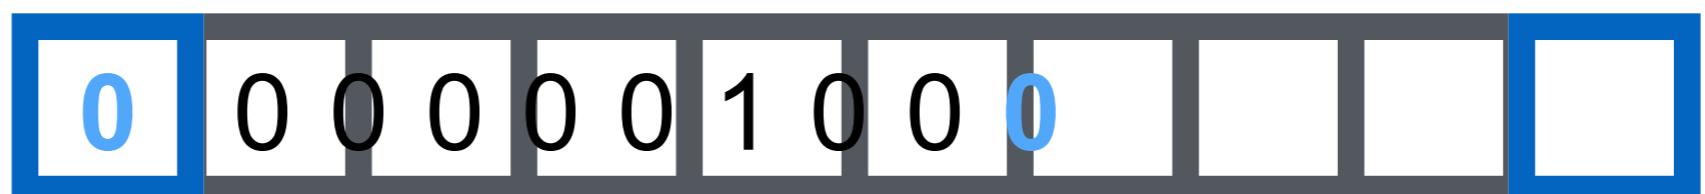
A filter:



After convolution*:



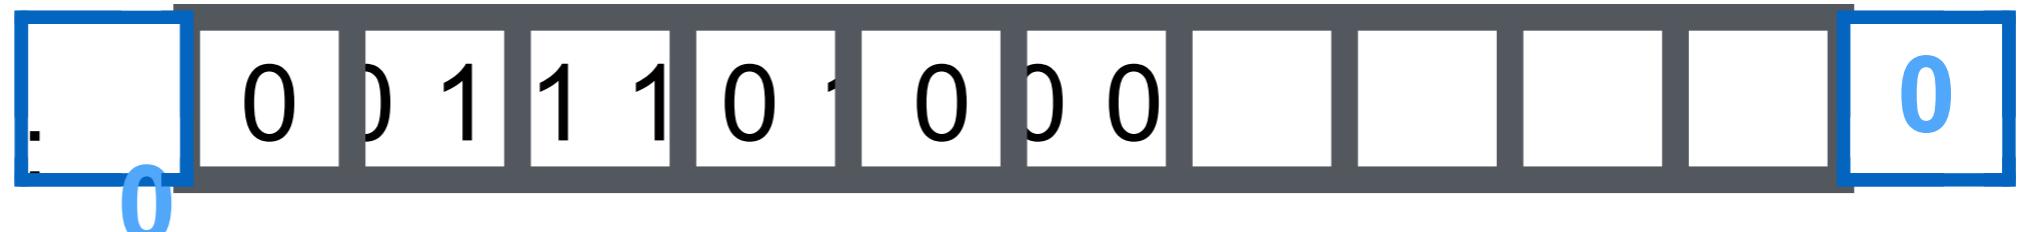
After ReLU:



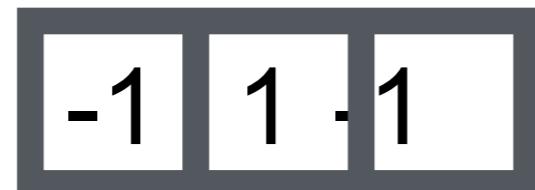
*correlation

Convolutional Layer: 1D example

A 1D image



A filter:



with bias +1

After
convolution*:

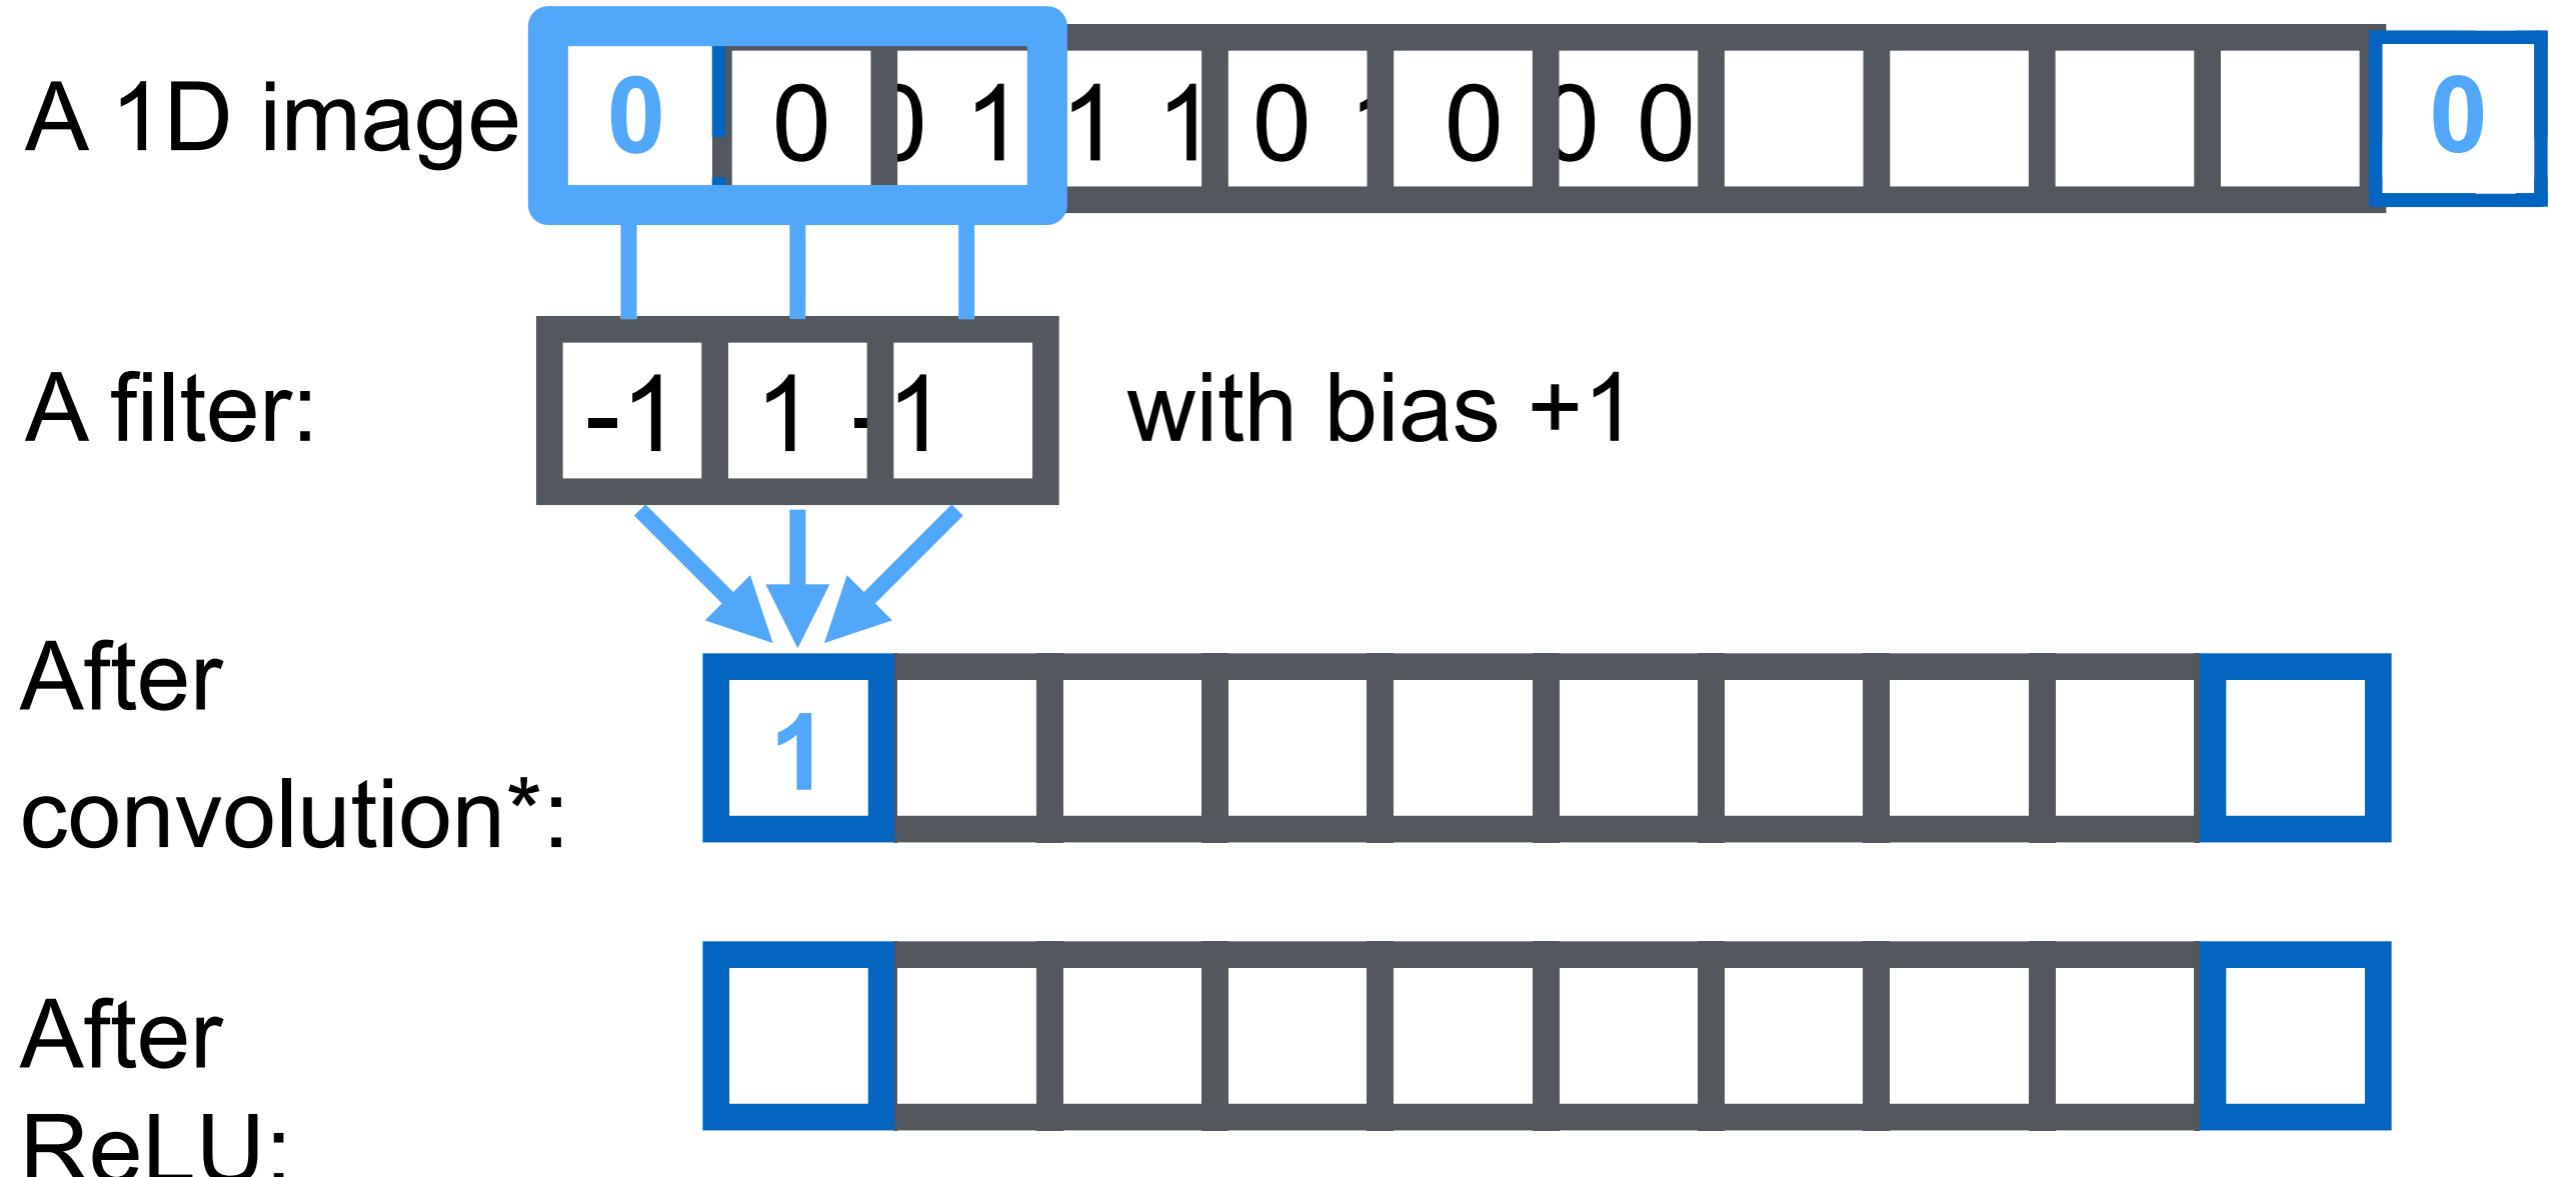


After ReLU:



*correlation

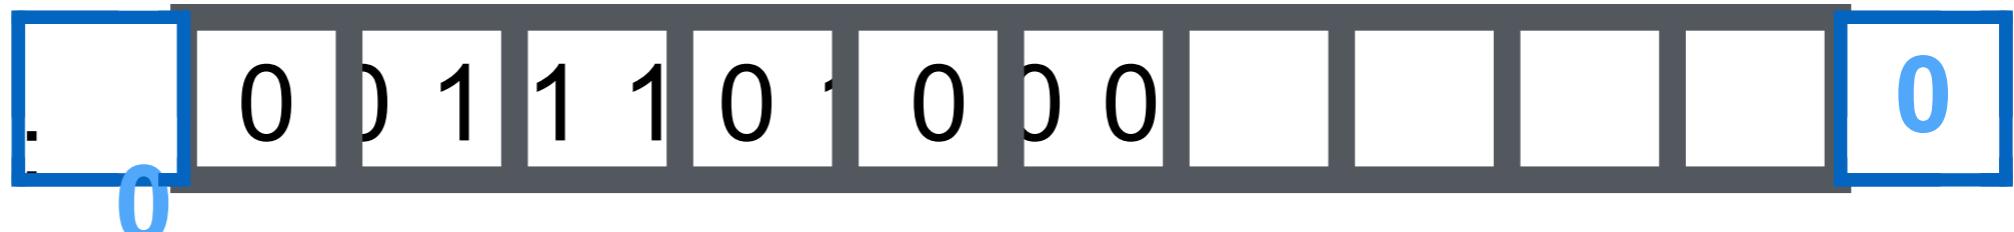
Convolutional Layer: 1D example



*correlation

Convolutional Layer: 1D example

A 1D image



A filter:



with bias +1

After
convolution*:



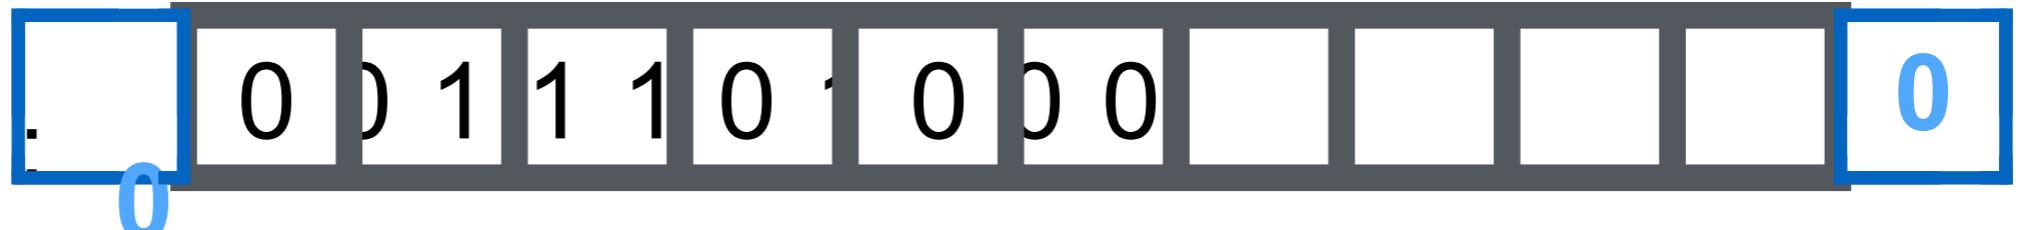
After
ReLU:



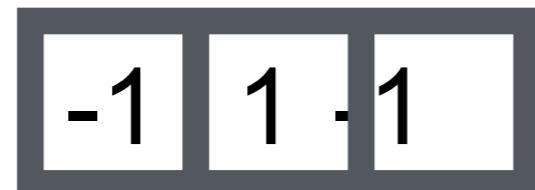
*correlation

Convolutional Layer: 1D example

A 1D image



A filter:



with bias +1

After convolution*:



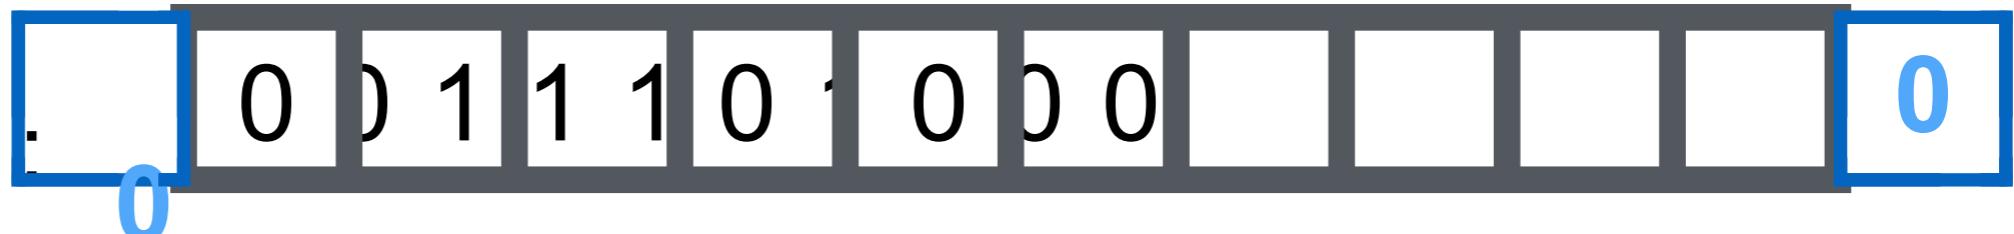
After ReLU:



*correlation

Convolutional Layer: 1D example

A 1D image



A filter:



with bias b

After
convolution*:



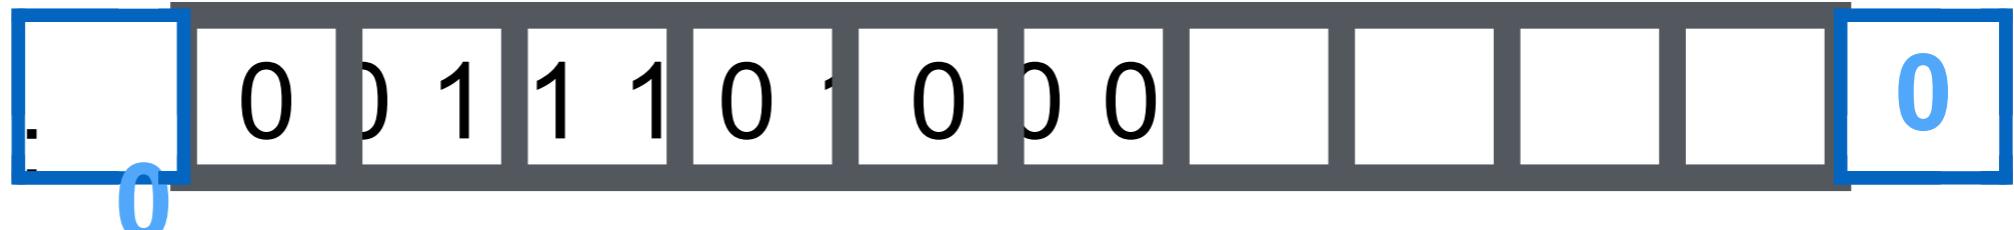
After ReLU:



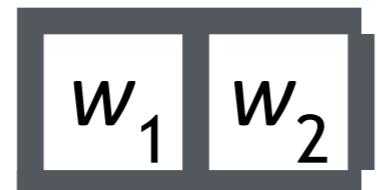
*correlation

Convolutional Layer: 1D example

A 1D image



A filter:



with bias b

After
convolution*:



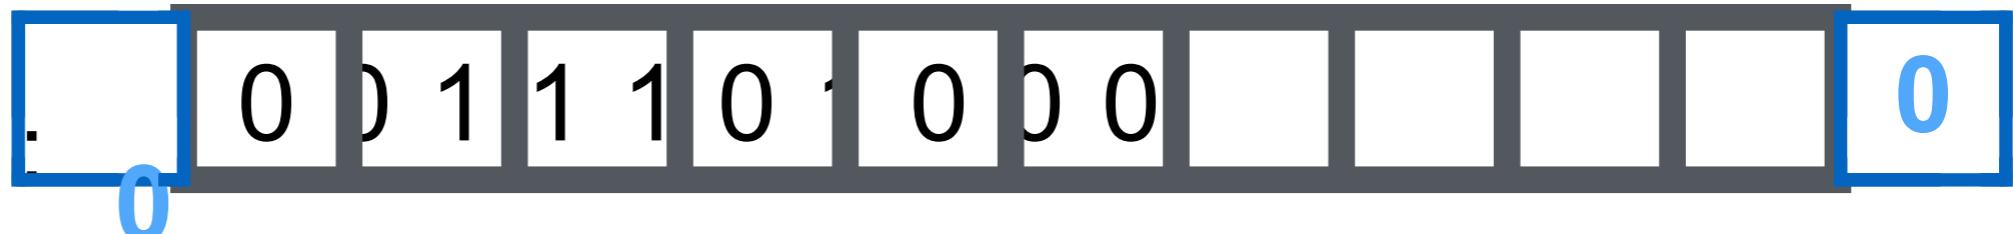
After ReLU:



*correlation

Convolutional Layer: 1D example

A 1D image



A filter:



with bias b

After
convolution*:



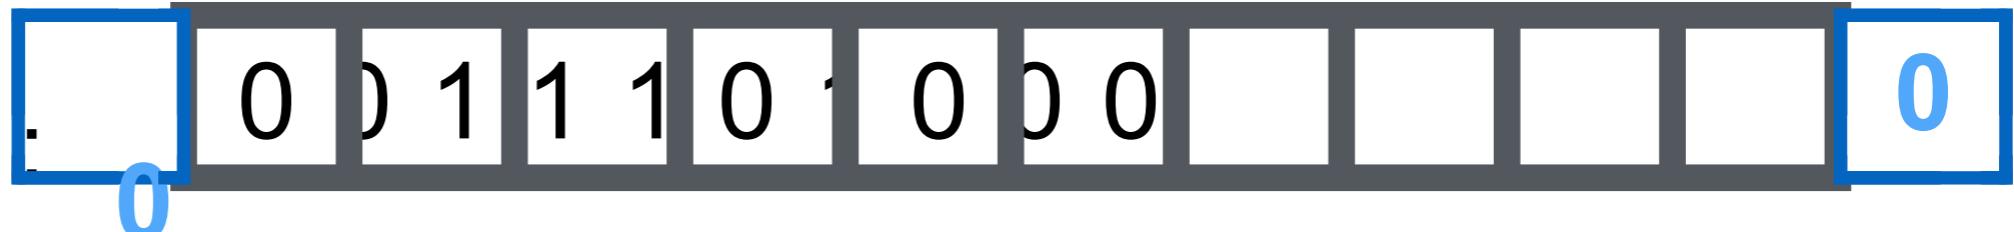
After ReLU:



*correlation

Convolutional Layer: 1D example

A 1D image



A filter:



with bias b

After convolution*:



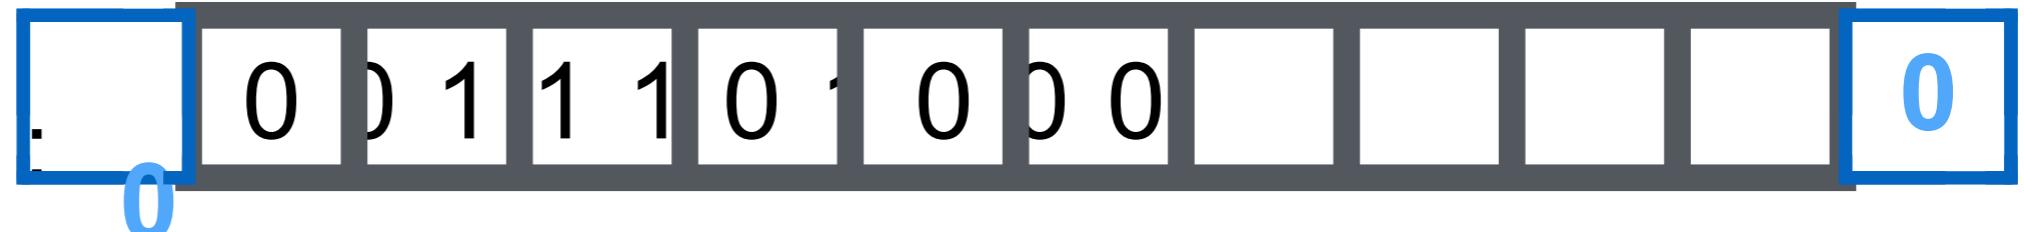
After ReLU:



- How many weights (including bias)? 4

Convolutional Layer: 1D example

A 1D image



A filter:



with bias b

After convolution*:



After ReLU:

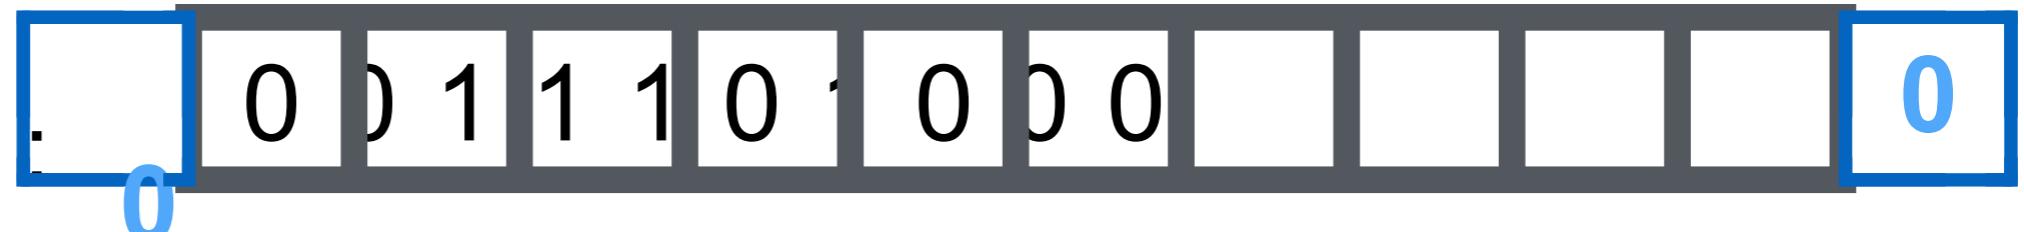


- How many weights (including bias)? 4
- How many weights (including biases) for fully connected layer with 10 inputs & 10 outputs?

*correlation

Convolutional Layer: 1D example

A 1D image



A filter:



with bias b

After

convolution*:



After ReLU:

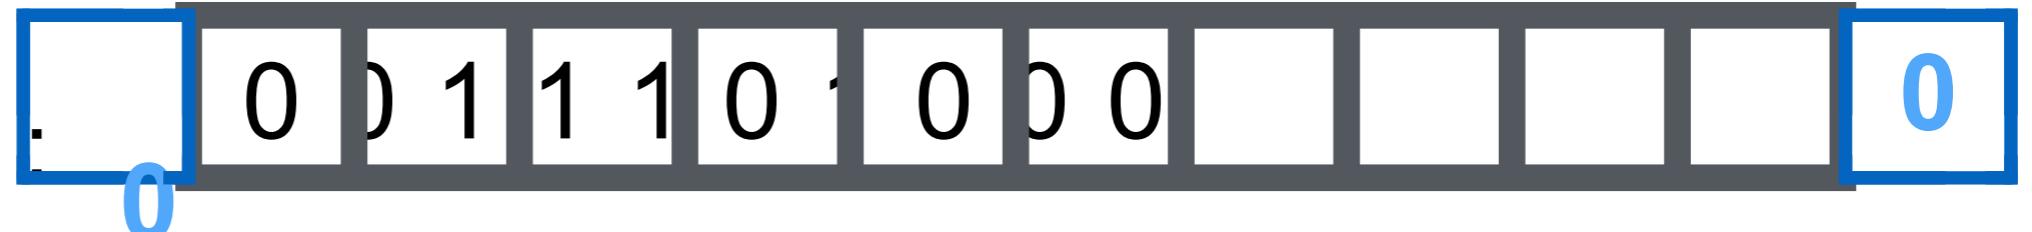


- How many weights (including bias)? 4
- How many weights (including biases) for fully connected layer with 10 inputs & 10 outputs? $10 \times 11 =$

*correlation

Convolutional Layer: 1D example

A 1D image



A filter:



with bias b

After
convolution*:



After ReLU:



- How many weights (including bias)? 4
- How many weights (including biases) for fully connected layer with 10 inputs & 10 outputs? $10 \times 11 = 110$

*correlation

Convolutional Layer: 2D example

A 2D
image:

1	0	1	0	0
1	0	1	0	1
1	1	1	0	0
1	0	1	0	1
1	0	1	0	1

Convolutional Layer: 2D example

A 2D
image:

1	0	1	0	0
1	0	1	0	1
1	1	1	0	0
1	0	1	0	1
1	0	1	0	1

Convolutional Layer: 2D example

A 2D
image:

1	0	1	0	0
1	0	1	0	1
1	1	1	0	0
1	0	1	0	1
1	0	1	0	1

A filter:

-1	-1	1
-1	1	-1
-1	-1	1

After
convolution:

Convolutional Layer: 2D example

A 2D
image:

1	0	1	0	0
1	0	1	0	1
1	1	1	0	0
1	0	1	0	1
1	0	1	0	1

A filter:

-1	-1	-1
-1	1	-1
-1	-1	1

After
convolution:

Convolutional Layer: 2D example

A 2D
image:

1	0	1	0	0
1	0	1	0	1
1	1	1	0	0
1	0	1	0	1
1	0	1	0	1

A filter:

-1	-1	1
-1	1	-1
-1	-1	1

After
convolution:



Convolutional Layer: 2D example

A 2D
image:

1	0	1	0	0
1	0	1	0	1
1	1	1	0	0
1	0	1	0	1
1	0	1	0	1

A filter:

-1	-1	1
-1	1	-1
-1	-1	1

After
convolution:



Convolutional Layer: 2D example

A 2D
image:

1		0	1	0	0
		1	0	1	0
		1	0	1	1
		1	1	1	0
		1	0	1	0
1	0	1	0	1	

A filter:

-1	-1	1
-1	1	-1
-1	-1	-1

-1

After
convolution:



Convolutional Layer: 2D example

A 2D
image:

1	0		1	0	0
			1	0	1
			1	0	1
			1	1	0
			1	0	1
1	0	1	0	1	

A filter:

-1	-1	-1
-1	1	-1
-1	-1	-1

-1

After
convolution:



Convolutional Layer: 2D example

A 2D
image:

1	0	1	0	0
1	0	1	0	1
1	1	1	0	0
1	0	1	0	1
1	0	1	0	1

A filter:

-1	-1	-1
-1	1	-1
-1	-1	-1

$$-1 + 0$$

After
convolution:



Convolutional Layer: 2D example

A 2D
image:

1	0	1	0	0
1	0	1	0	1
1	1	1	0	0
1	0	1	0	1
1	0	1	0	1

$$-1 + 0 + -1$$

A filter:

-1	-1	-1
-1	1	-1
-1	-1	-1

After
convolution:



Convolutional Layer: 2D example

A 2D
image:

1	0	1	0	0
1	0	1	0	1
1	1	1	0	0
1	0	1	0	1
1	0	1	0	1

A filter:

-1	-1	-1
-1	1	1
-1	-1	-1

$$\begin{aligned}-1 + 0 + -1 \\ + -1\end{aligned}$$

After
convolution:



Convolutional Layer: 2D example

A 2D
image:

1	0	1	0	0
1	0	1	0	1
1	1	1	0	0
1	0	1	0	1
1	0	1	0	1

A filter:

-1	-1	1
-1	1	-1
-1	-1	1

$$\begin{aligned}-1 + 0 + -1 \\ + -1 + 0\end{aligned}$$

After
convolution:



Convolutional Layer: 2D example

A 2D
image:

1	0	1	0	0
1	0	1	0	1
1	1	1	0	0
1	0	1	0	1
1	0	1	0	1

A filter:

-1	-1	1
-1	1	-1
-1	-1	1

$$\begin{aligned} & -1 + 0 + -1 \\ & + -1 + 0 + -1 \end{aligned}$$

After
convolution:



Convolutional Layer: 2D example

A 2D
image:

1	0	1	0	0
1	0	1	0	1
1	1	1	0	0
1	0	1	0	1
1	0	1	0	1

A filter:

-1	-1	1
-1	1	1
-1	-1	1

$$\begin{aligned} & -1 + 0 + -1 \\ & + -1 + 0 + -1 \\ & + -1 \end{aligned}$$

After
convolution:



Convolutional Layer: 2D example

A 2D
image:

1	0	1	0	0
1	0	1	0	1
1	1	1	0	0
1	0	1	0	1
1	0	1	0	1

A filter:

-1	-1	1
-1	1	1
-1	-1	-1

$$\begin{aligned} & -1 + 0 + -1 \\ & + -1 + 0 + -1 \\ & + -1 + -1 \end{aligned}$$

After
convolution:



Convolutional Layer: 2D example

A 2D
image:

1	0	1	0	0
1	0	1	0	1
1	1	1	0	0
1	0	1	0	1
1	0	1	0	1

A filter:

-1	-1	1
-1	1	-1
-1	-1	-1

$$\begin{aligned} & -1 + 0 + -1 \\ & + -1 + 0 + -1 \\ & + -1 + -1 + -1 \end{aligned}$$

After
convolution:



Convolutional Layer: 2D example

A 2D
image:

1	0	1	0	0
1	0	1	0	1
1	1	1	0	0
1	0	1	0	1
1	0	1	0	1

A filter:

-1	-1	1
-1	1	-1
-1	-1	-1

$$\begin{aligned} & -1 + 0 + -1 \\ & + -1 + 0 + -1 \\ & + -1 + -1 + -1 \\ & = -7 \end{aligned}$$

After
convolution:



Convolutional Layer: 2D example

A 2D
image:

1	0	1	0	0
1	0	1	0	1
1	1	1	0	0
1	0	1	0	1
1	0	1	0	1

A filter:

-1	-1	1
-1	1	-1
-1	-1	1

$$\begin{aligned}-1 + 0 + -1 \\ + -1 + 0 + -1 \\ + -1 + -1 + -1 \\ = -7\end{aligned}$$

After
convolution:

-7

Convolutional Layer: 2D example

A 2D
image:

1	0	1	0	0
1	0	1	0	1
1	1	1	0	0
1	0	1	0	1
1	0	1	0	1

A filter:

-1	-1	1
-1	1	-1
-1	-1	1

After
convolution:



Convolutional Layer: 2D example

A 2D
image:

1	0	1	0	0
1	0	1	0	1
1	1	1	0	0
1	0	1	0	1
1	0	1	0	1

A filter:

-1	-1	1
-1	1	1
-1	-1	1

After
convolution:

-7	-2
----	----

Convolutional Layer: 2D example

A 2D
image:

1	0	1	0	0
1	0	1	0	1
1	1	1	0	0
1	0	1	0	1
1	0	1	0	1

A filter:

-1	-1	1
-1	1	1
-1	-1	1

-7	-2	4
----	----	---

After
convolution:

Convolutional Layer: 2D example

A 2D
image:

1	0	1	0	0
1	0	1	0	1
1	1	1	0	0
1	0	1	0	1
1	0	1	0	1

-1	1	-1	
-1	1	-1	
-1	-1	-1	

After
convolution:

-7	-2	-4	
-5			

Convolutional Layer: 2D example

A 2D
image:

1	0	1	0	0
1	0	1	0	1
1	1	1	0	0
1	0	1	0	1
1	0	1	0	1

A filter:

-1	1	-1	
-1	1	-1	
-1	-1	-1	

After
convolution:

-7	-2	-4	
-5			
-2			

Convolutional Layer: 2D example

A 2D
image:

1	0	1	0	0		
1	0	1	0	1		
1	1	1	0	0		
1	0	1	0	1		
1	0	1	0	1		

A filter:

-1	-1	1
-1	1	1
-1	-1	1

After
convolution:

-7	-2	4
-5	-2	5

Convolutional Layer: 2D example

A 2D
image:

1	0	1	0	0
1	0	1	0	1
1	1	1	0	0
1	0	1	0	1
1	0	1	0	1

A filter:

-1	1	-1
-1	1	-1
-1	-1	-1

After
convolution:

-7	-2	-4
-5	-2	5
-7		

Convolutional Layer: 2D example

A 2D
image:

1	0	1	0	0
1	0	1	0	1
1	1	1	0	0
1	0	1	0	1
1	0	1	0	1

-1	-1	-1
-1	1	-1
-1	-1	-1

After
convolution:

-7	-2	-4
-5	-2	-5
-7	-2	

Convolutional Layer: 2D example

A 2D
image:

1	0	1	0	0
1	0	1	0	1
1	1	1	0	0
1	0	1	0	1
1	0	1	0	1

-1	-1	-1
-1	1	-1
-1	-1	-1

After
convolution:

-7	-2	-4
-5	-2	-5
-7	2	-5

Convolutional Layer: 2D example

A 2D
image:

1	0	0	0	
1	0	0	1	
1	1	0	0	
1	0	0	1	
1	0	0	1	

A filter:

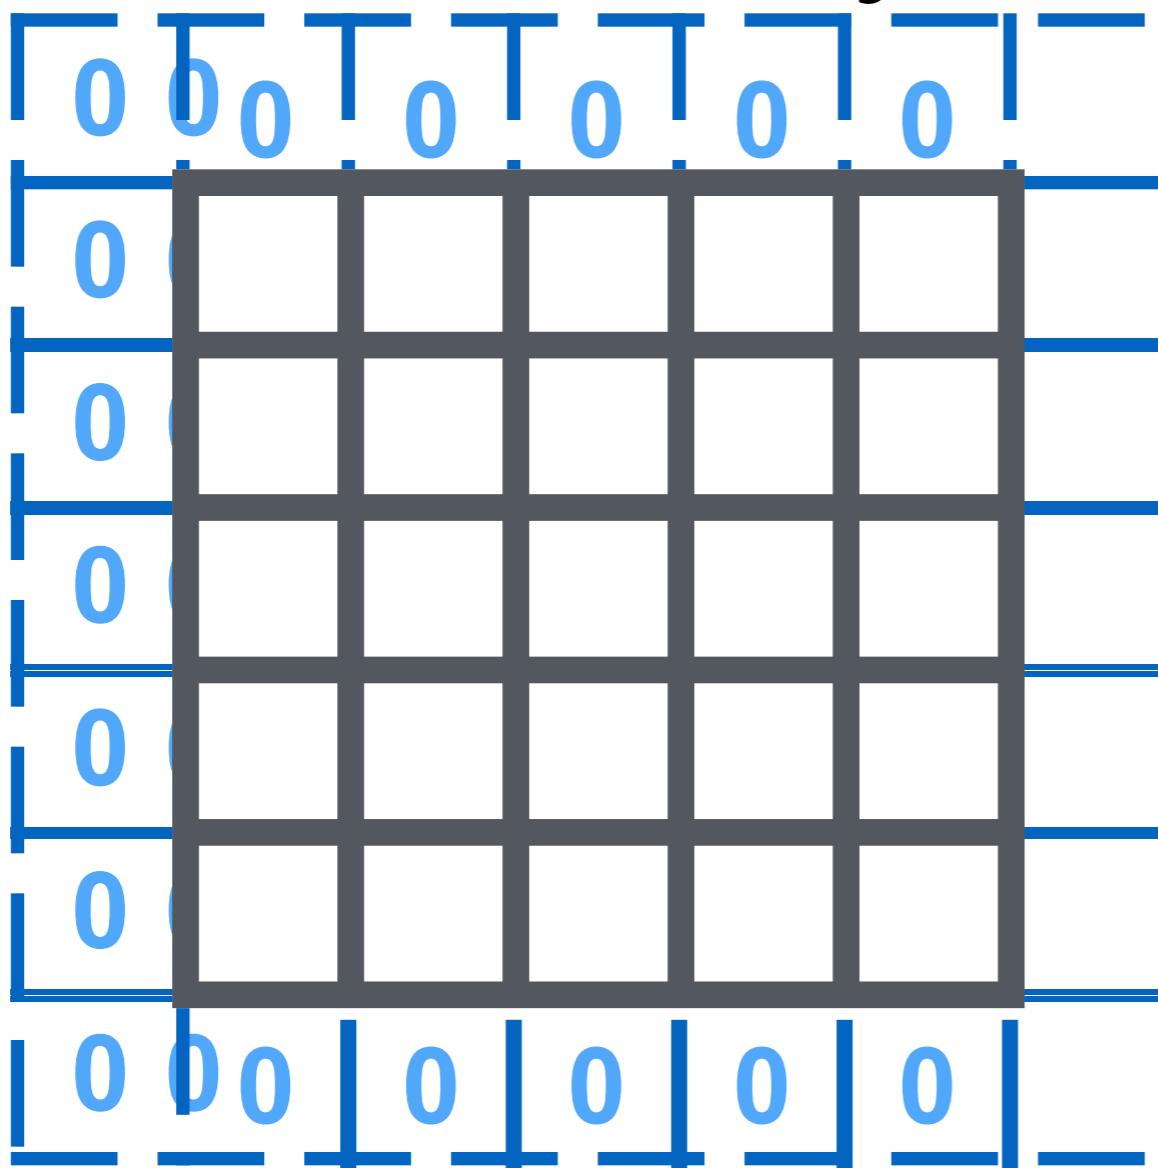
-1	-1	-1
-1	1	-1
-1	-1	-1

After
convolution:

-7	-2	-4
-5	-2	-5
-7	-2	-5

Convolutional Layer: 2D example

A 2D
image:



A filter:

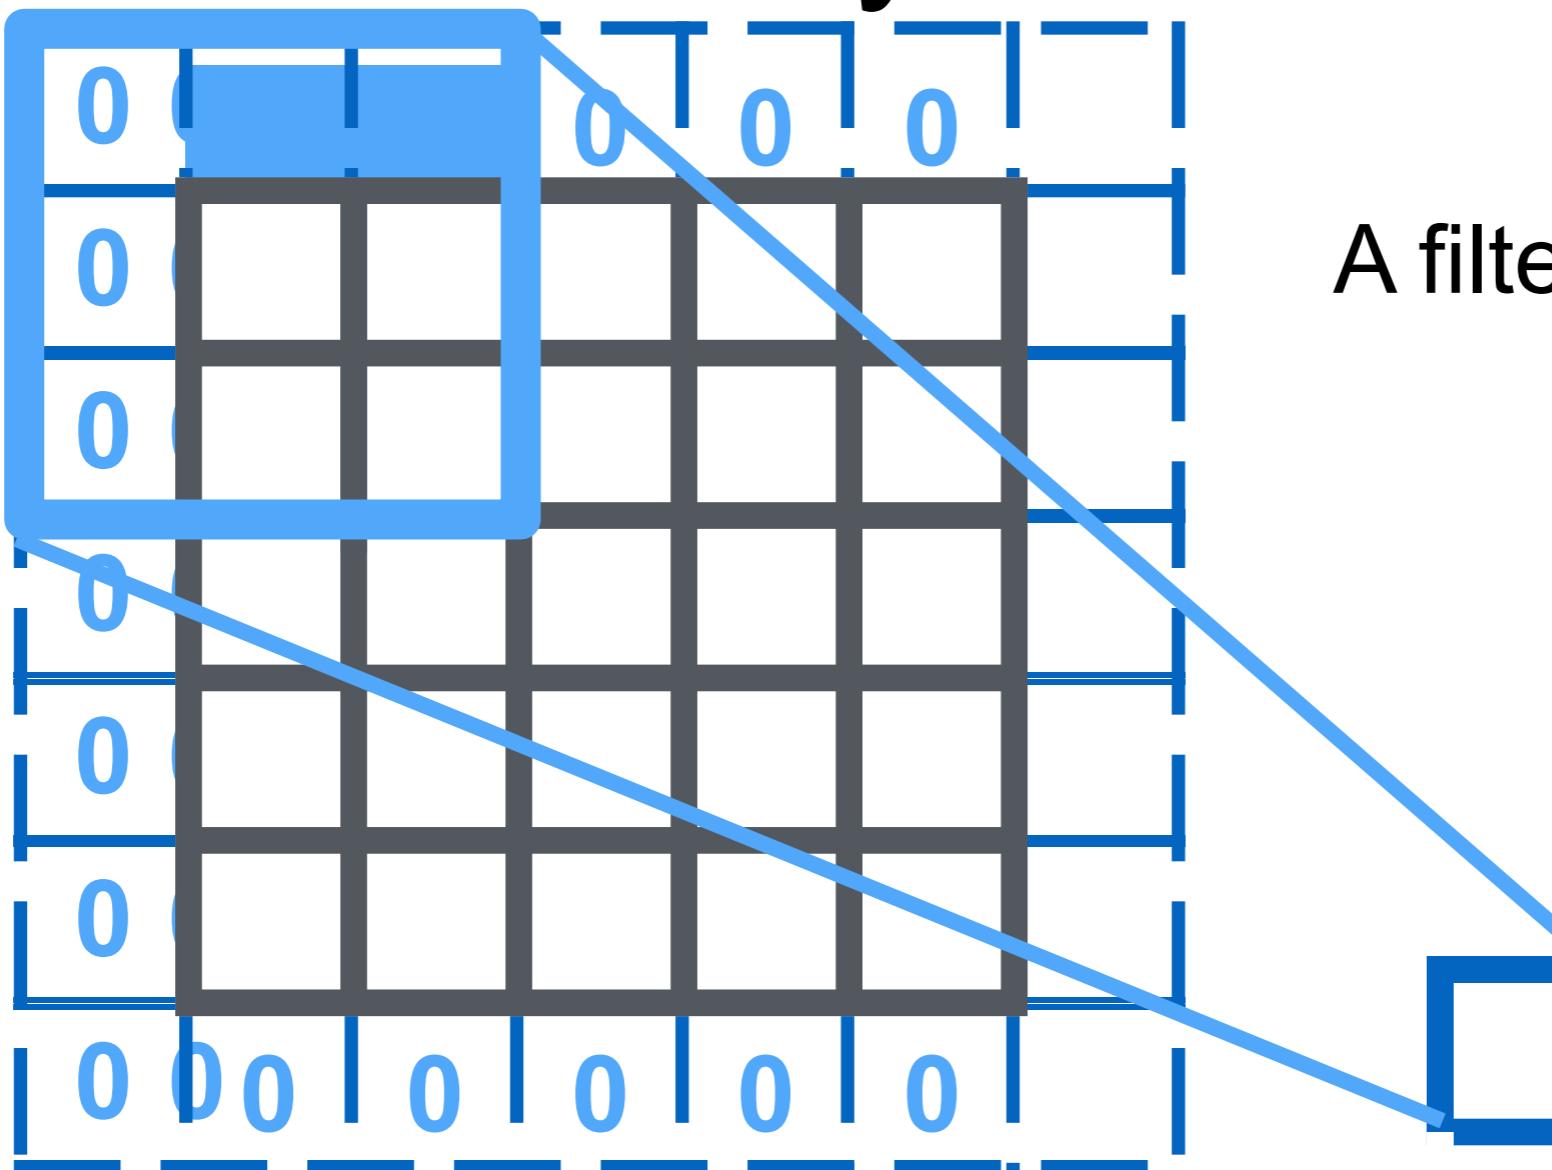
-1	-1	1
-1	1	1
-1	-1	1

After
convolution:

	-7	2	-4
-5	-2	5	
-7	-2	5	

Convolutional Layer: 2D example

A 2D
image:



A filter:

-1	-1	1
-1	1	1
-1	-1	1

After
convolution:

-7	2	-4
-5	-2	5
-7	-2	5

Convolutional Layer: 2D example

A 2D image:

0	0	0	0	0	0	0
0	1	0	1	0	0	0
0	1	0	1	0	1	0
0	1	1	1	0	0	0
0	1	0	1	0	1	0
0	1	0	1	0	1	0
0	0	0	0	0	0	0

A filter:

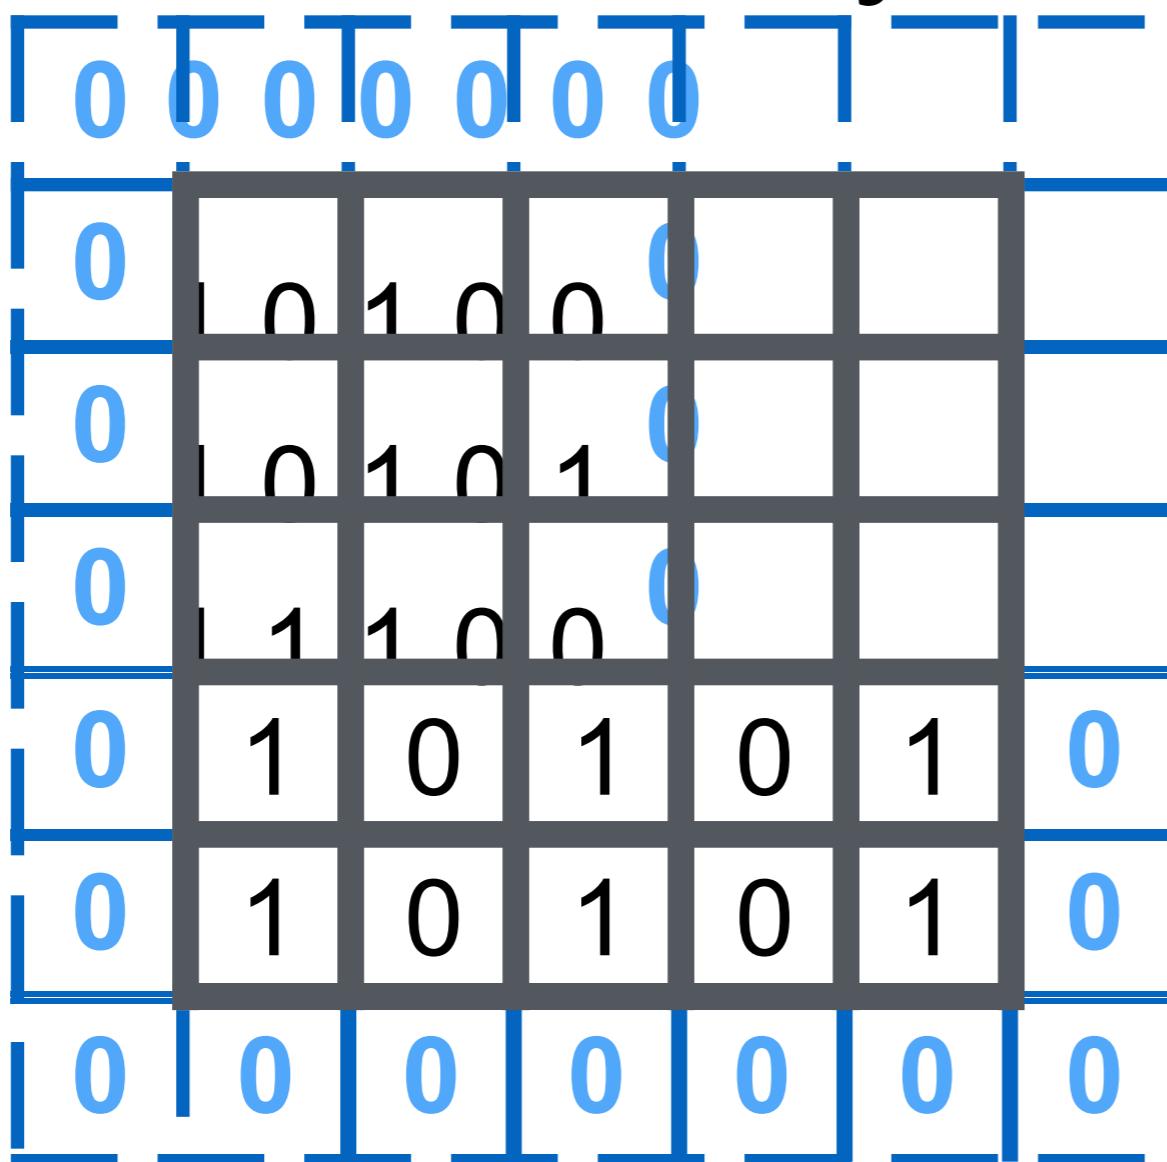
-1	-1	1
-1	1	-1
-1	-1	-1

After
convolution:

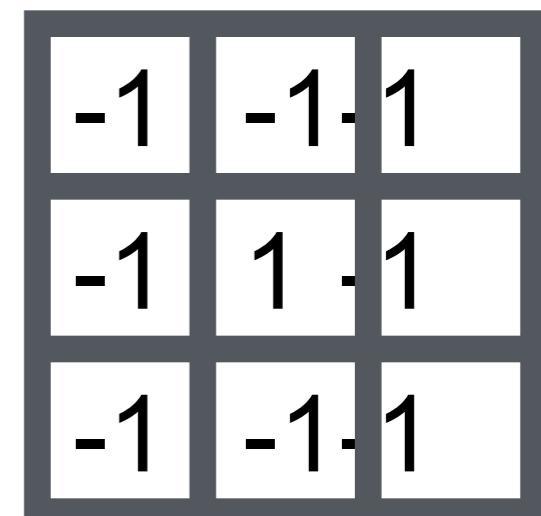
-7	-2	-4
-5	-2	-5
-7	-2	-5

Convolutional Layer: 2D example

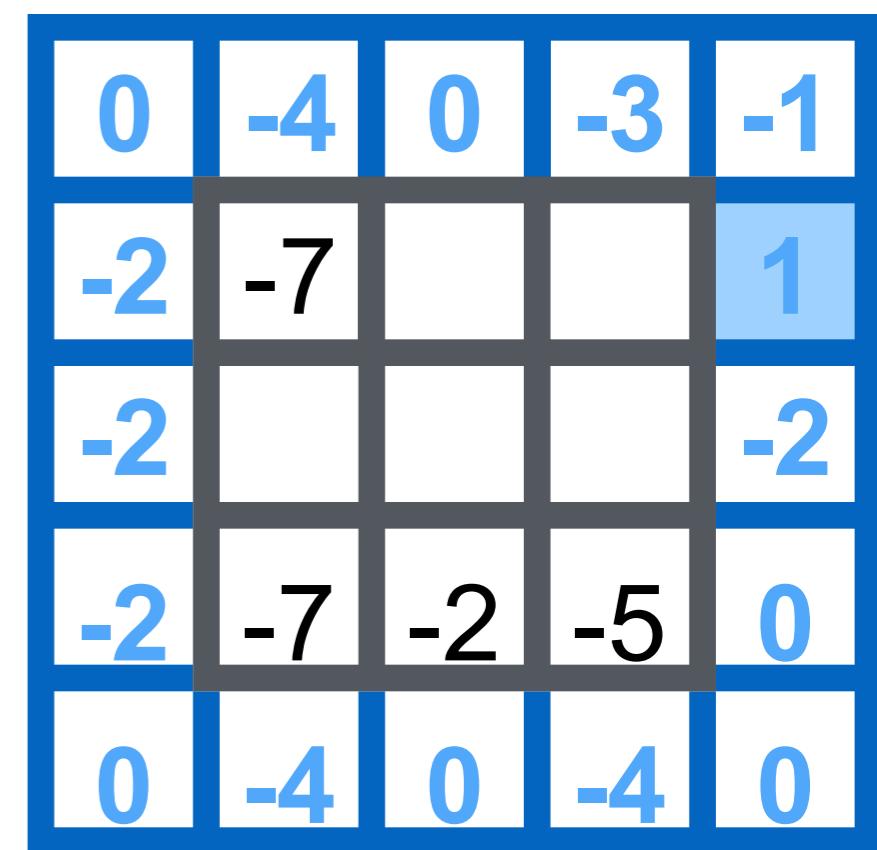
A 2D
image:



A filter



After
convolution
& ReLU:



Convolutional Layer: 2D example

A 2D image:

0	0	0	0	0	0	0
0	1	0	0	0	0	0
0	1	0	1	0	0	0
0	1	1	0	0	0	0
0	1	0	1	0	1	0
0	1	0	1	0	1	0
0	0	0	0	0	0	0

A filter:

-1	-1	1
-1	1	-1
-1	-1	1

Detecting isolated 1's.

How do we detect all the 1's?

-- Filter - 1 surrounded by 0's.

How do we detect edges?

-- Exercise?

After convolution & ReLU:

0	0	0	0	0
0	0			1
0				0
0				0
0	0	0	0	0

Convolutional Layer: 2D example

A 2D image:

0	0	0	0	0	0	0
0	1	0	1	0	0	0
0	1	0	1	0	1	0
0	1	1	1	0	0	0
0	1	0	1	0	1	0
0	1	0	1	0	1	0
0	0	0	0	0	0	0

A filter:

-1	-1	1
-1	1	-1
-1	-1	1

Detecting isolated 1's.

How do we detect all the 1's?

-- Filter - 1 surrounded by 0's.

How do we detect edges?

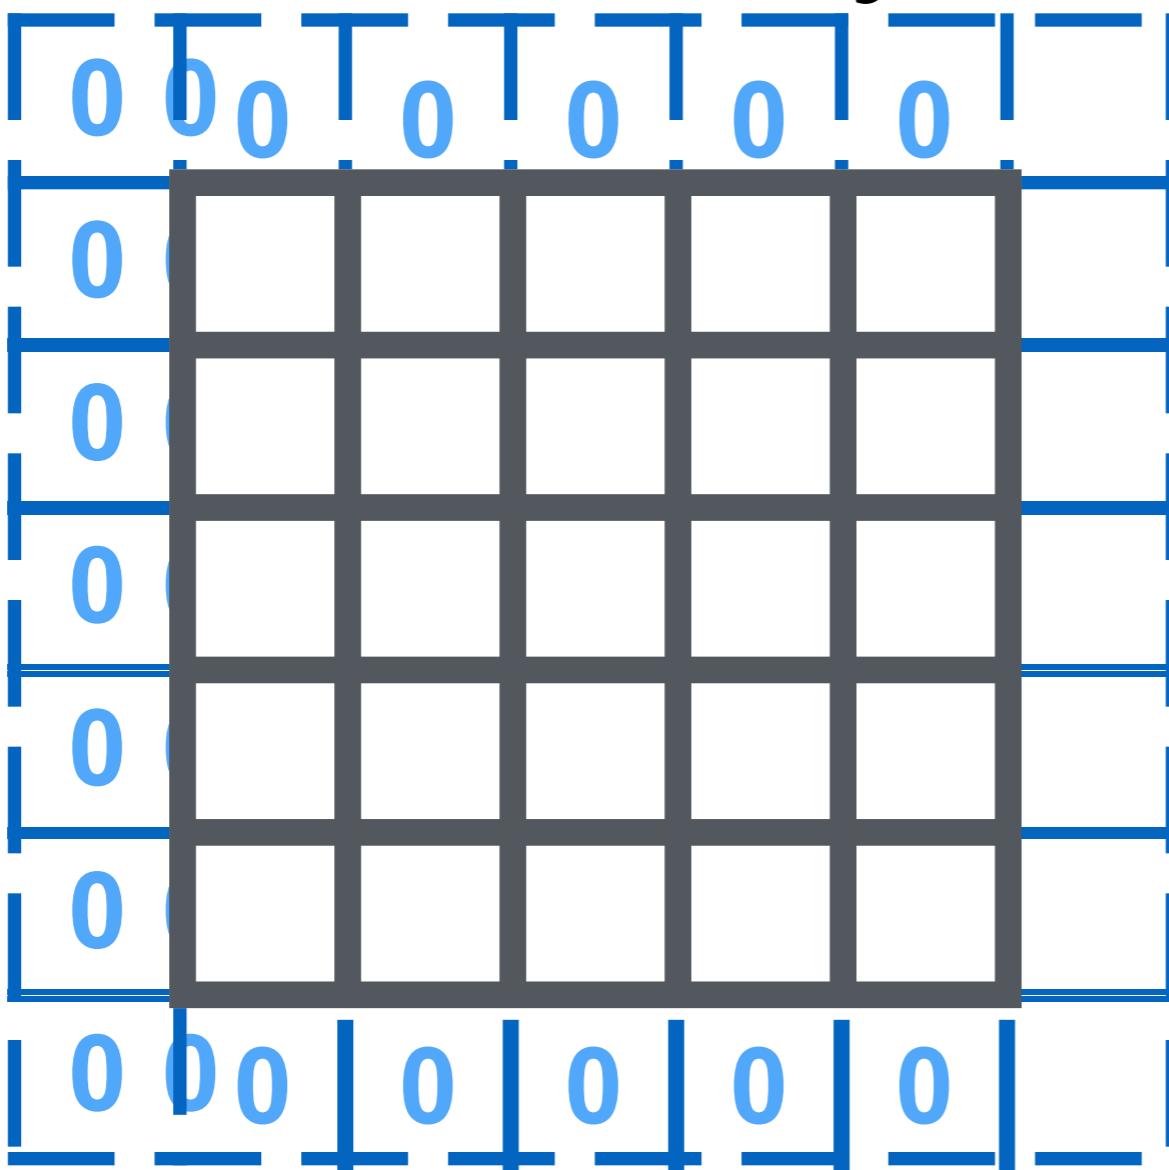
-- Exercise?

After convolution & ReLU:

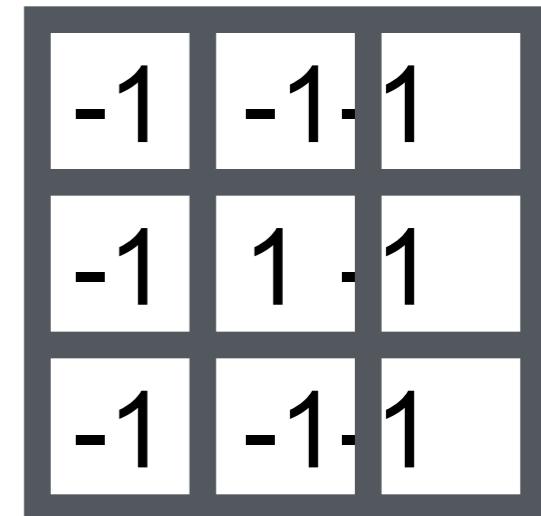
0	0	0	0	0
0	0			1
0				0
0				0
0	0	0	0	0

Convolutional Layer: 2D example

A 2D
image:

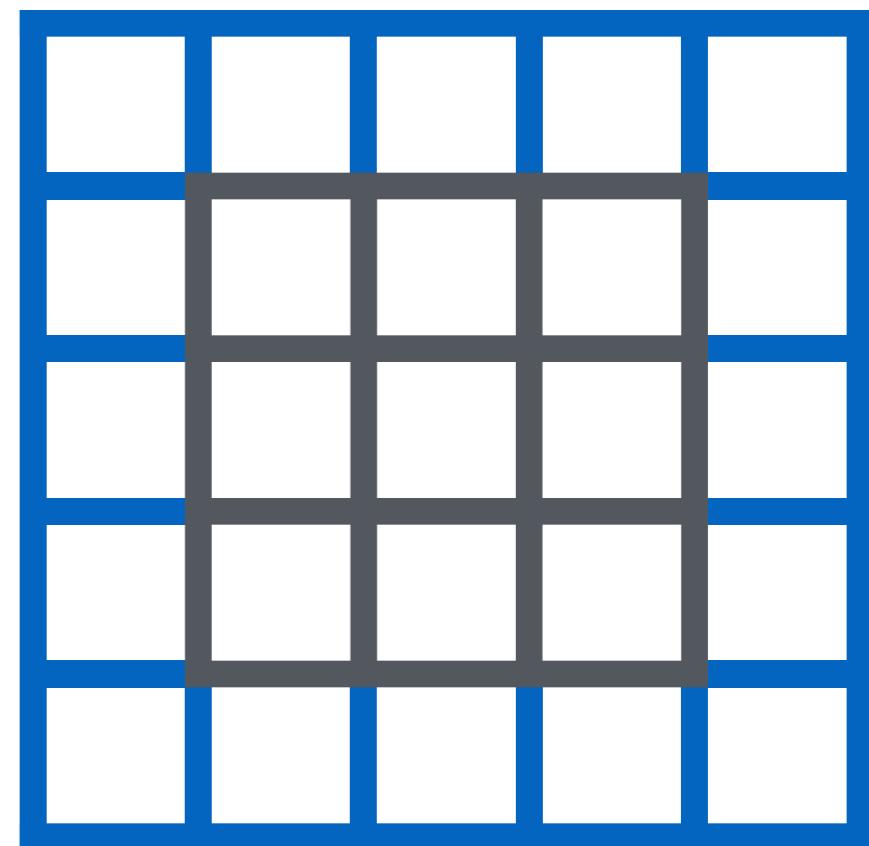


A filter:



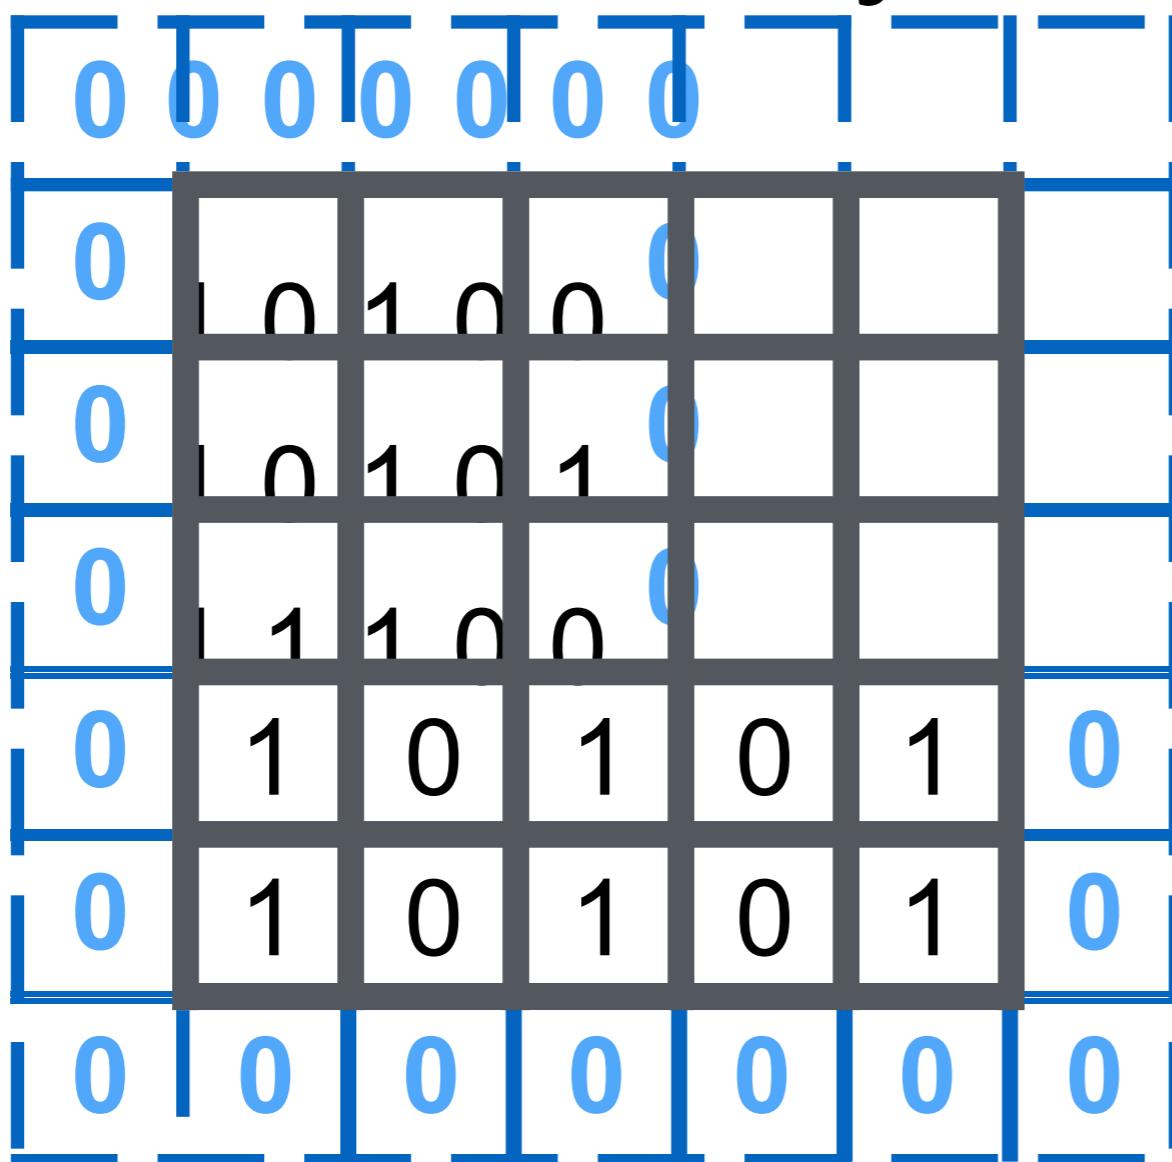
with bias 2

After
convolution:

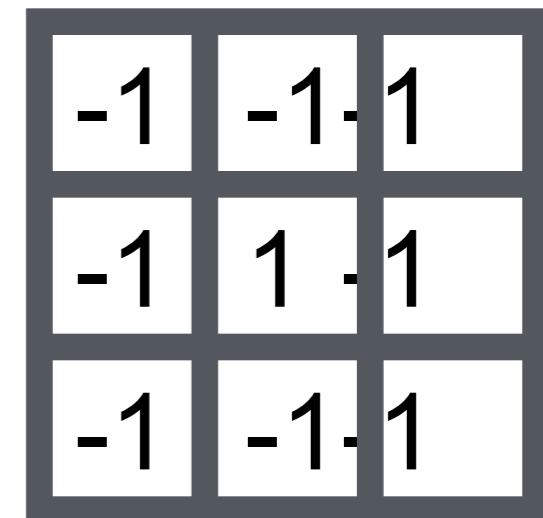


Convolutional Layer: 2D example

A 2D
image:



A filter:

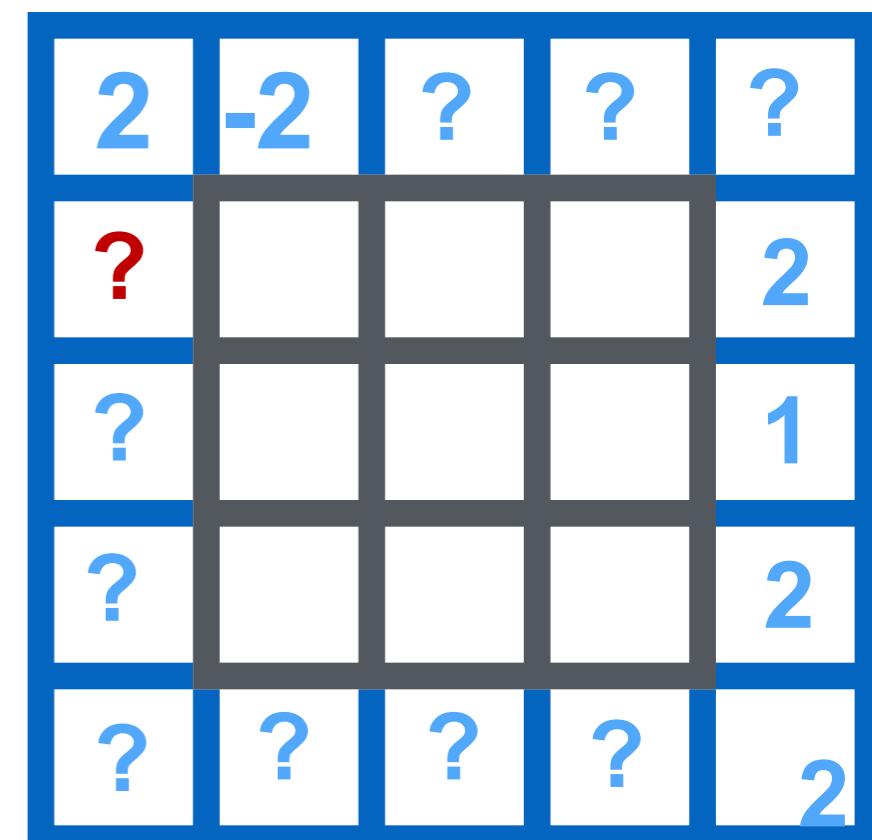


with bias 2

What does it detect?

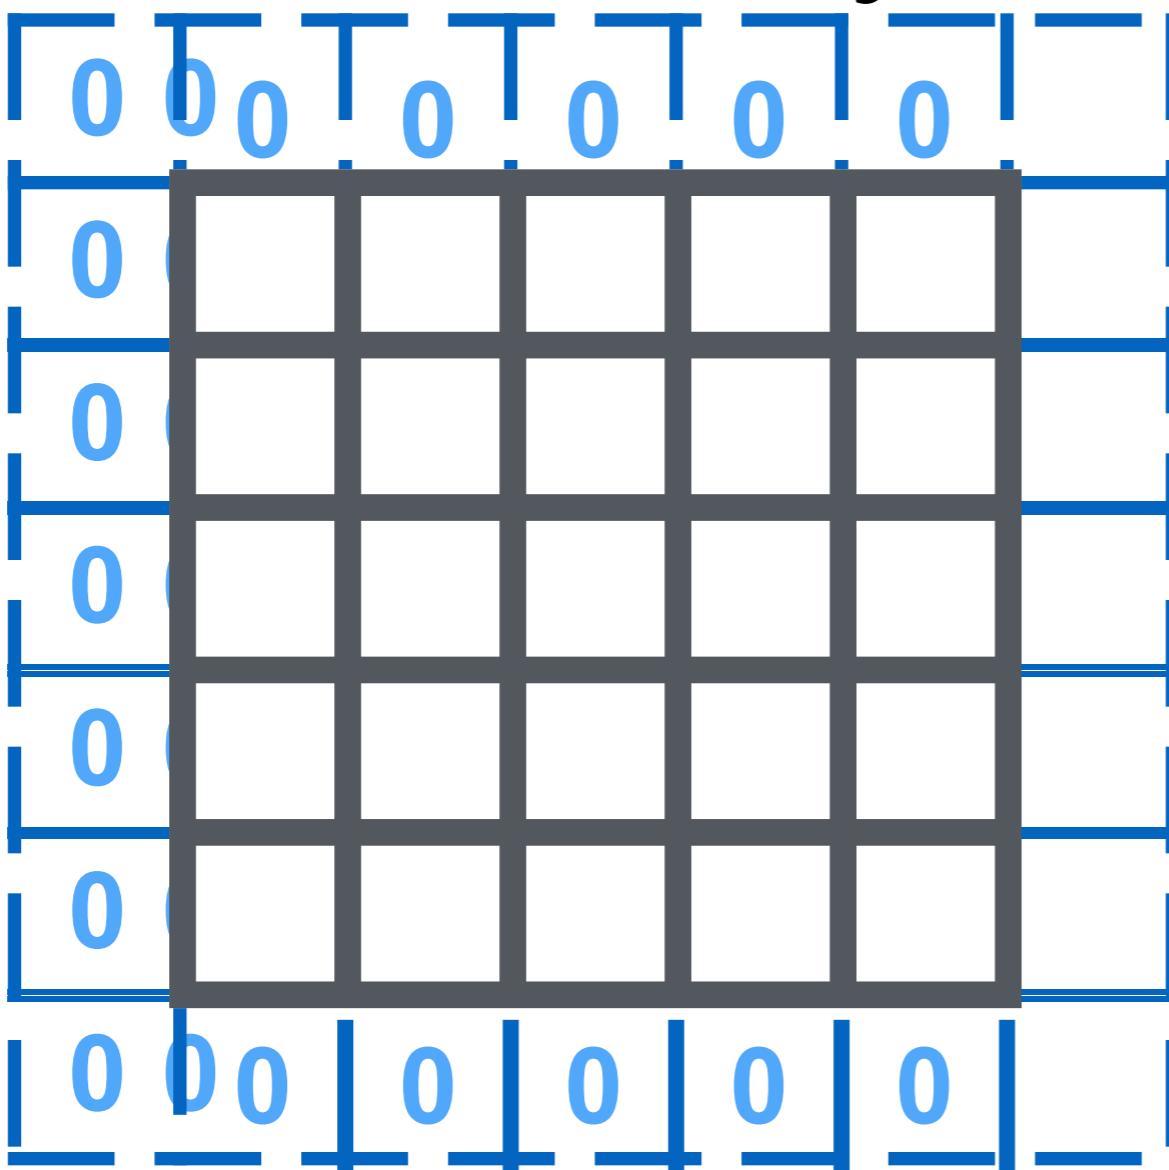
- ## ■ Corner points and ■ isolated ones

After
convolution:

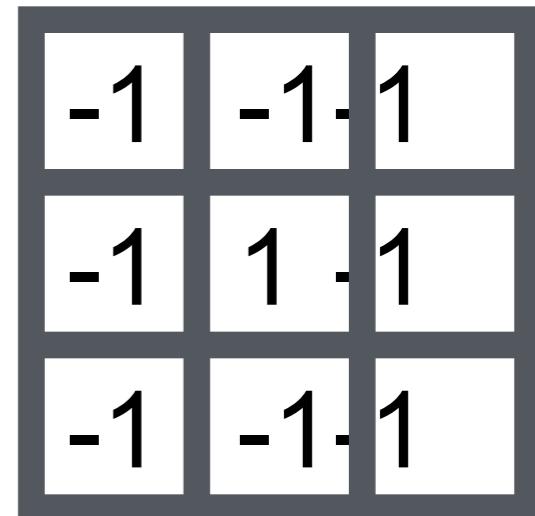


Convolutional Layer: 2D example

A 2D
image:

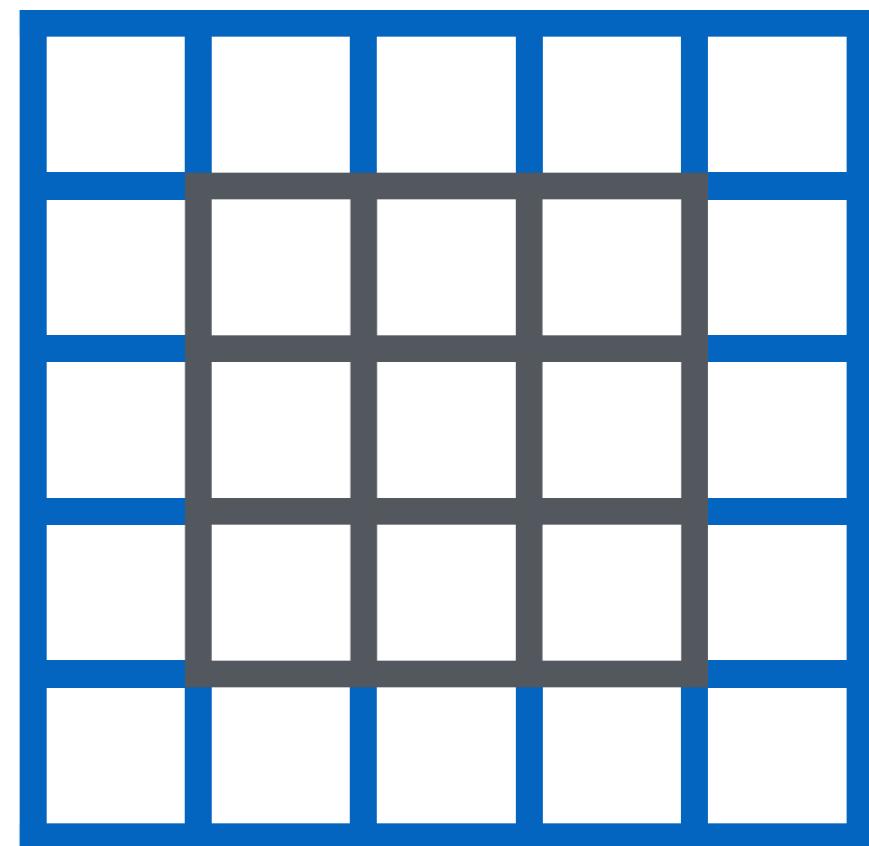


A filter:



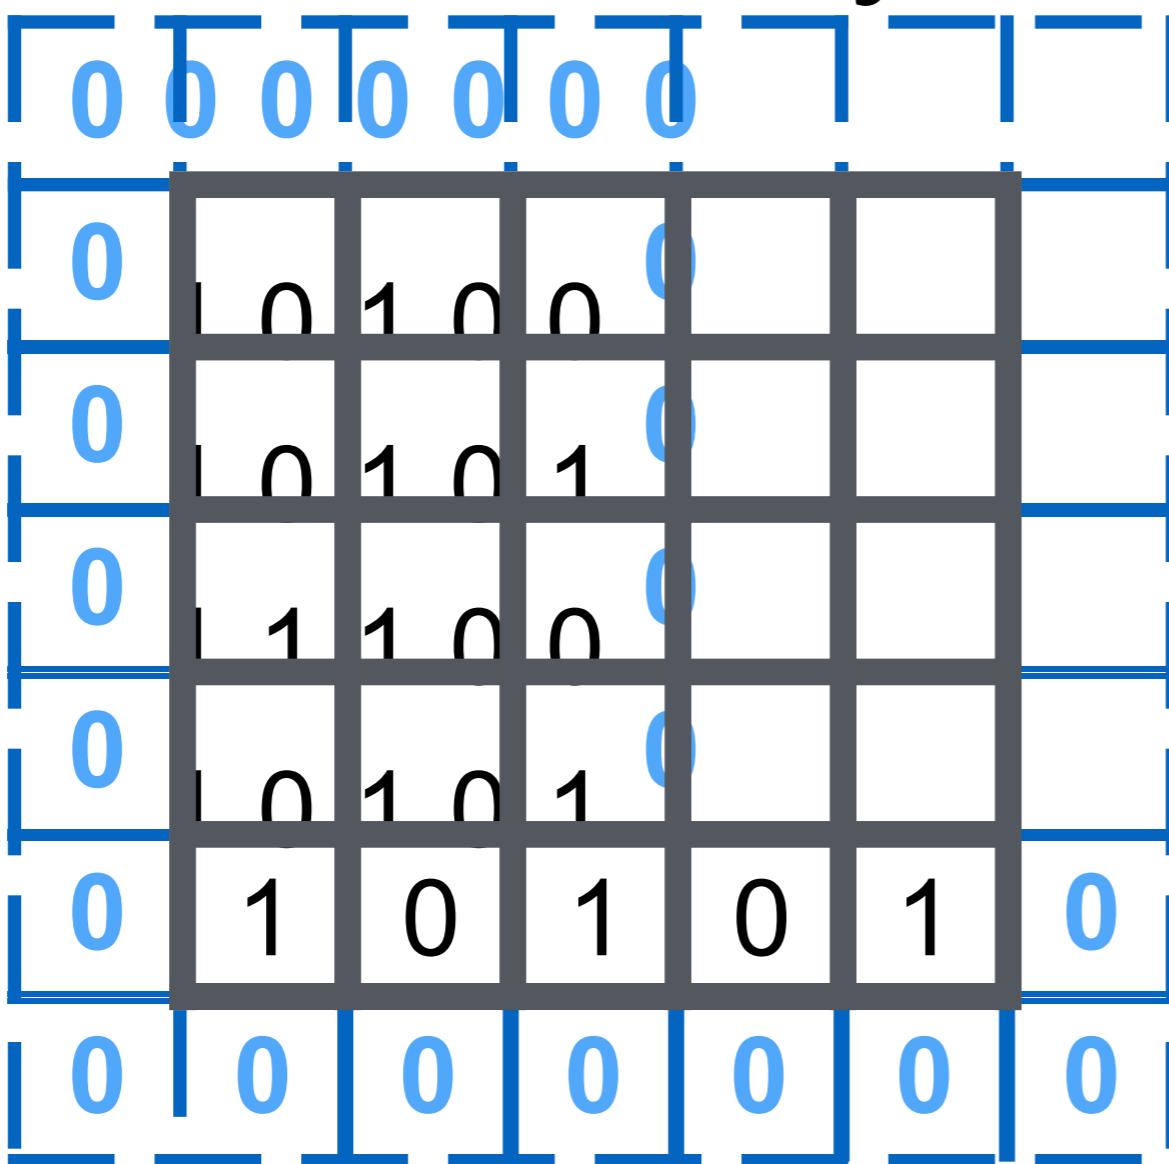
with bias 2

After
convolution:

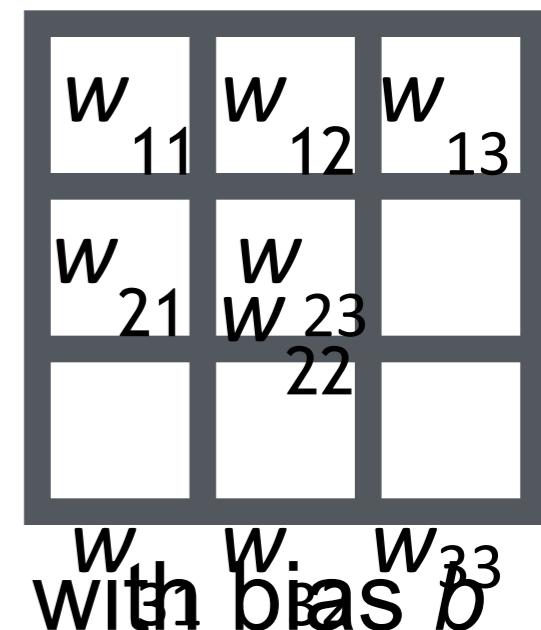


Convolutional Layer: 2D example

A 2D
image:



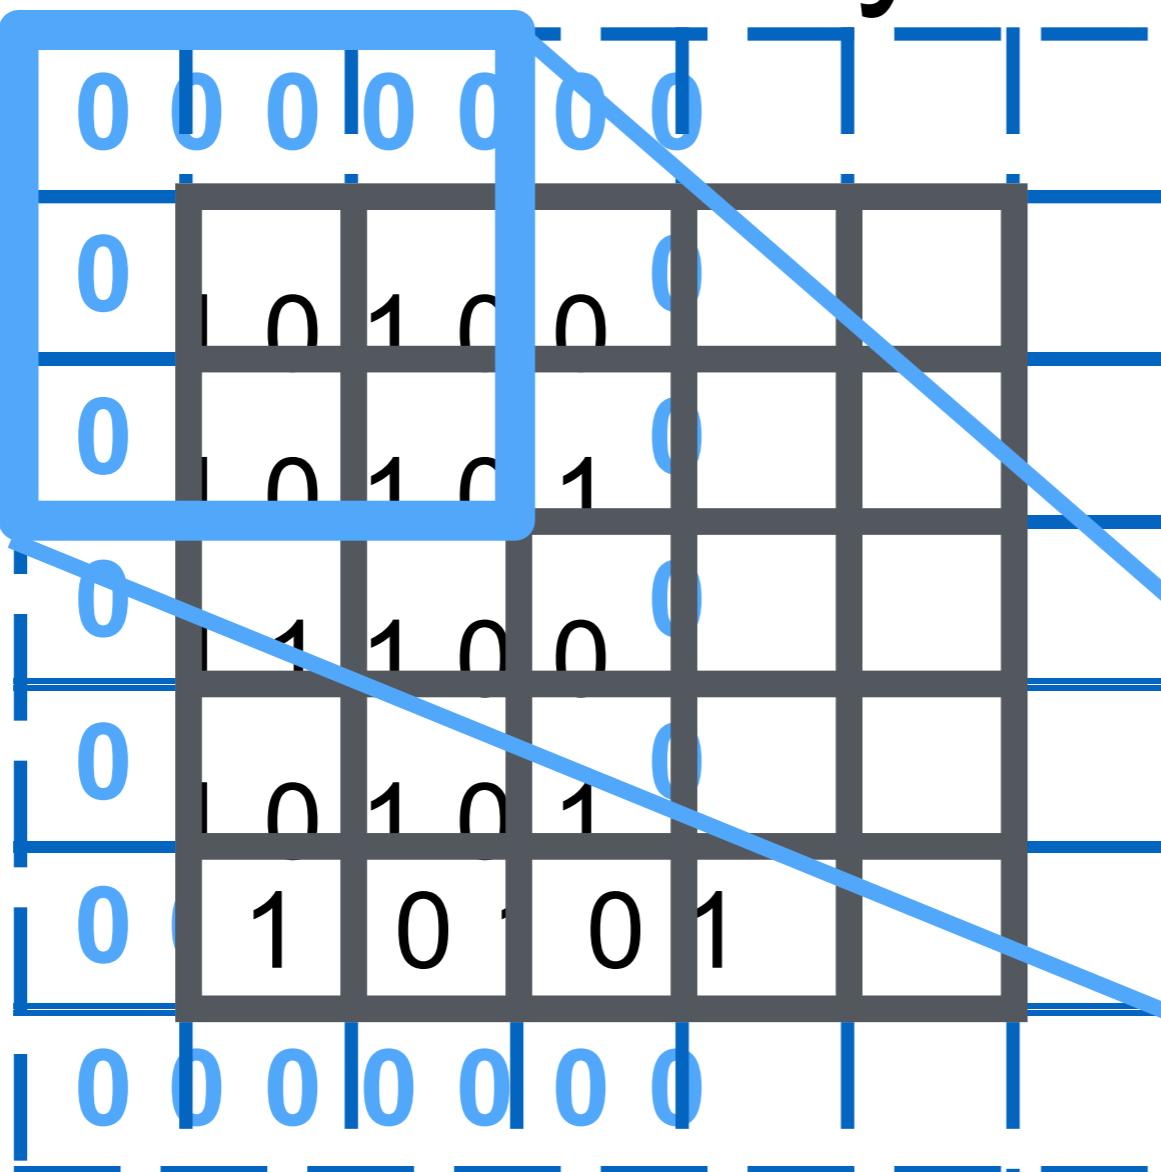
A filter:



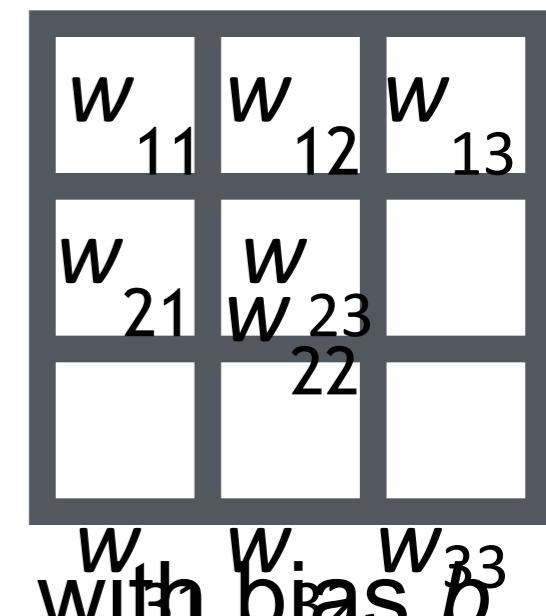
After
convolution:

Convolutional Layer: 2D example

A 2D
image:



A filter:



After
convolution:

Convolutional Layer: 2D example

A 2D
image:

1	0	1	0	0
1	0	1	0	1
1	1	1	0	0
1	0	1	0	1
1	0	1	0	1

A filter:

w_{11}	w_{12}	w_{13}
w_{21}	w_{23}	w_{22}

w_1 w_2 w_3
with bias b

What is the output size?

Convolutional Layer: 3D example

A 3D
image:



[<https://helpx.adobe.com/photoshop/key-concepts/skew.html>]

Convolutional Layer: 3D example

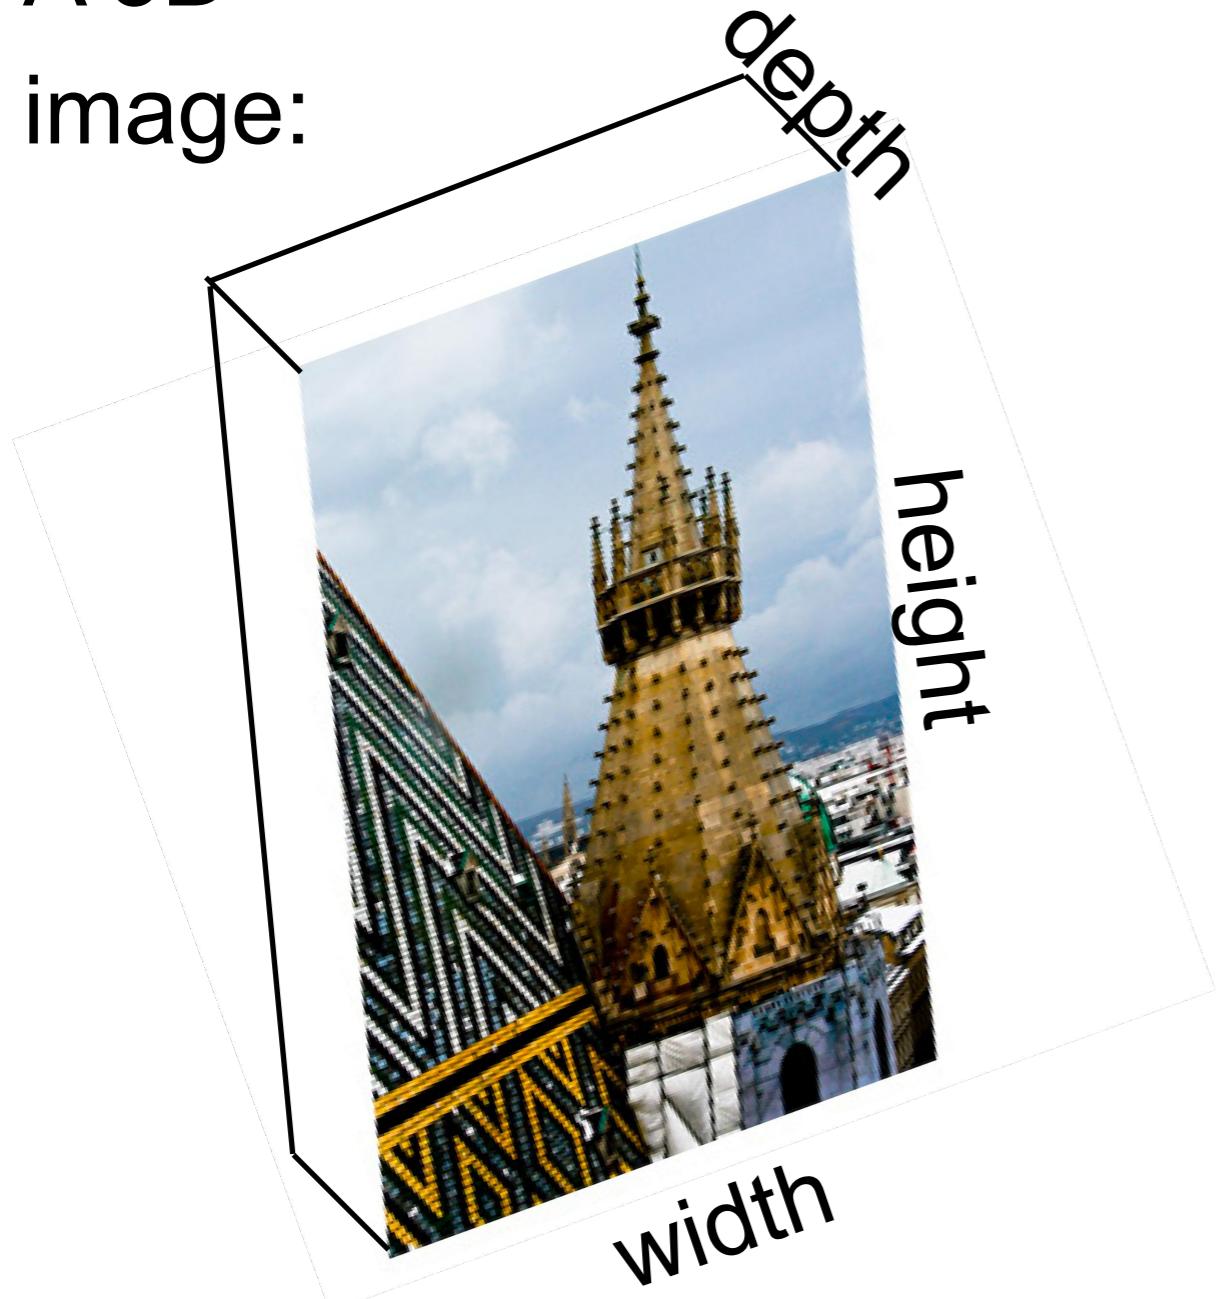
A 3D
image:



[<https://helpx.adobe.com/photoshop/key-concepts/skew.html>]

Convolutional Layer: 3D example

A 3D
image:

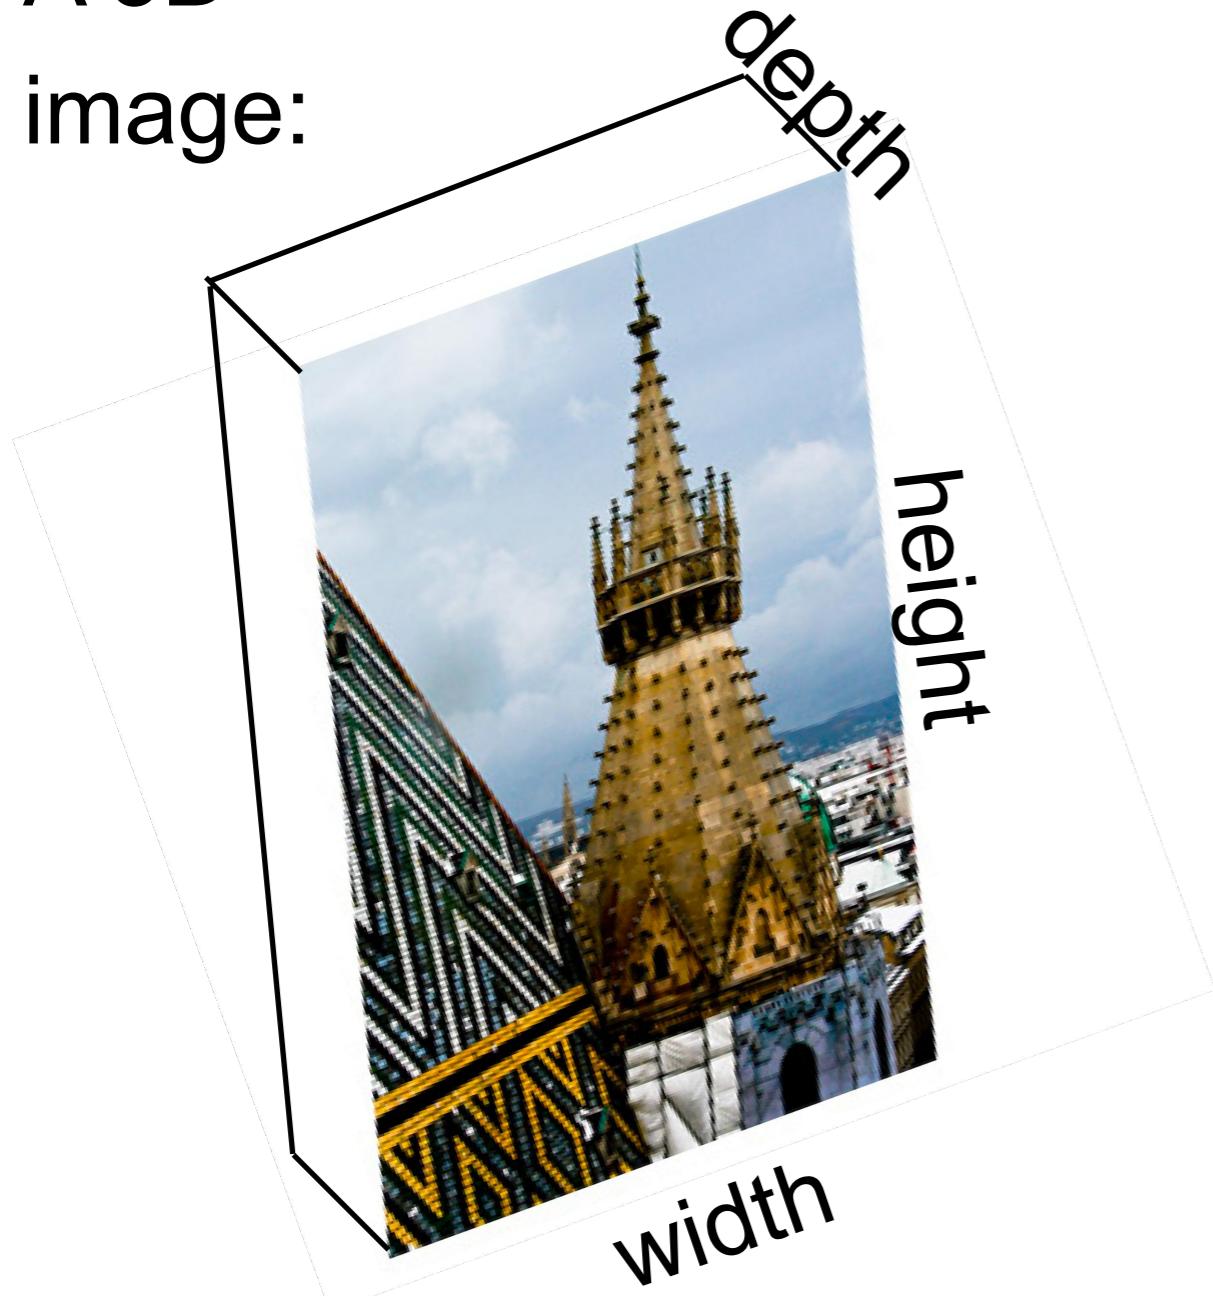


[<https://helpx.adobe.com/photoshop/key-concepts/skew.html>]

Convolutional Layer: 3D example

A 3D

image:



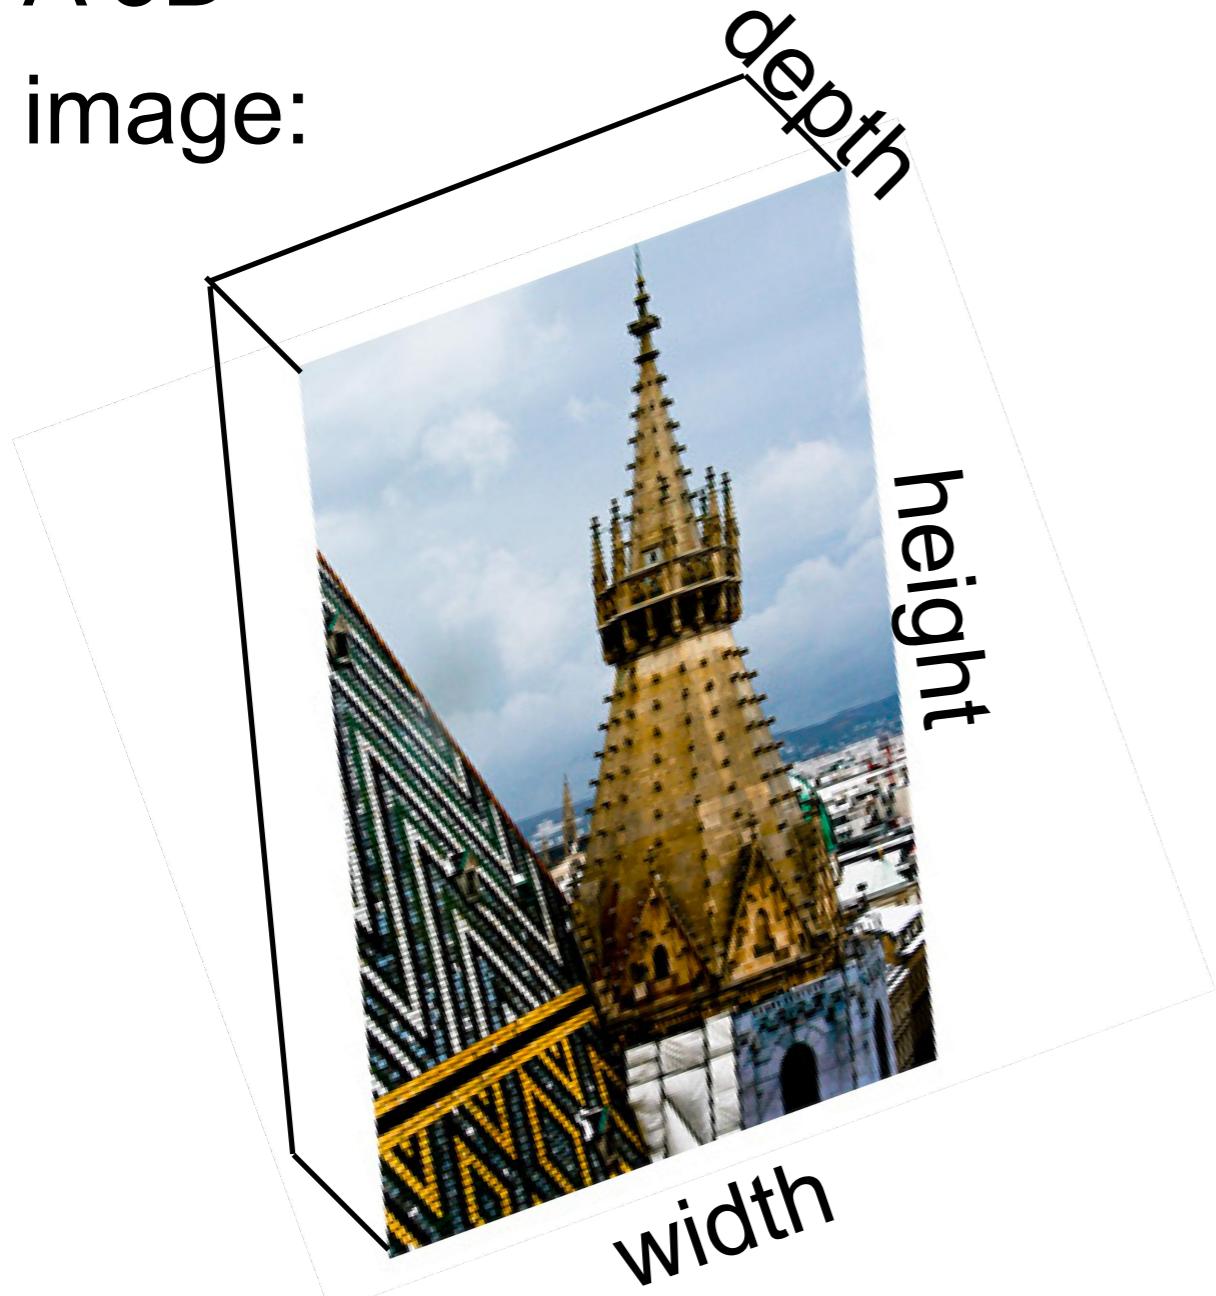
- Tensor: generalization of a matrix
- E.g. 1D: vector, 2D: matrix

[<https://helpx.adobe.com/photoshop/key-concepts/skew.html>]

Convolutional Layer: 3D example

A 3D

image:

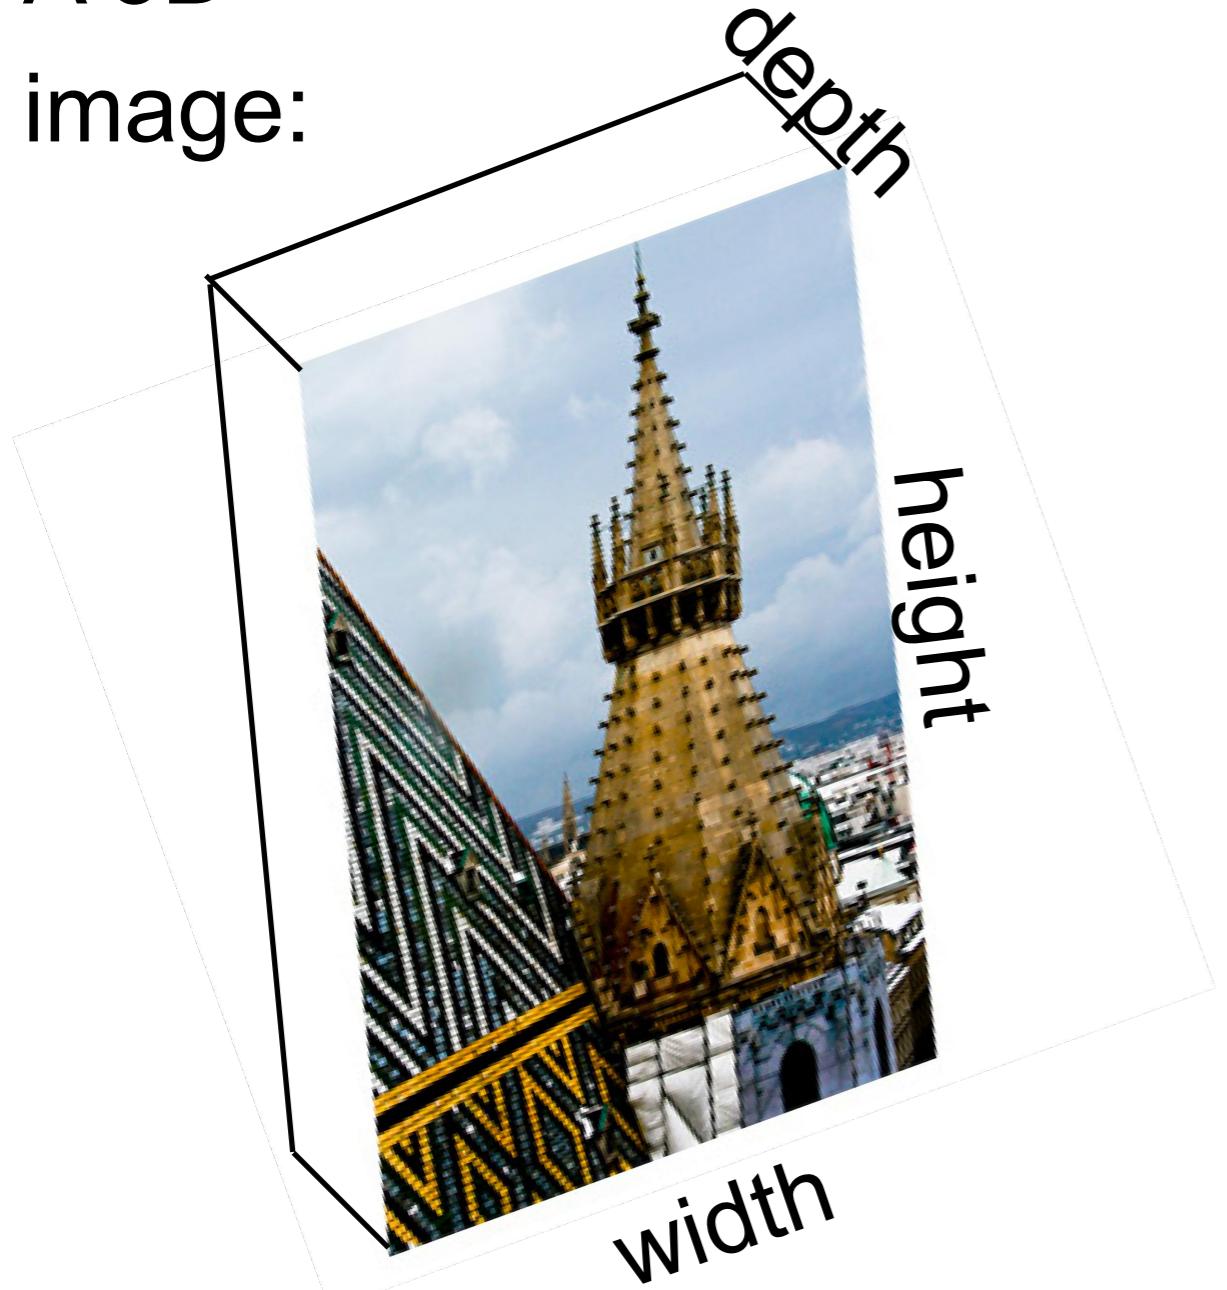


- Tensor: generalization of a matrix
- E.g. 1D: vector, 2D: matrix

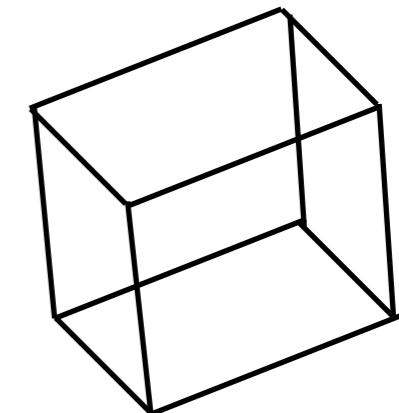
[<https://helpx.adobe.com/photoshop/key-concepts/skew.html>]

Convolutional Layer: 3D example

A 3D
image:



A filter:



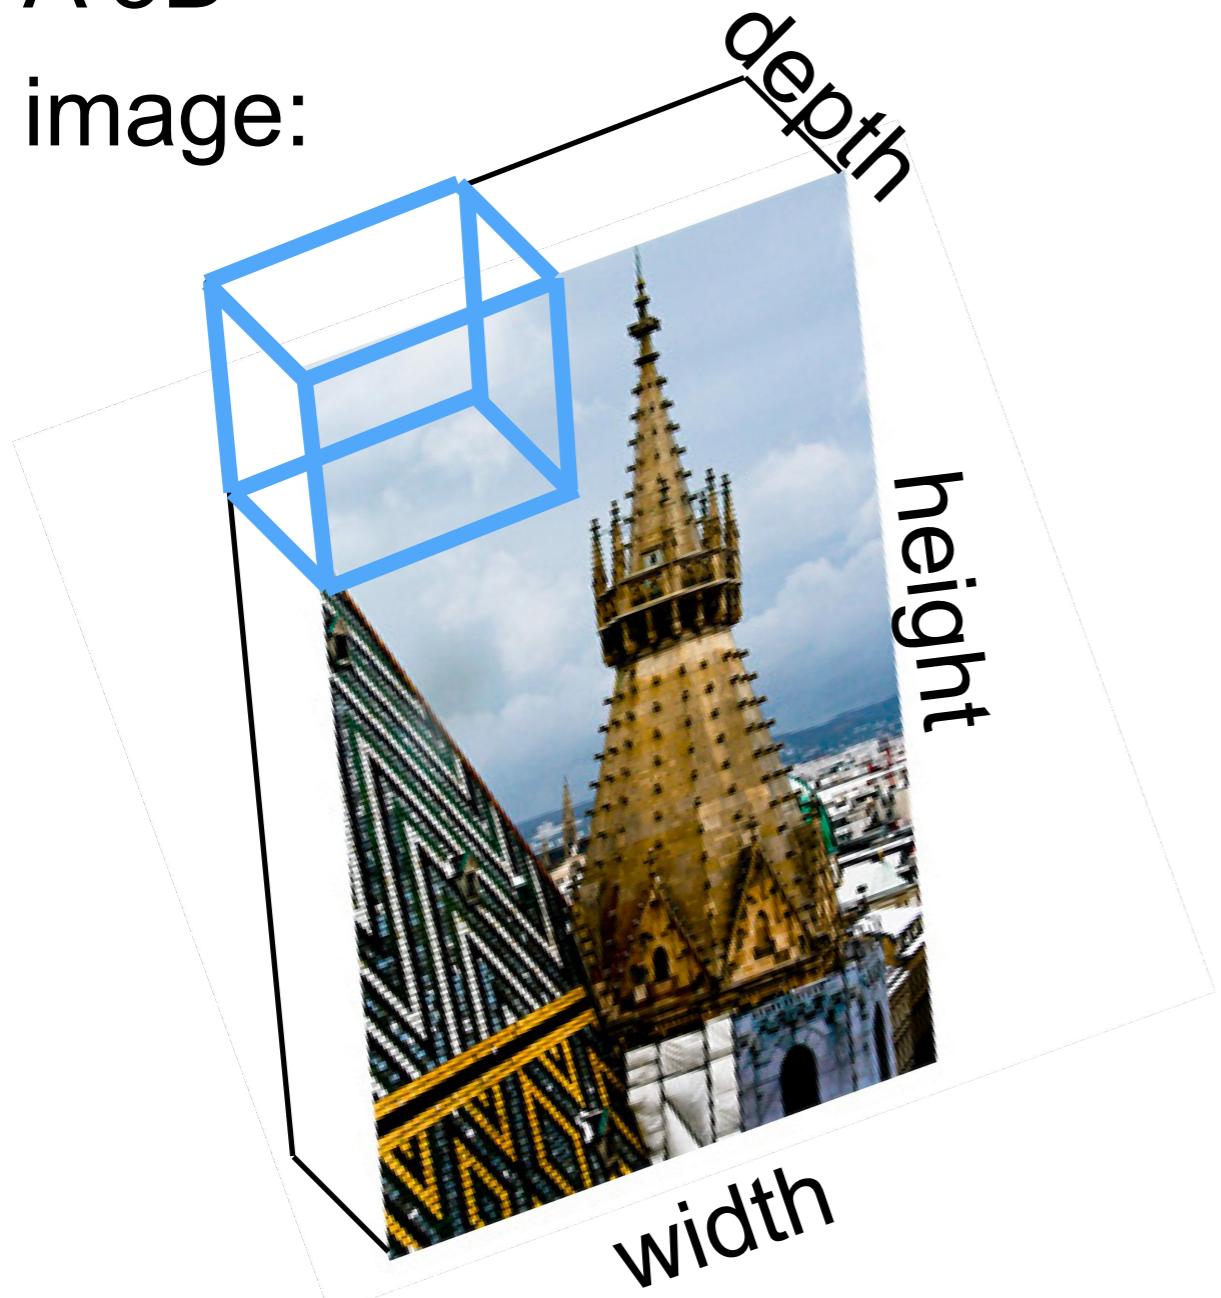
- Tensor: generalization of a matrix
- E.g. 1D: vector, 2D: matrix

[<https://helpx.adobe.com/photoshop/key-concepts/skew.html>]

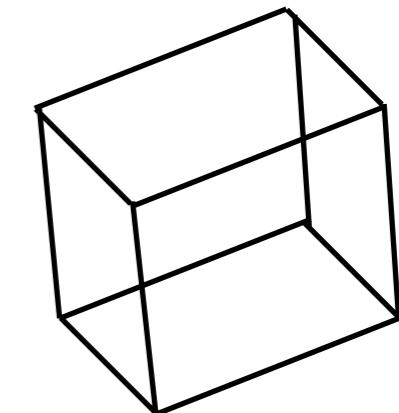
Convolutional Layer: 3D example

A 3D

image:



A filter:



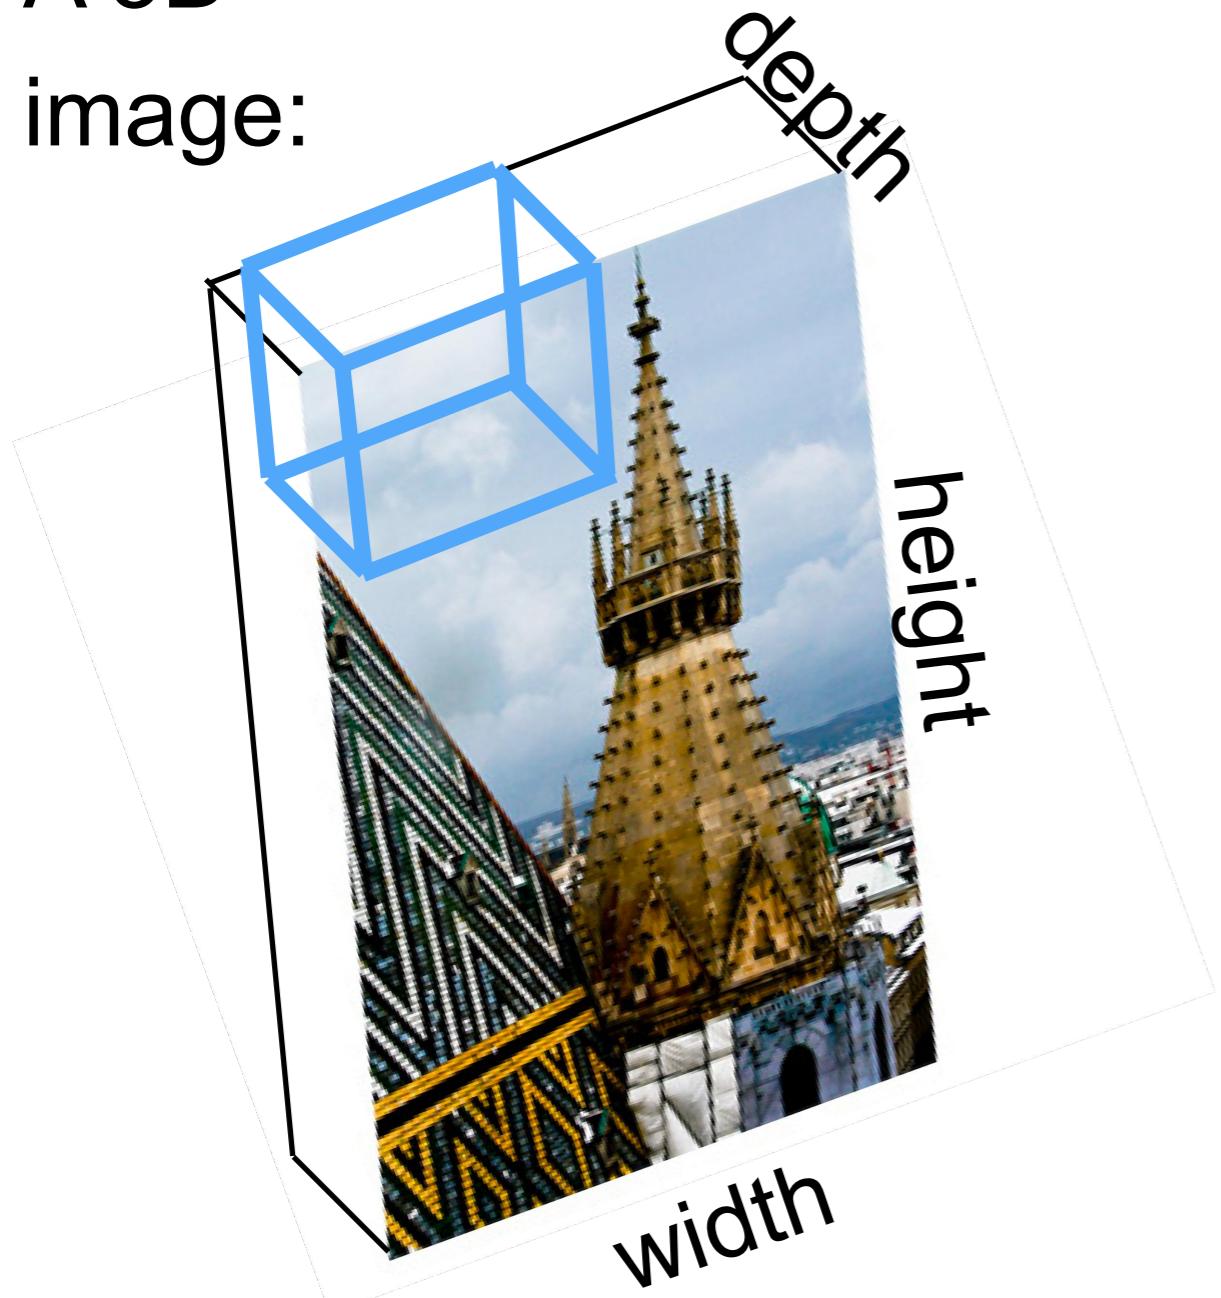
- Tensor: generalization of a matrix
- E.g. 1D: vector, 2D: matrix

[<https://helpx.adobe.com/photoshop/key-concepts/skew.html>]

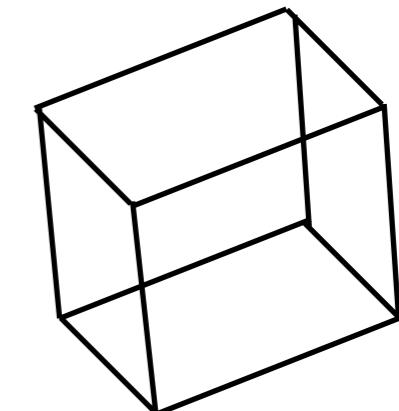
Convolutional Layer: 3D example

A 3D

image:



A filter:



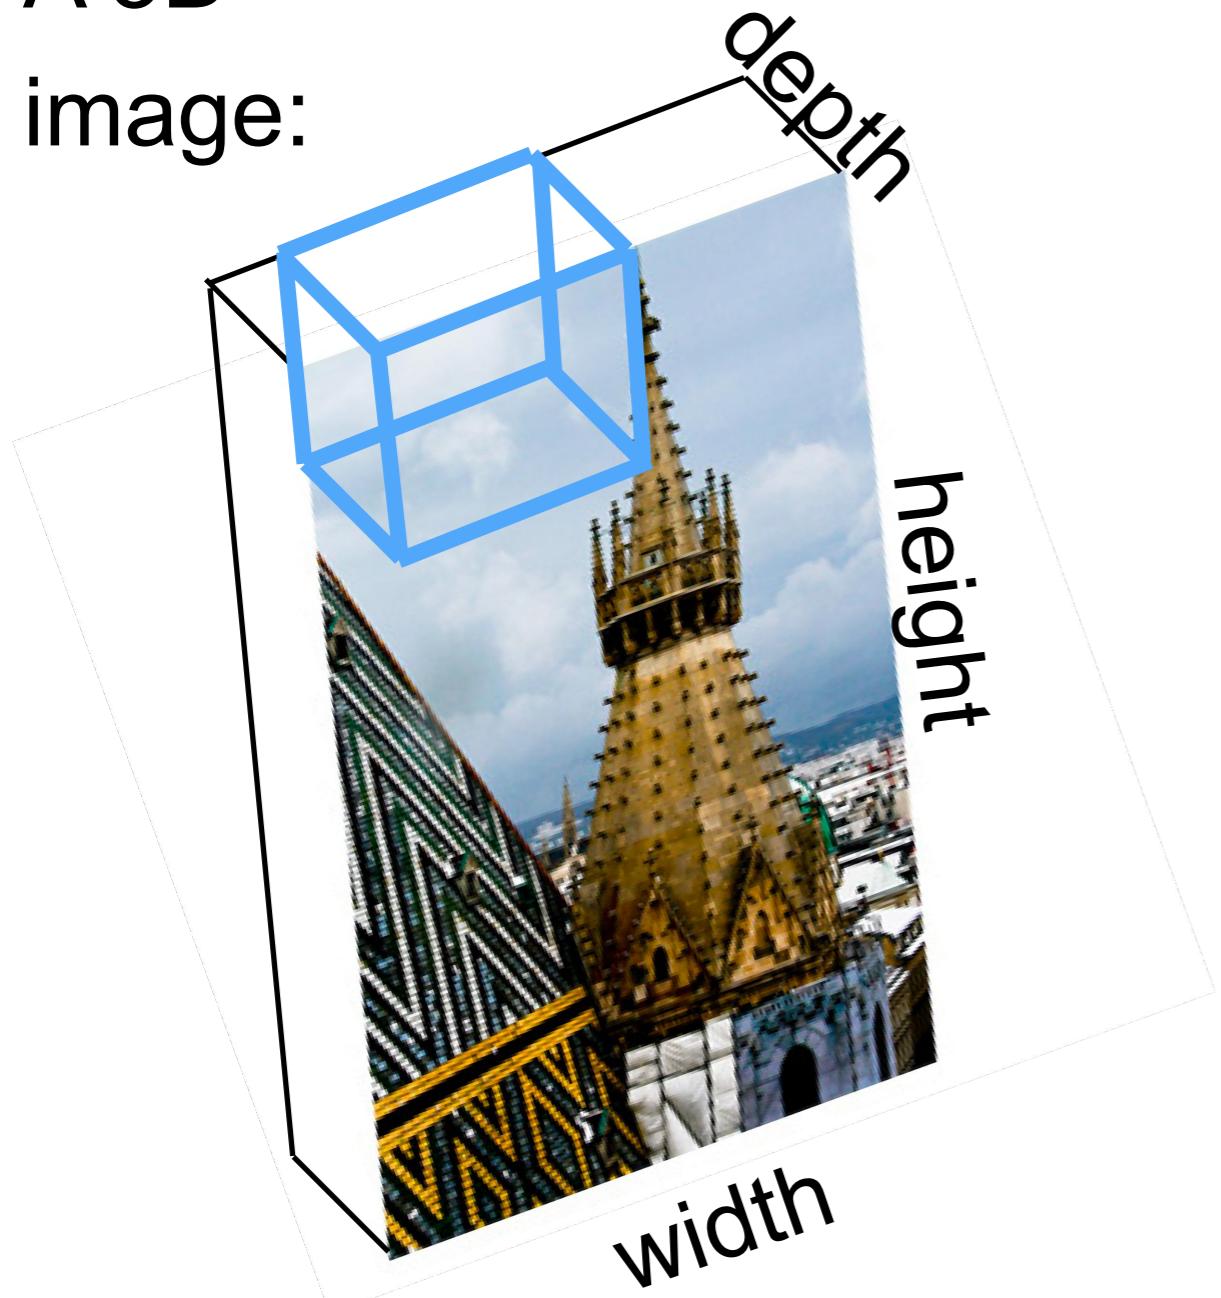
- Tensor: generalization of a matrix
- E.g. 1D: vector, 2D: matrix

[<https://helpx.adobe.com/photoshop/key-concepts/skew.html>]

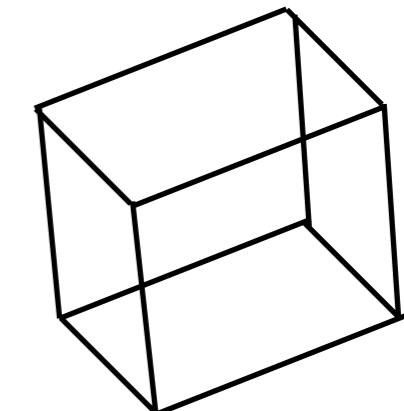
Convolutional Layer: 3D example

A 3D

image:



A filter:



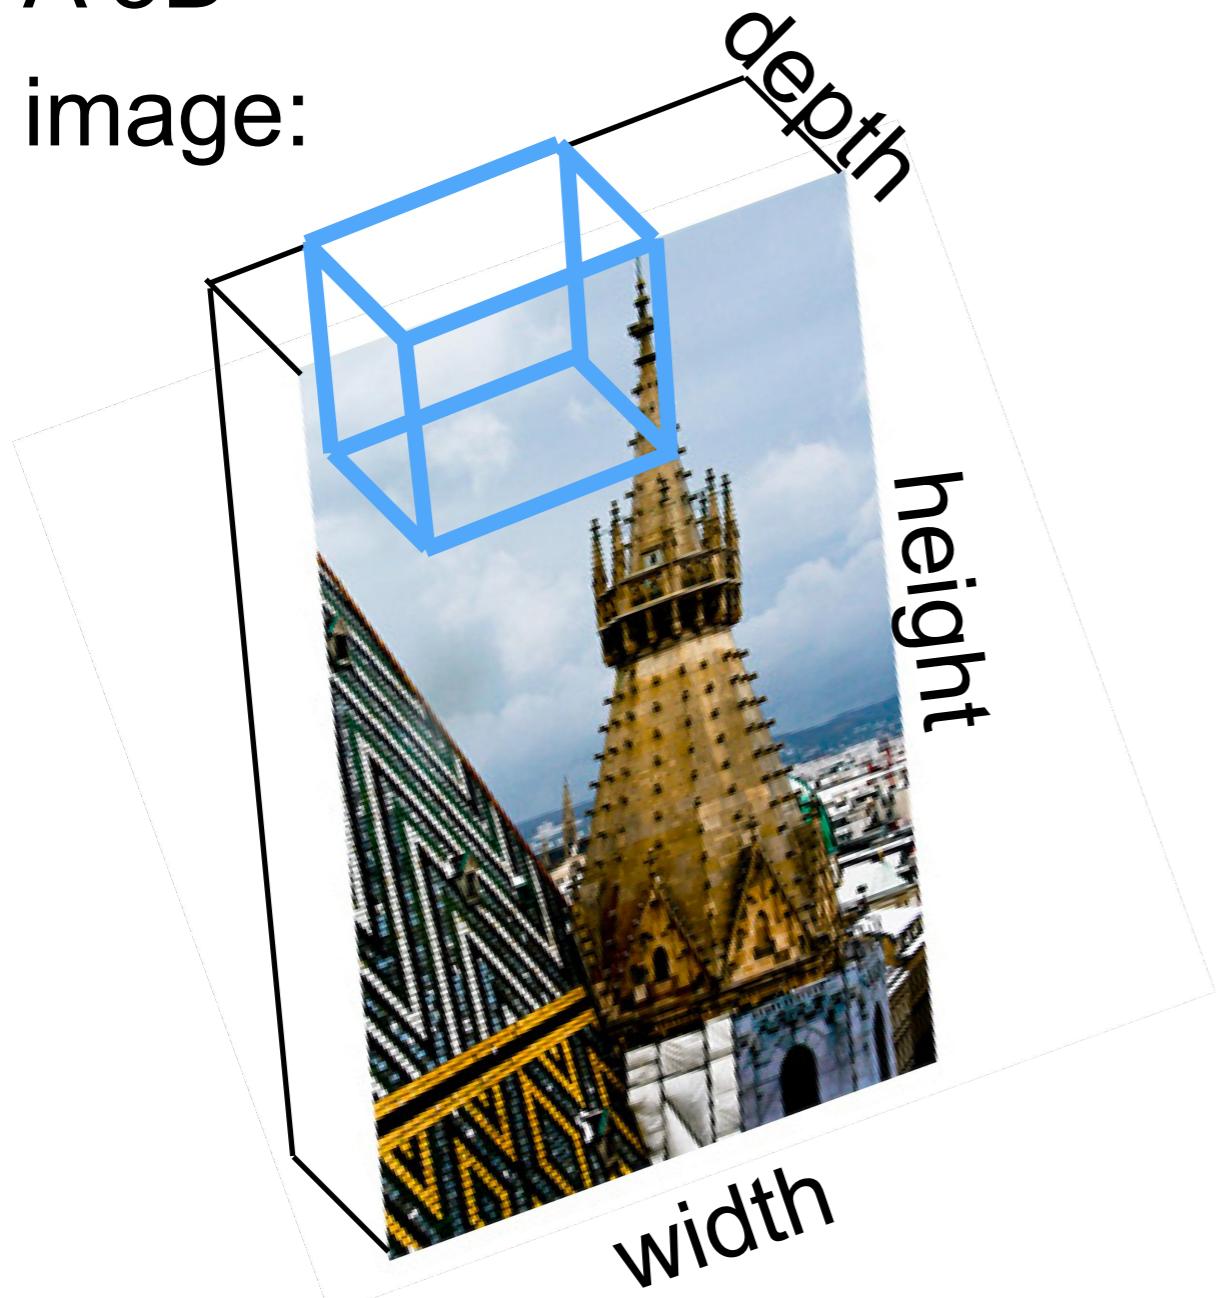
- Tensor: generalization of a matrix
- E.g. 1D: vector, 2D: matrix

[<https://helpx.adobe.com/photoshop/key-concepts/skew.html>]

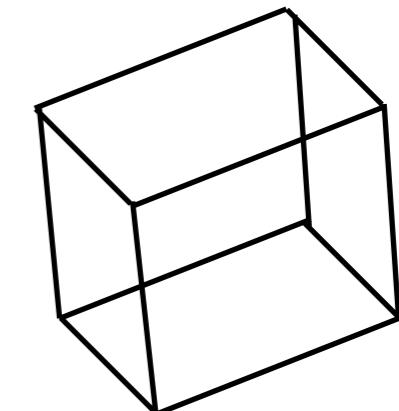
Convolutional Layer: 3D example

A 3D

image:



A filter:

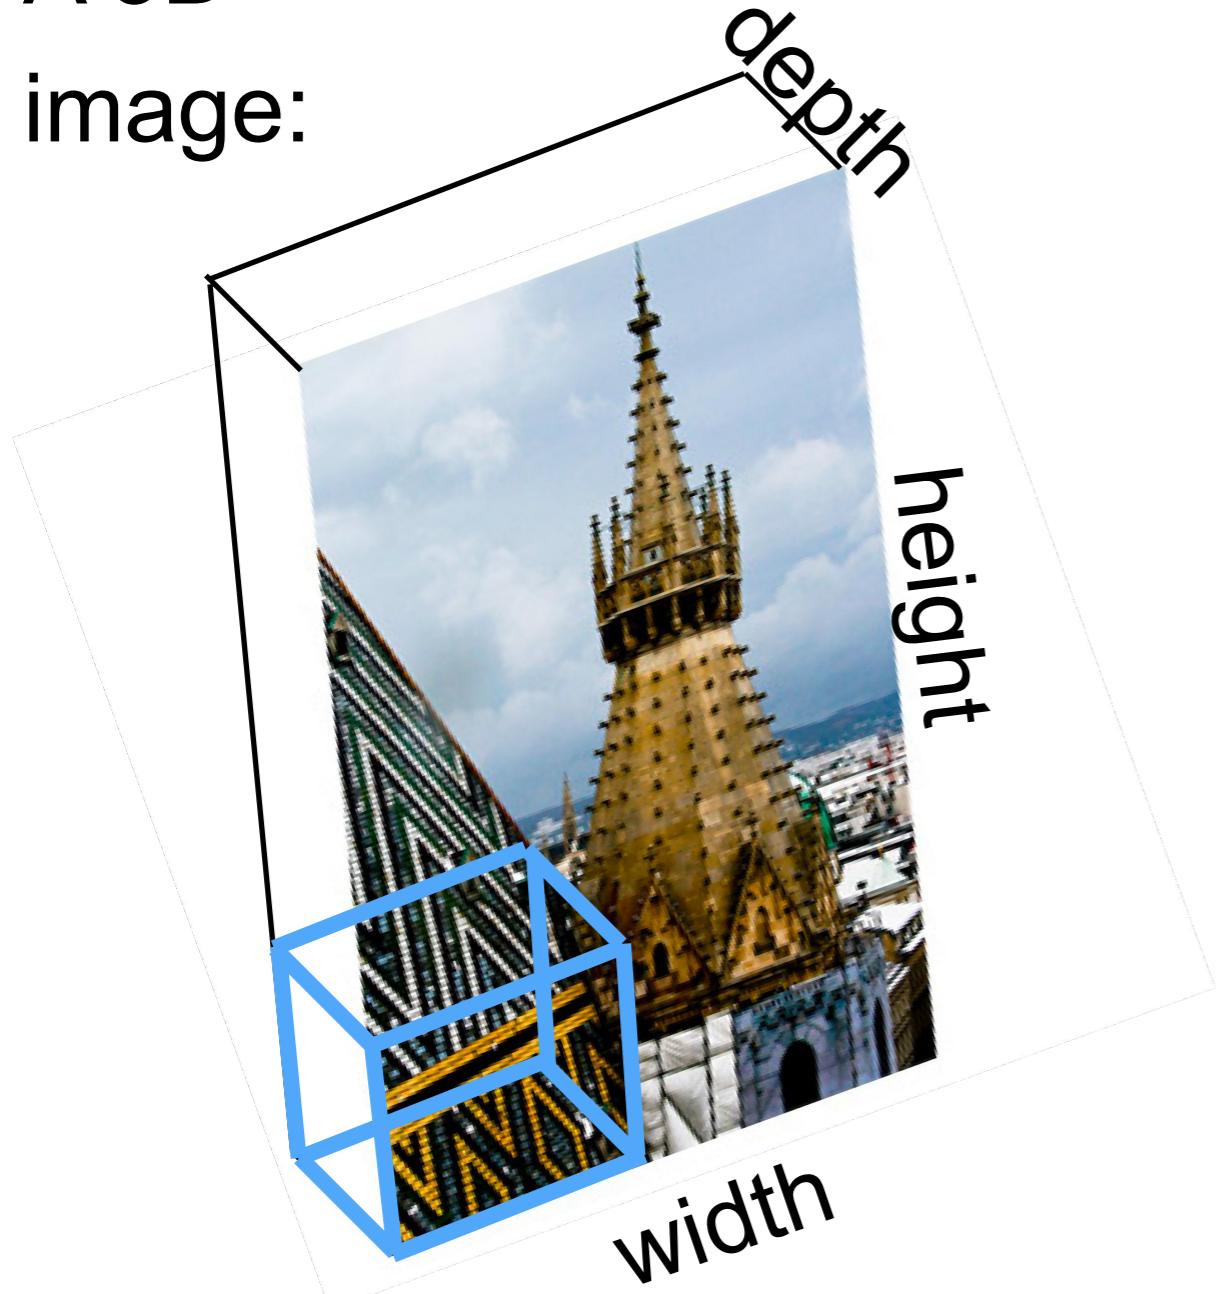


- Tensor: generalization of a matrix
- E.g. 1D: vector, 2D: matrix

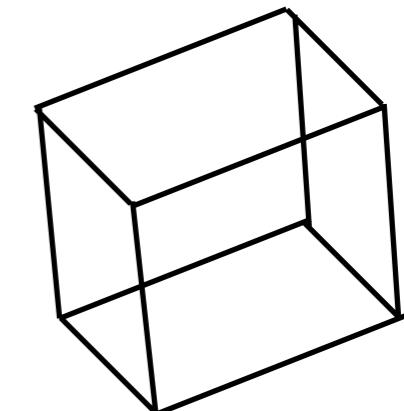
[<https://helpx.adobe.com/photoshop/key-concepts/skew.html>]

Convolutional Layer: 3D example

A 3D
image:



A filter:



- Tensor: generalization of a matrix
- E.g. 1D: vector, 2D: matrix

[<https://helpx.adobe.com/photoshop/key-concepts/skew.html>]

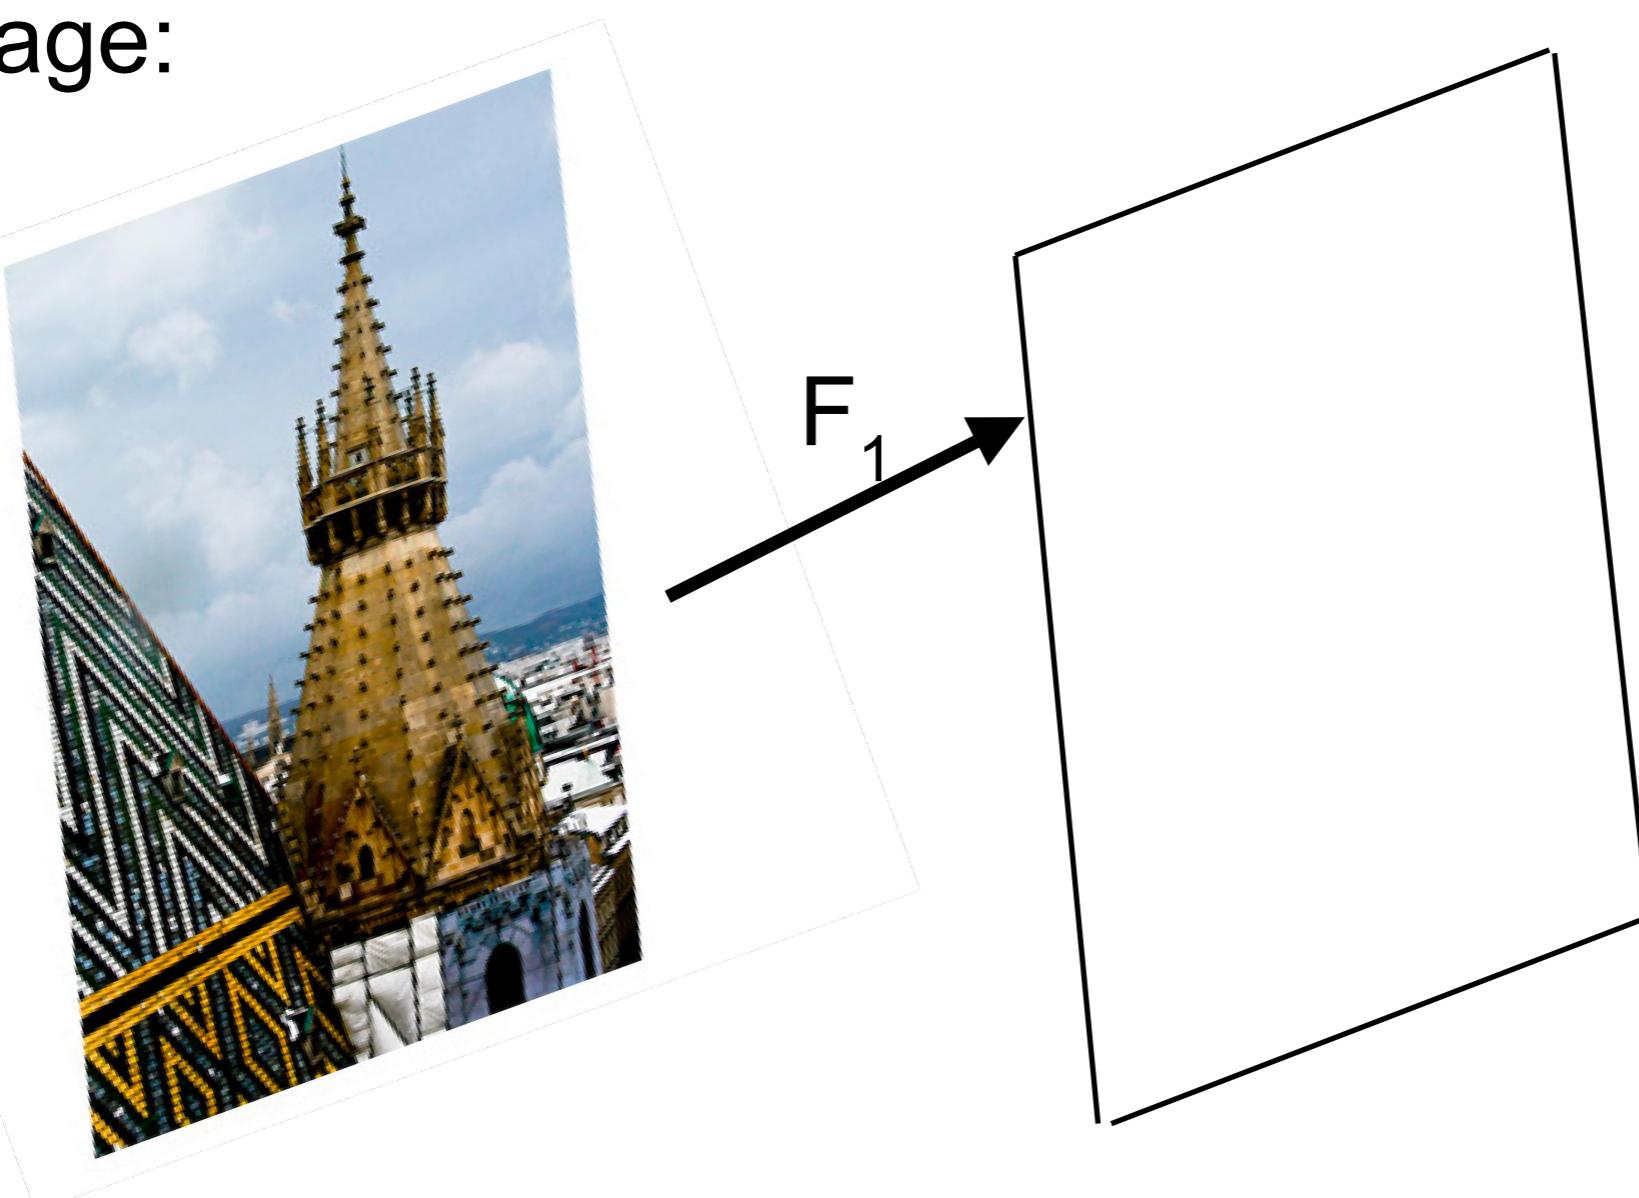
Convolutional Layer: multiple filters

An
image:



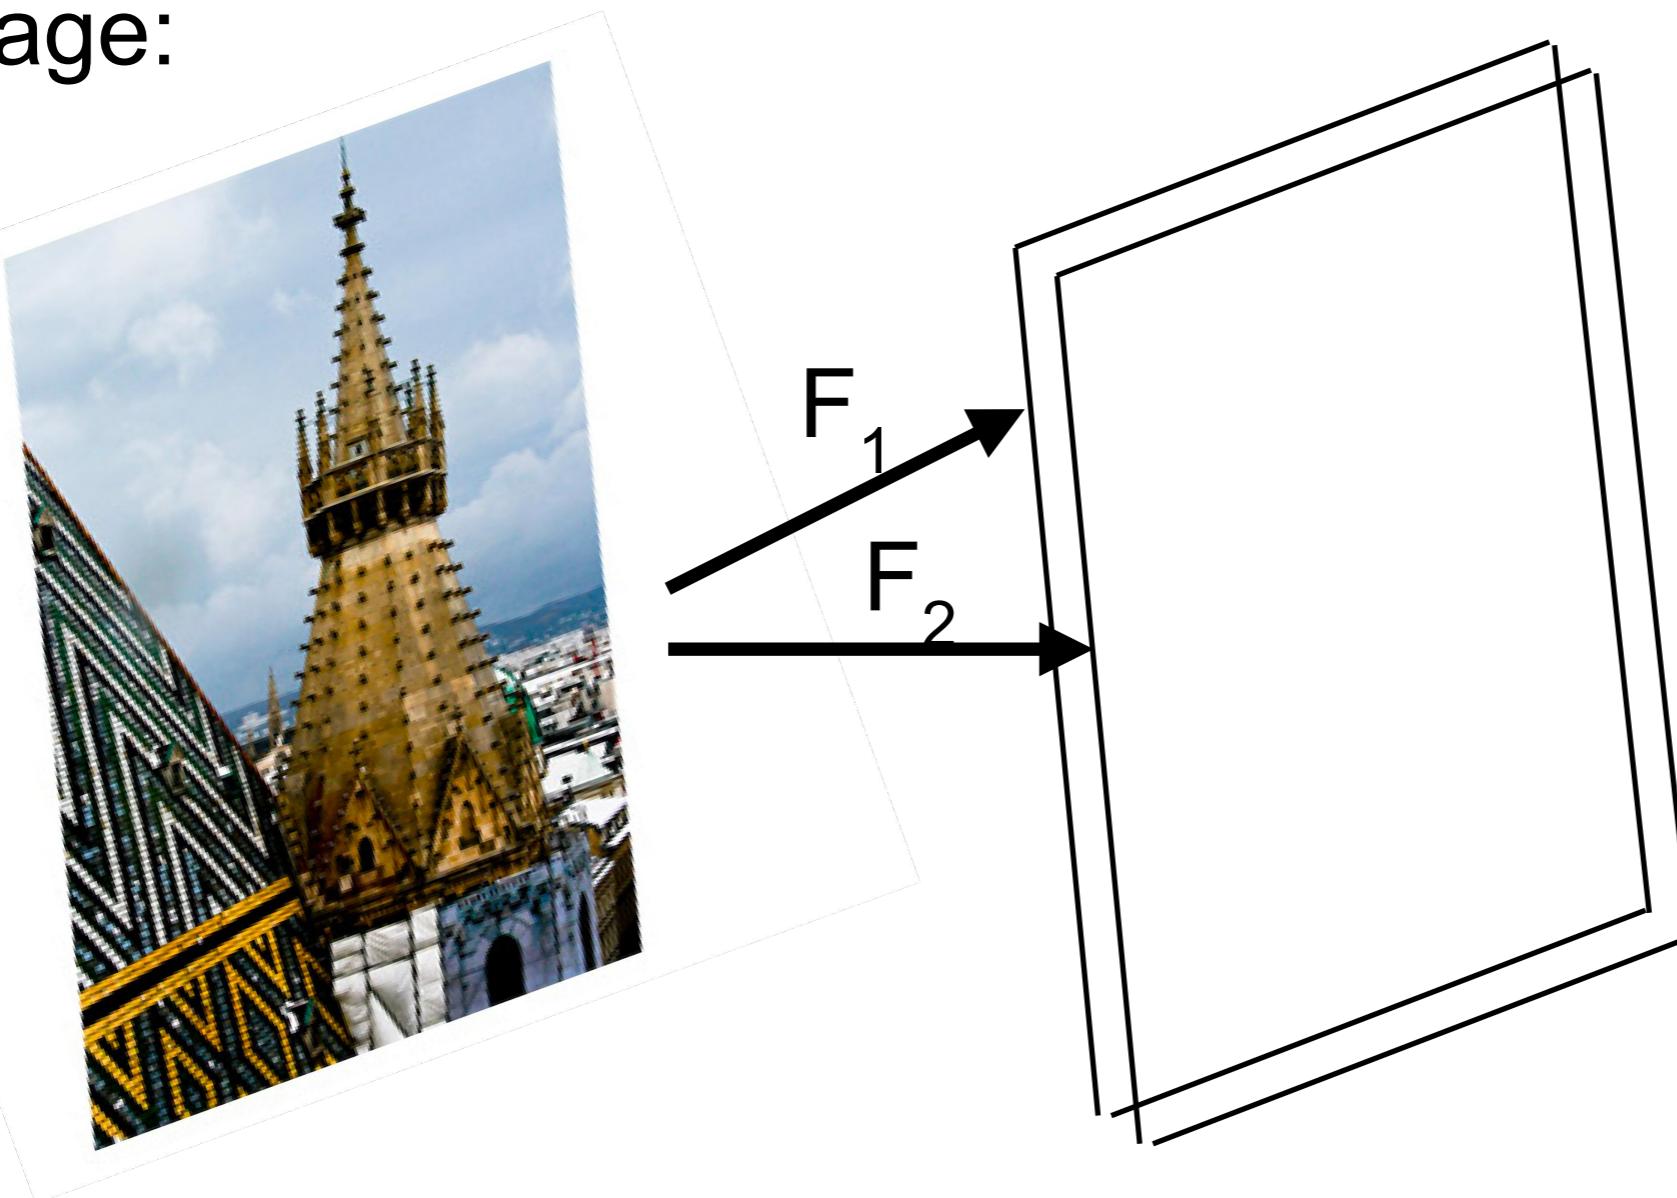
Convolutional Layer: multiple filters

An
image:



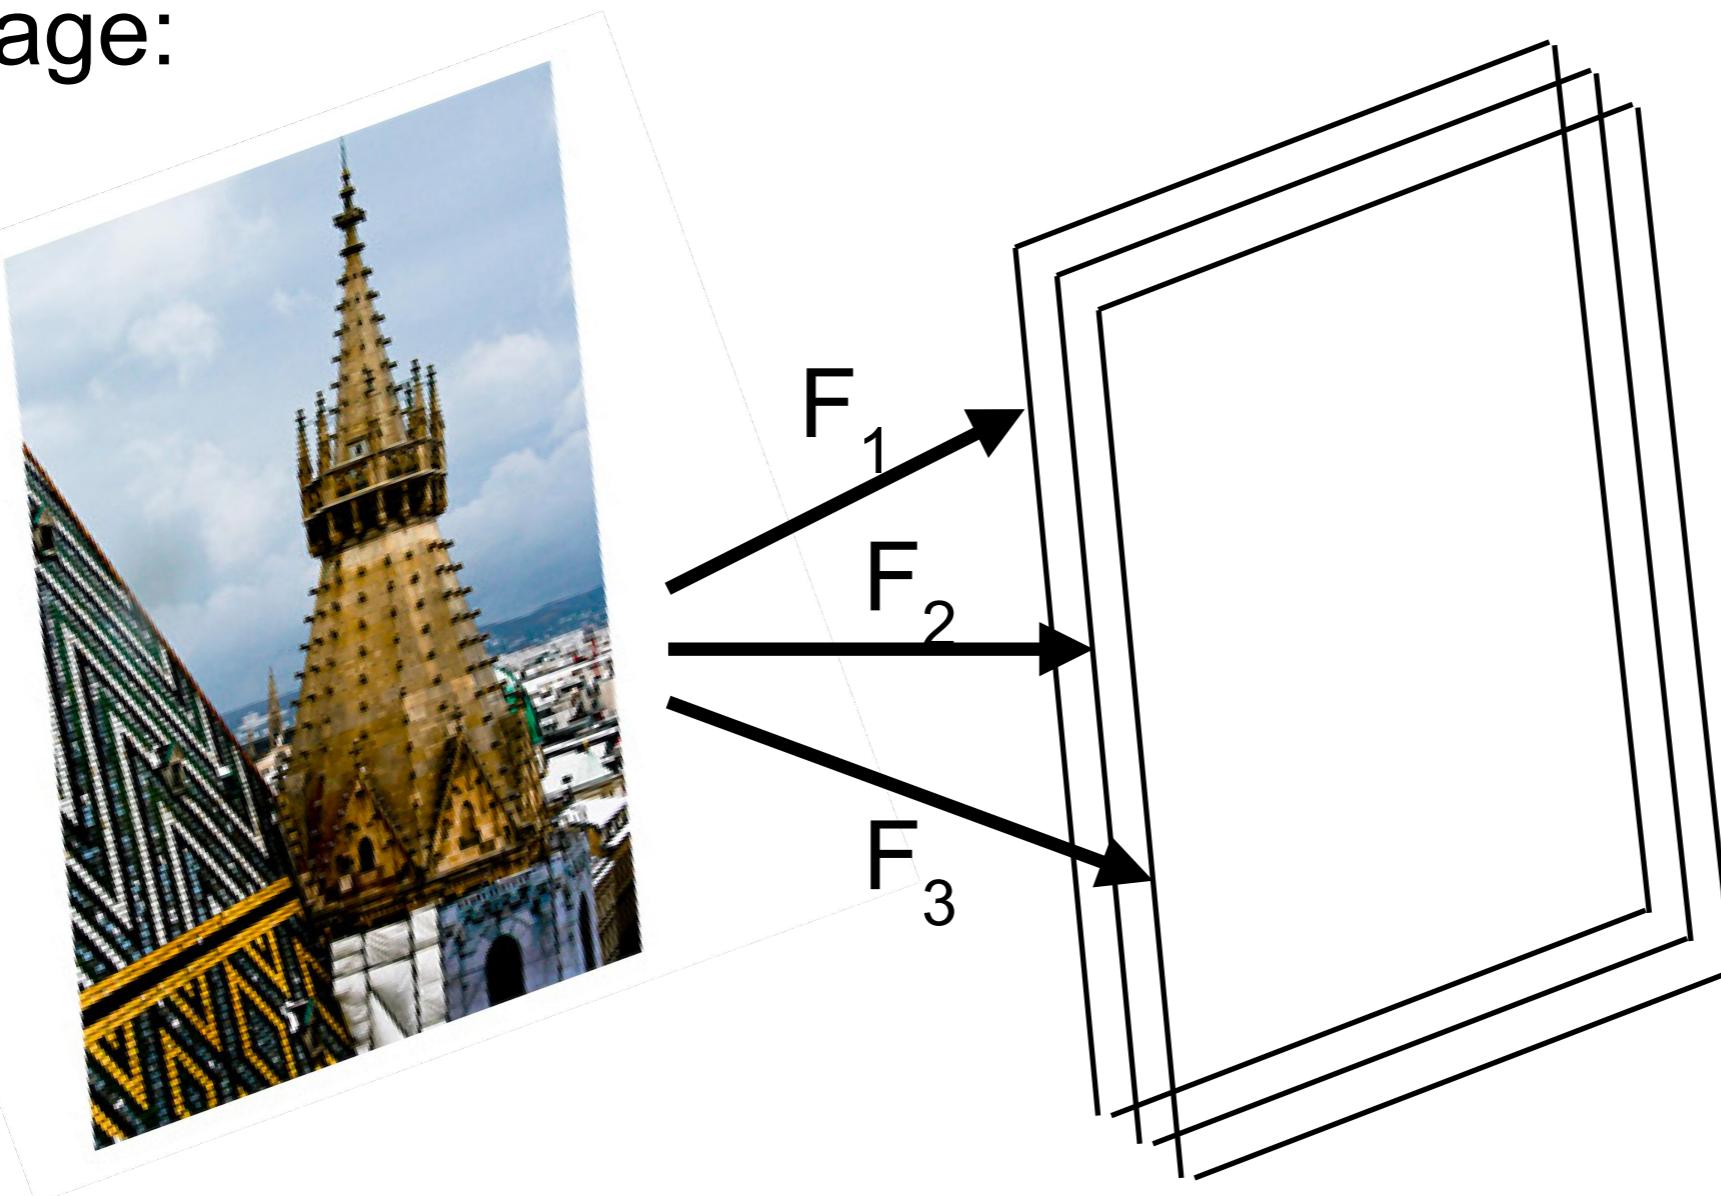
Convolutional Layer: multiple filters

An
image:



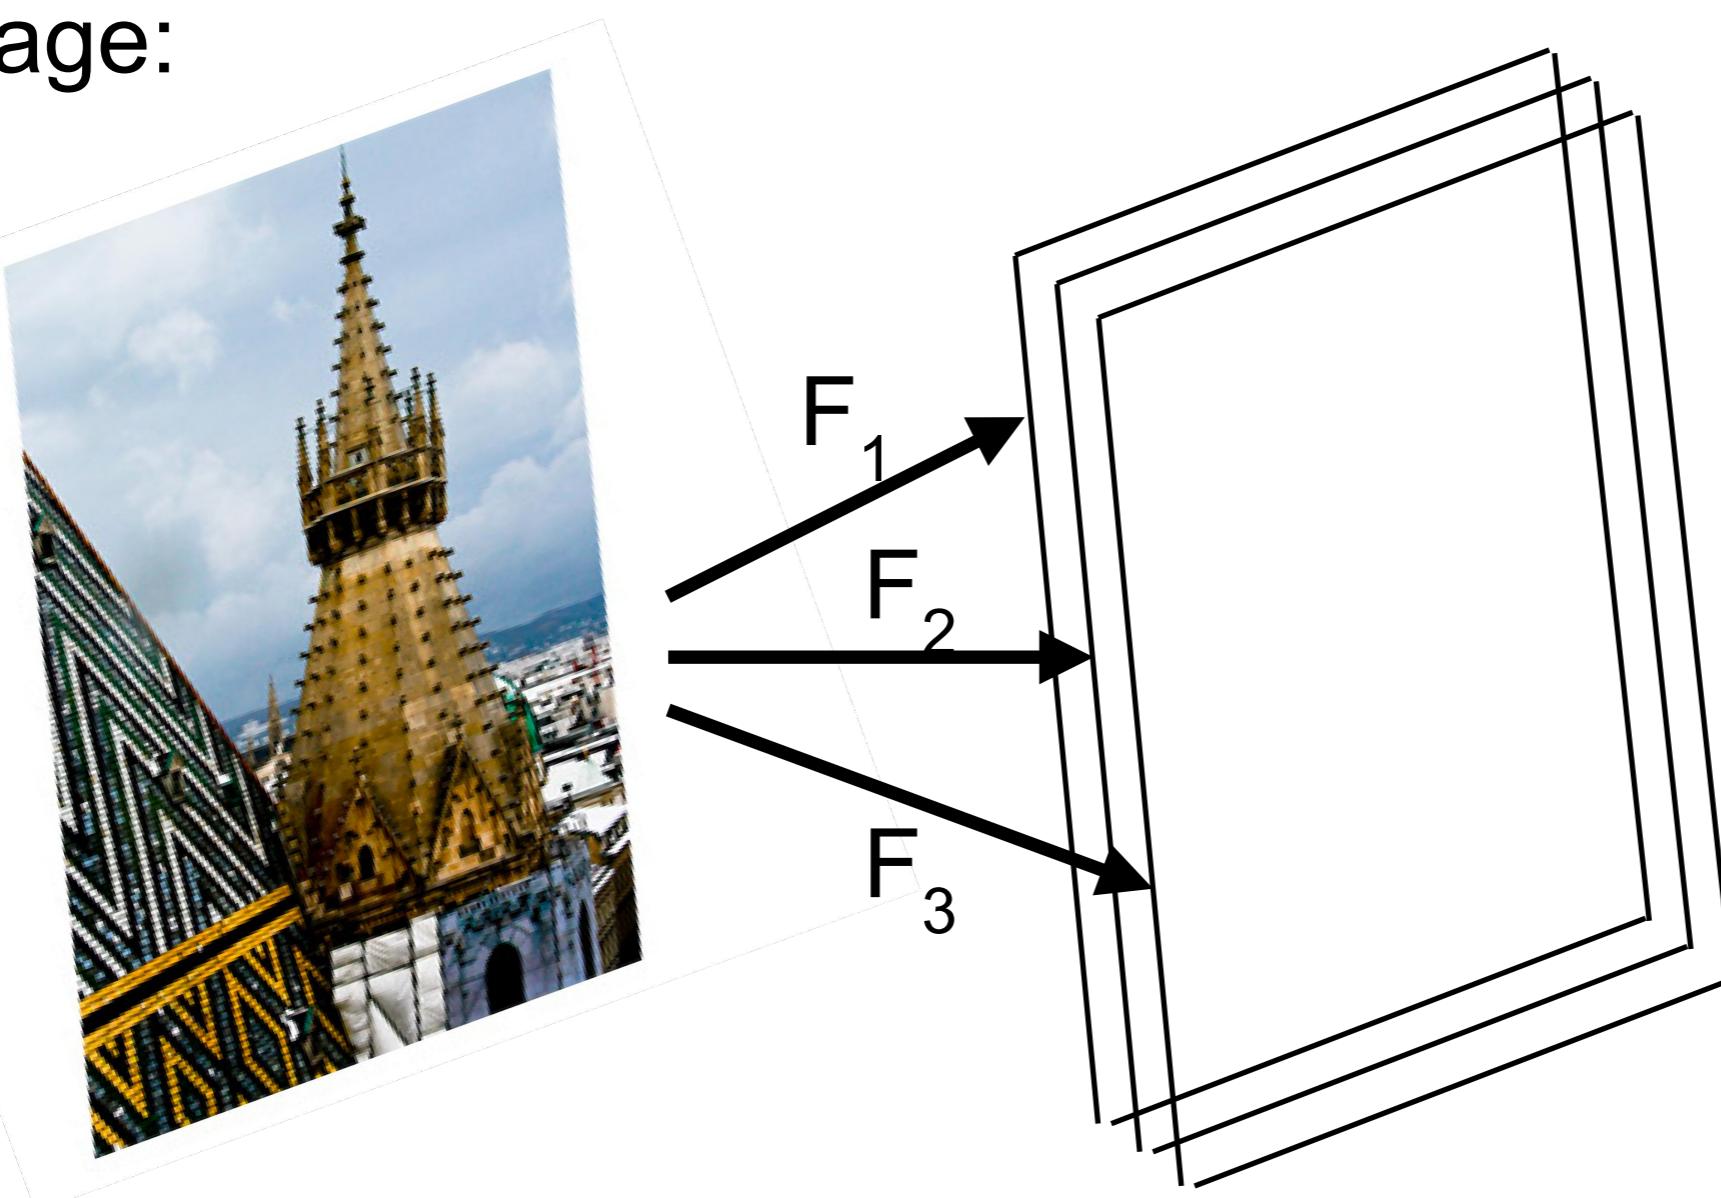
Convolutional Layer: multiple filters

An
image:



Convolutional Layer: multiple filters

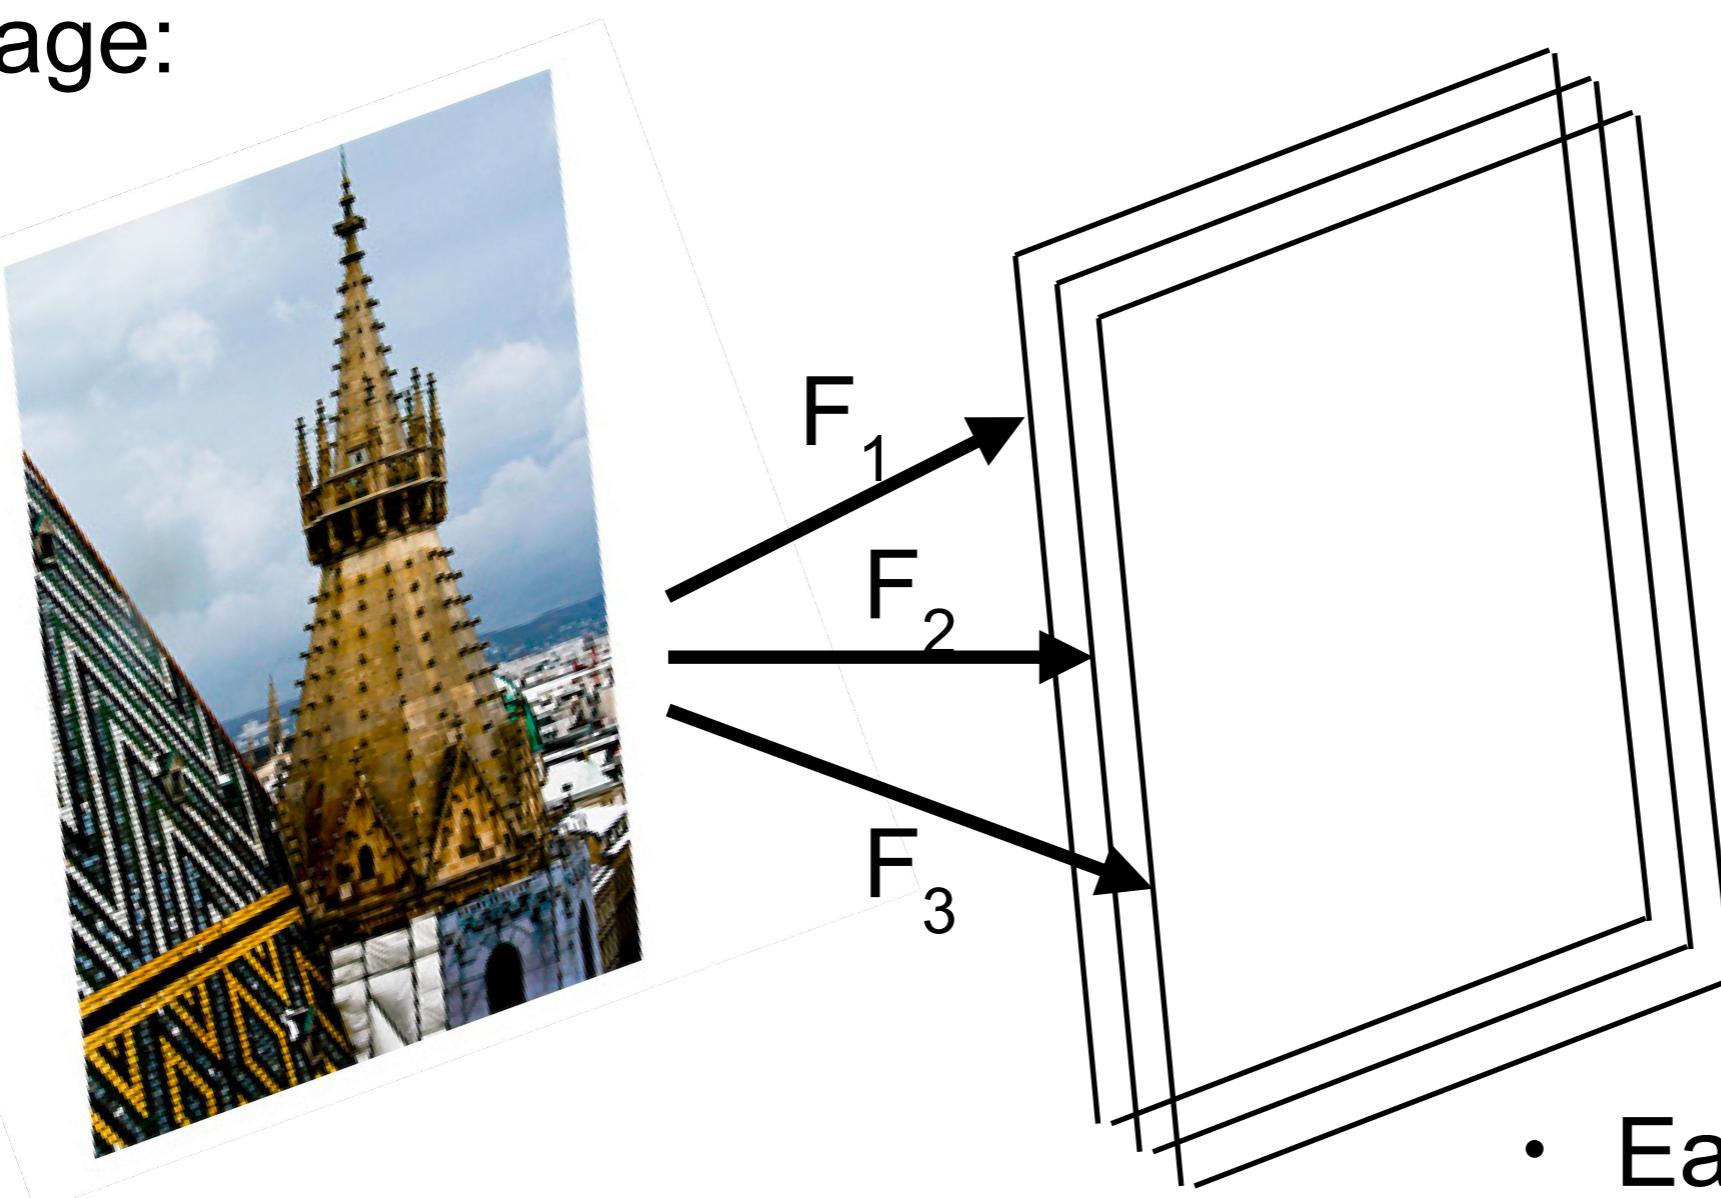
An
image:



- Collection of filters in the layer: *filter bank*

Convolutional Layer: multiple filters

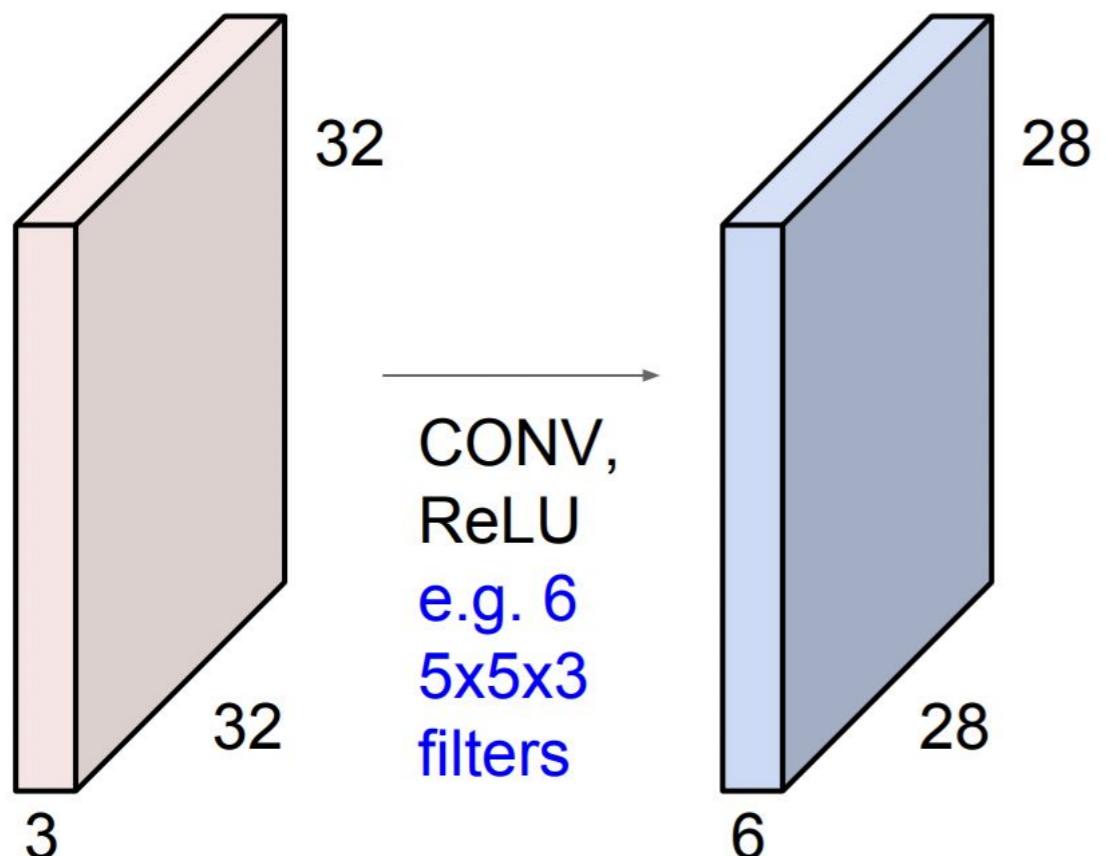
An
image:



- Collection of filters in the layer: *filter bank*
- Each resulting image is a *channel*

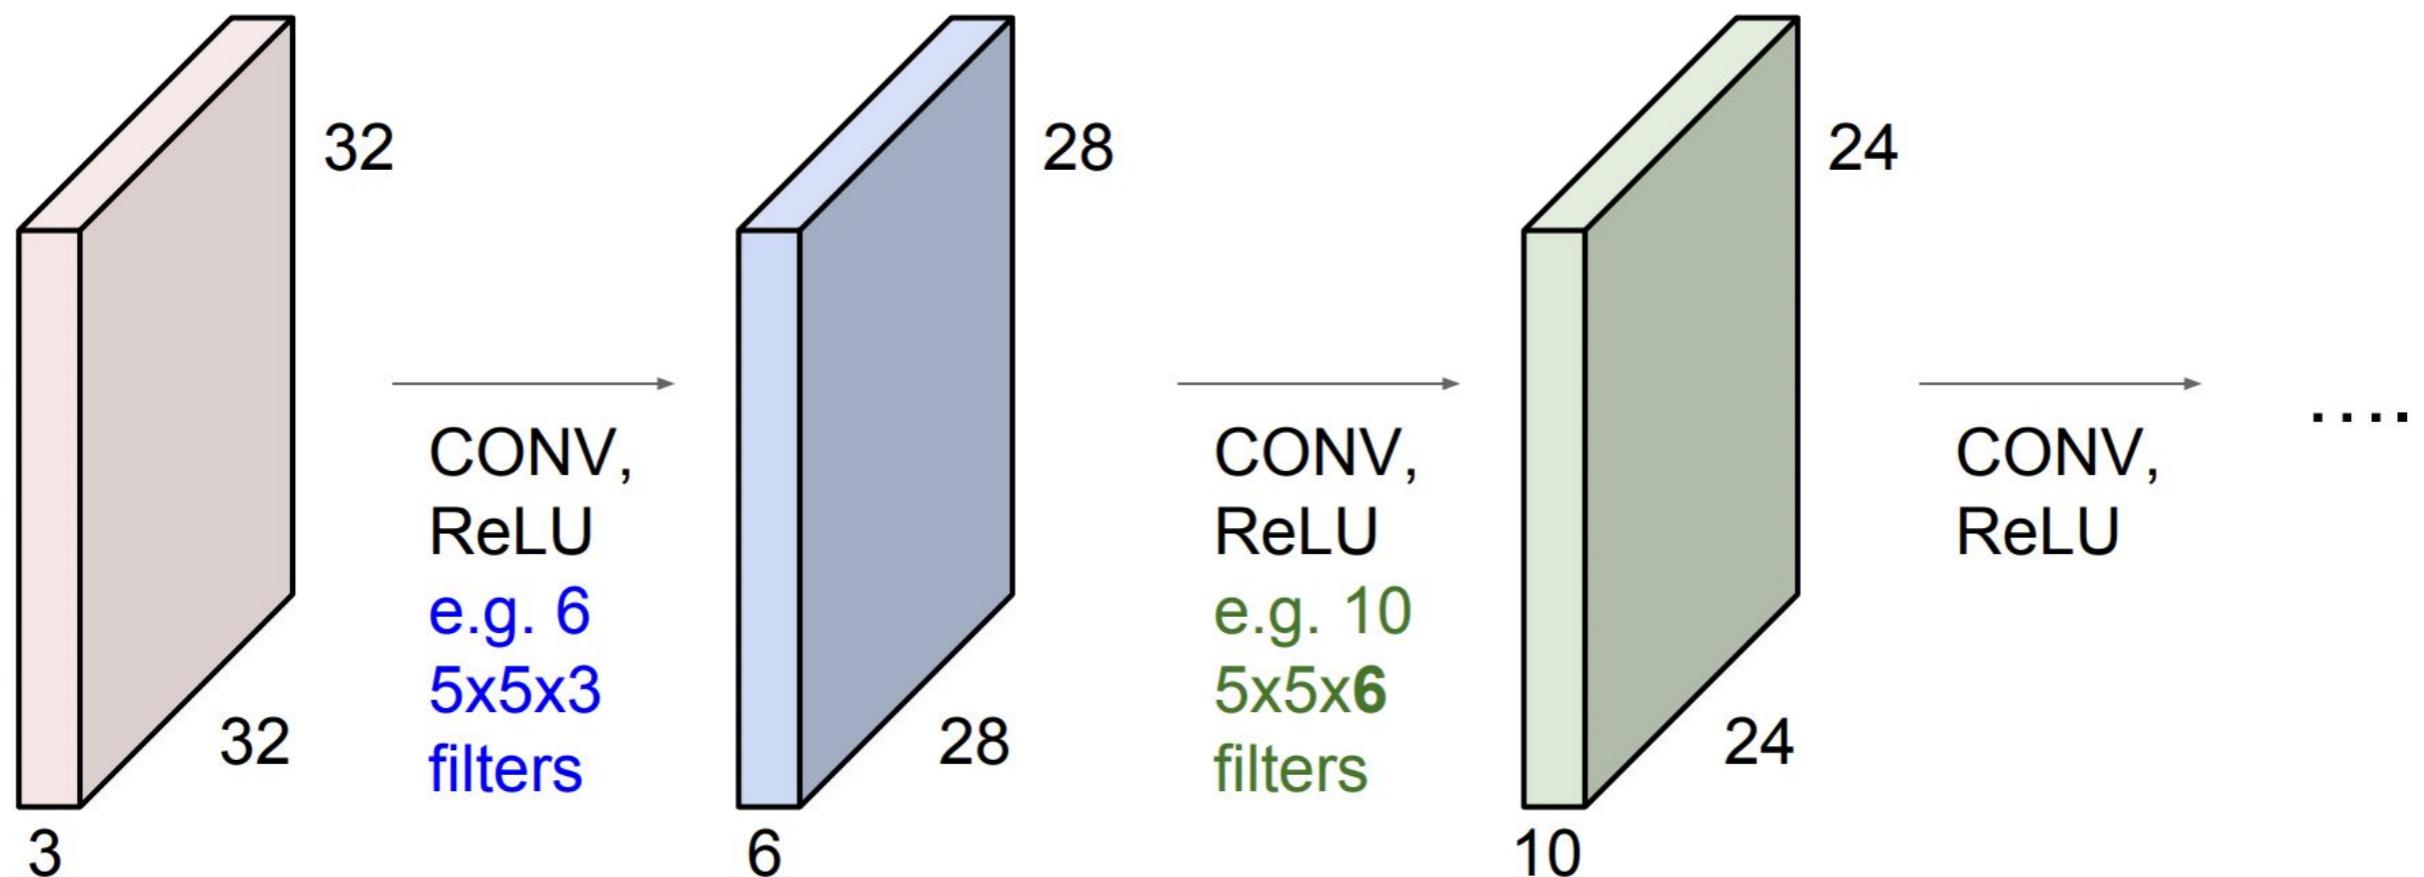
Conv Nets

Preview: ConvNet is a sequence of Convolution Layers, interspersed with activation functions



Conv Nets

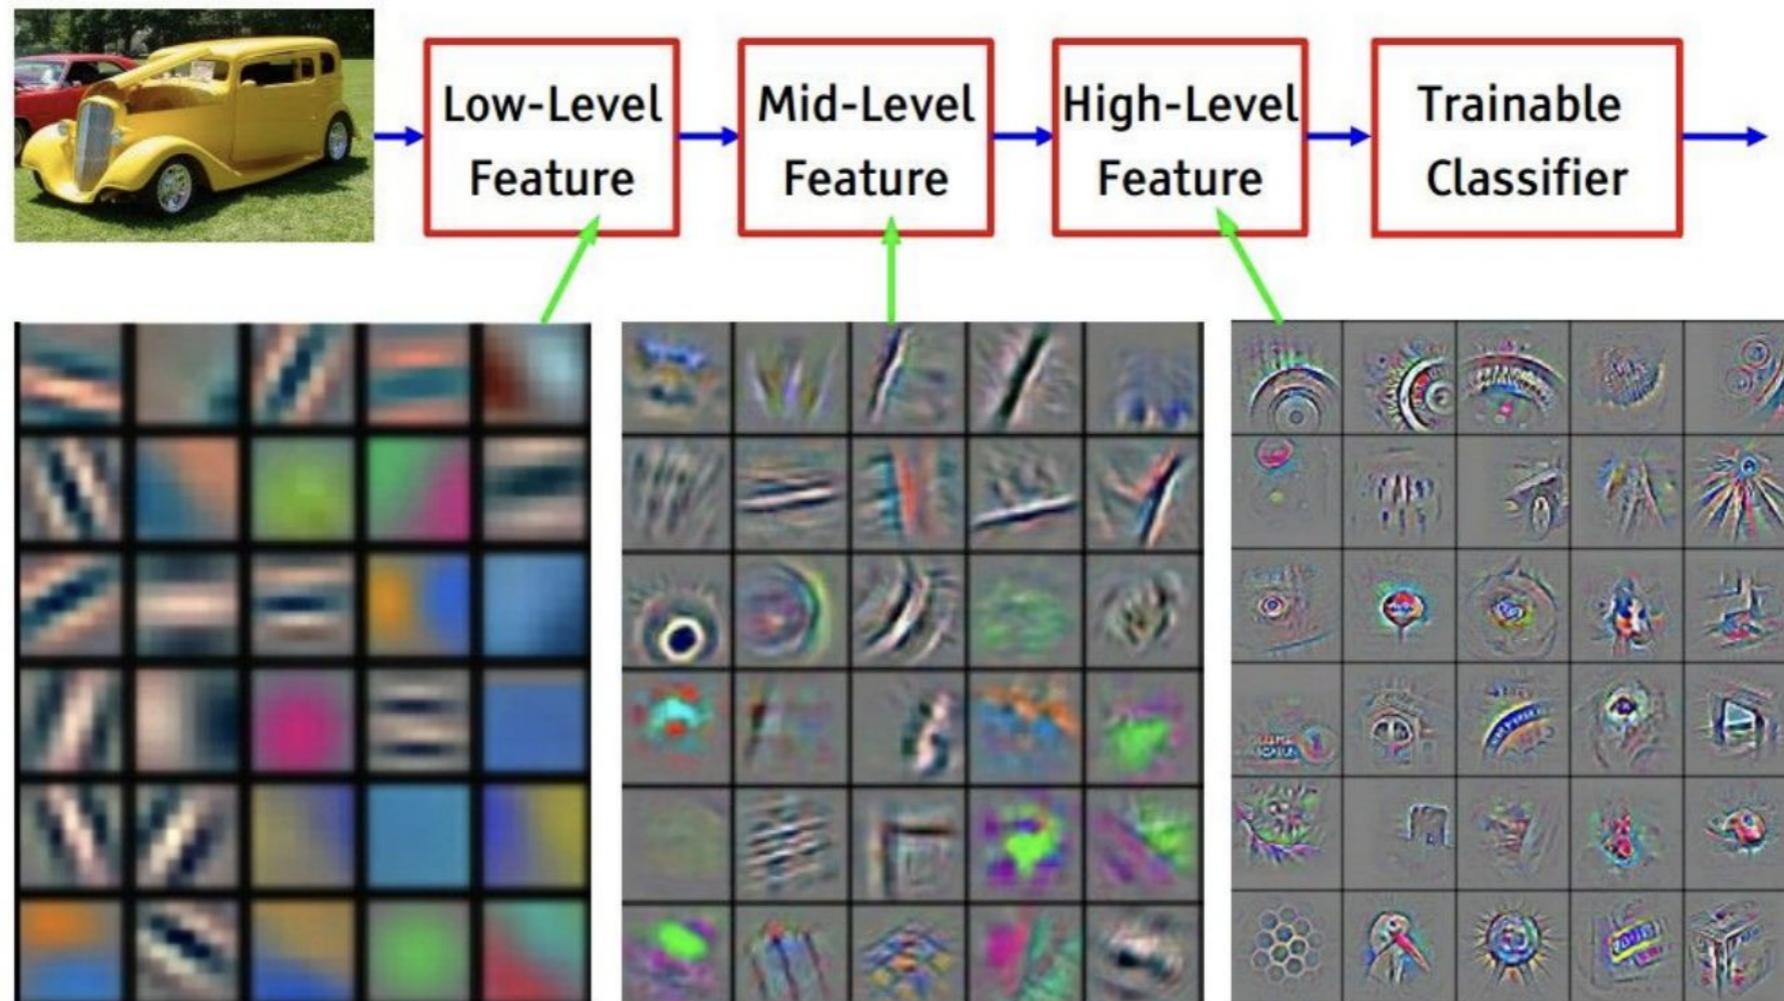
Preview: ConvNet is a sequence of Convolutional Layers, interspersed with activation functions



Conv Nets Learn Hierarchy

Preview

[From recent Yann LeCun slides]



Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]

Max pooling layer: 2D example

Output from the
convolutional
layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling layer: 2D example

Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)

Max pooling layer: 2D example

Output from the convolutional layer & ReLU:

0	0	0	0	0	0	0
0	0	0	0	1	0	
0	0	0	0	0	0	
0	1	0	0	0	0	
0	0	0	0	0	0	
0	0	0	0	0	0	

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)

After max pooling:

Max pooling layer: 2D example

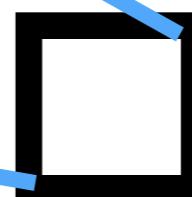
Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)

After max
pooling:



Max pooling layer: 2D example

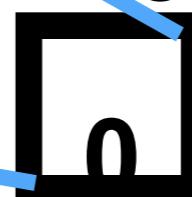
Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)

After max
pooling:



Max pooling layer: 2D example

Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	0	1
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)

After max
pooling:

0	1
---	---

Max pooling layer: 2D example

Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)

After max pooling:

0	1	
---	---	--

Max pooling layer: 2D example

Output from the convolutional layer & ReLU:

0	0	0	0	0	0	0
0	0	0	0	0	1	0
0	0	0	0	0	0	0
0	1	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)

After max pooling:

0	0	1	1
---	---	---	---

Max pooling layer: 2D example

Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)

After max pooling:

0	0	1	1
1	1	1	1
1	1	0	0
1	1	0	0

Max pooling layer: 2D example

Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)

After max pooling:

0	0	1	1
1	1	1	1
1	1	0	0
1	1	0	0

Max pooling layer: 2D example

Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)

After max pooling:

0	0	1	1
1	1	1	1
1	1	0	0
1	1	0	0

Max pooling layer: 2D example

Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)

After max pooling:

0	0	1	1
1	1	1	1
1	1	0	0
1	1	0	0

Max pooling layer: 2D example

Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)
- stride 1

After max
pooling

1	0	10	1	1	1
1	1	0	0		
1	1	0	0		

Max pooling layer: 2D example

Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)
- stride 3

After max pooling:

Max pooling layer: 2D example

Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)
- stride 3

After max
pooling

Max pooling layer: 2D example

Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)
- stride 3

After max
pooling:



Max pooling layer: 2D example

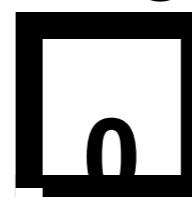
Output from the convolutional layer & ReLU:

0	0	0	0	0	0	0
0	0	0	0	1	0	
0	0	0	0	0	0	
0	1	0	0	0	0	
0	0	0	0	0	0	
0	0	0	0	0	0	
0	0	0	0	0	0	

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)
- stride 3

After max pooling:



Max pooling layer: 2D example

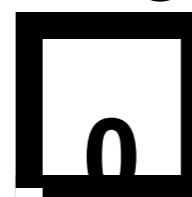
Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)
- stride 3

After max pooling:



Max pooling layer: 2D example

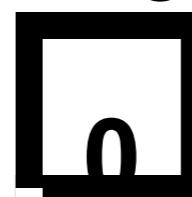
Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)
- stride 3

After max pooling:



Max pooling layer: 2D example

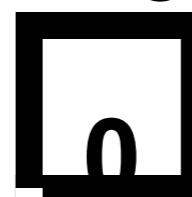
Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)
- stride 3

After max pooling:



Max pooling layer: 2D example

Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)
- stride 3

After max pooling:

0

Max pooling layer: 2D example

Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)
- stride 3

After max pooling:

0	1
---	---

Max pooling layer: 2D example

Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)
- stride 3

After max pooling:

0	1
---	---

Max pooling layer: 2D example

Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)
- stride 3

After max pooling:

0	1
---	---

Max pooling layer: 2D example

Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)
- stride 3

After max pooling:

0	1
---	---

Max pooling layer: 2D example

Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)
- stride 3

After max pooling:

0	1
---	---

Max pooling layer: 2D example

Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)
- stride 3

After max pooling:

0	1
---	---

Max pooling layer: 2D example

Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)
- stride 3

After max pooling:

0	1

Max pooling layer: 2D example

Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)
- stride 3

After max pooling:

0	1
1	

Max pooling layer: 2D example

Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)
- stride 3

After max pooling:

0	1
1	

Max pooling layer: 2D example

Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)
- stride 3

After max pooling:

0	1
1	

Max pooling layer: 2D example

Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)
- stride 3

After max pooling:

0	1
1	

Max pooling layer: 2D example

Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)
- stride 3

After max pooling:

0	1
1	

Max pooling layer: 2D example

Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)
- stride 3

After max pooling:

0	1
1	

Max pooling layer: 2D example

Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)
- stride 3

After max pooling:

0	1
1	

Max pooling layer: 2D example

Output from the convolutional layer & ReLU:

0	0	0	0	0	0
0	0	0	0	1	0
0	0	0	0	0	0
0	1	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Max pooling: returns max of its arguments

- size 3x3 (“size 3”)
- stride 3

After max pooling:

0	1
1	

- Can use stride with filters too
- No weights in max pooling

Convolution (General Formulae)

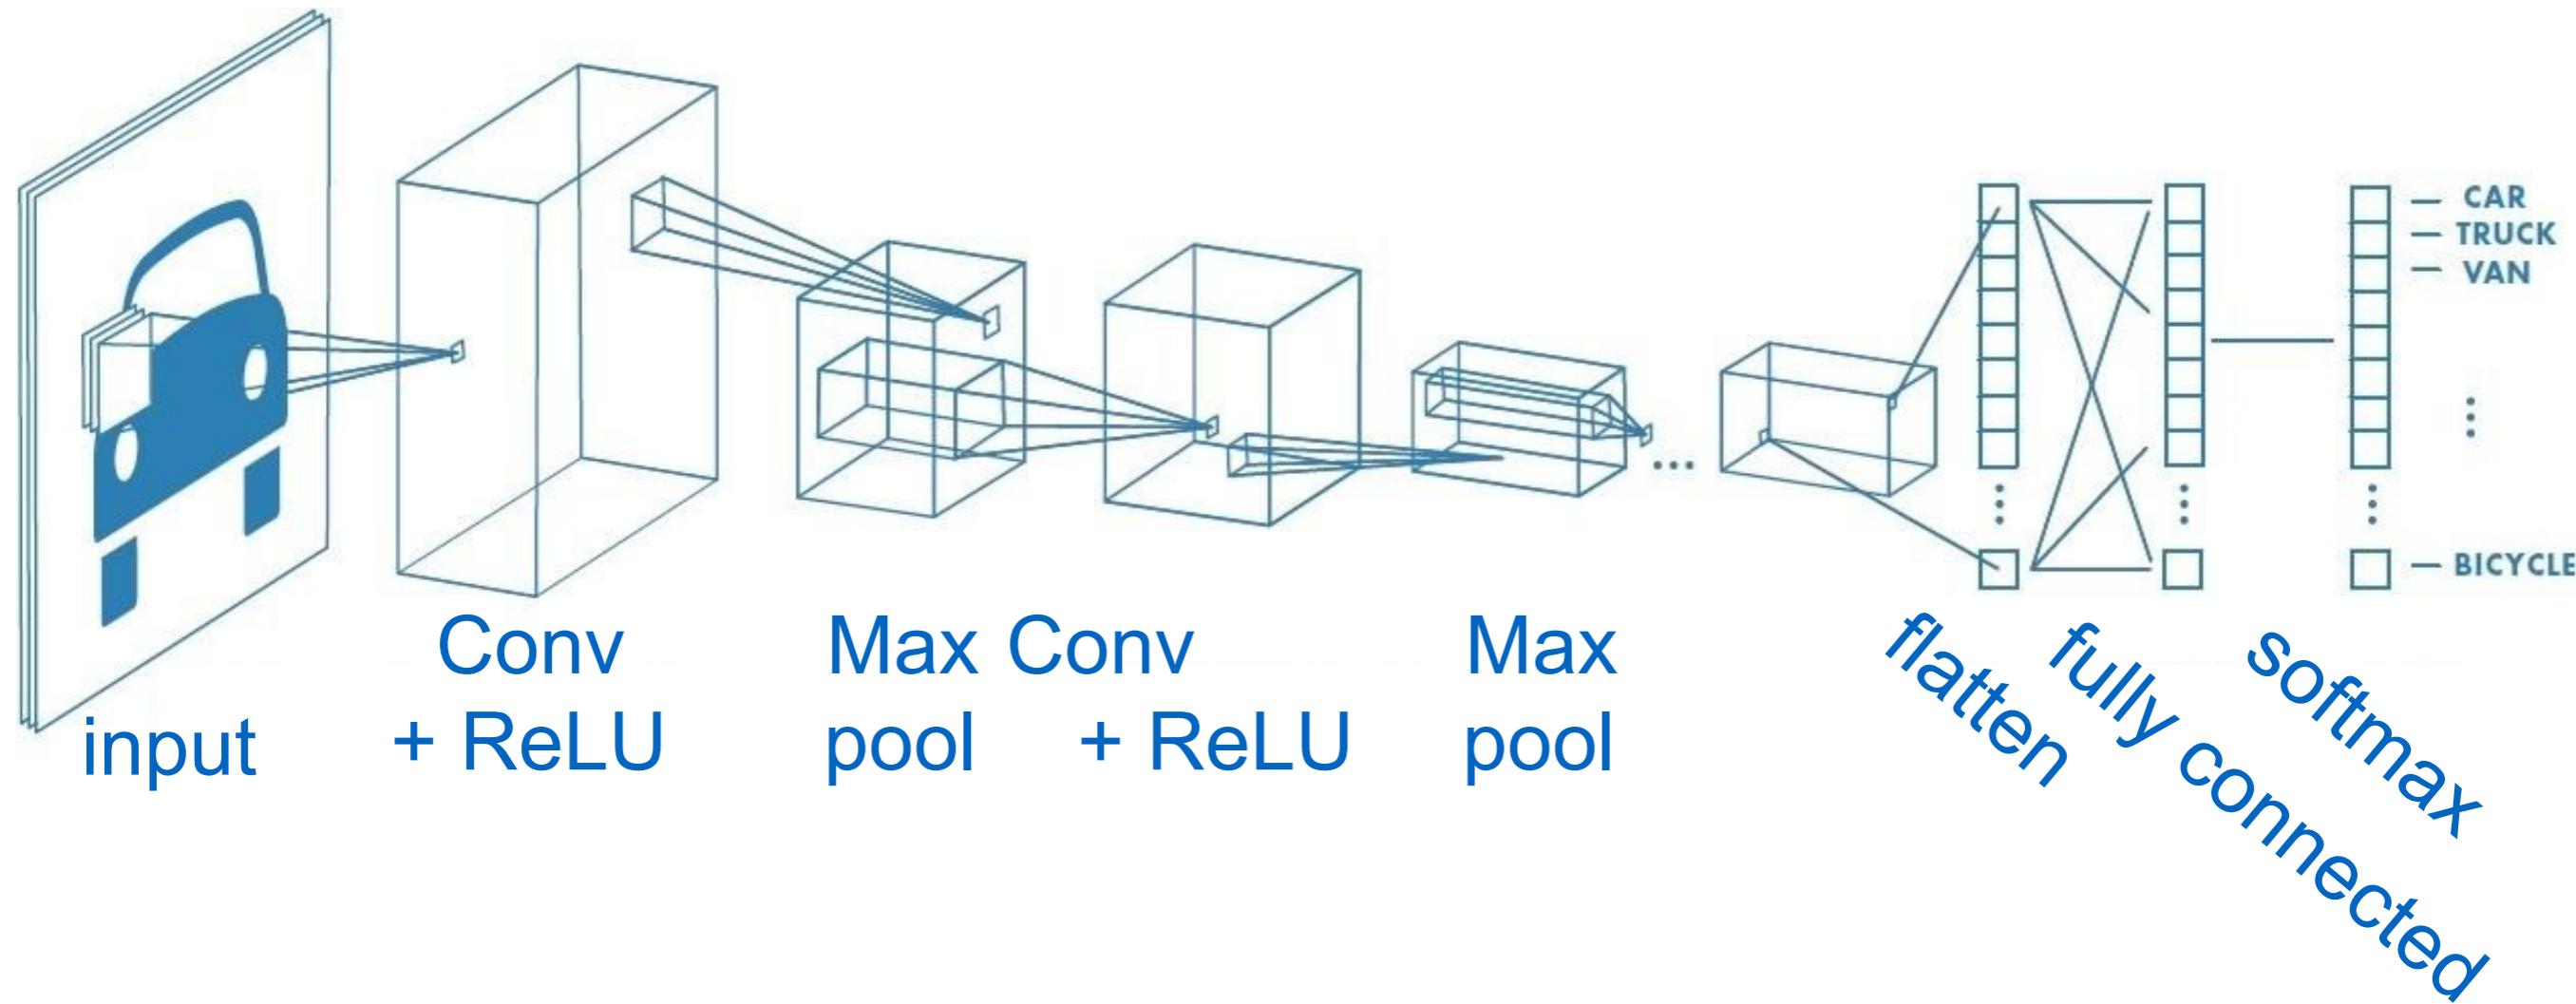
Summary. To summarize, the Conv Layer:

- Accepts a volume of size $W_1 \times H_1 \times D_1$
- Requires four hyperparameters:
 - Number of filters K ,
 - their spatial extent F ,
 - the stride S ,
 - the amount of zero padding P .
- Produces a volume of size $W_2 \times H_2 \times D_2$ where:
 - $W_2 = (W_1 - F + 2P)/S + 1$
 - $H_2 = (H_1 - F + 2P)/S + 1$ (i.e. width and height are computed equally by symmetry)
 - $D_2 = K$
- With parameter sharing, it introduces $F \cdot F \cdot D_1$ weights per filter, for a total of $(F \cdot F \cdot D_1) \cdot K$ weights and K biases.
- In the output volume, the d -th depth slice (of size $W_2 \times H_2$) is the result of performing a valid convolution of the d -th filter over the input volume with a stride of S , and then offset by d -th bias.

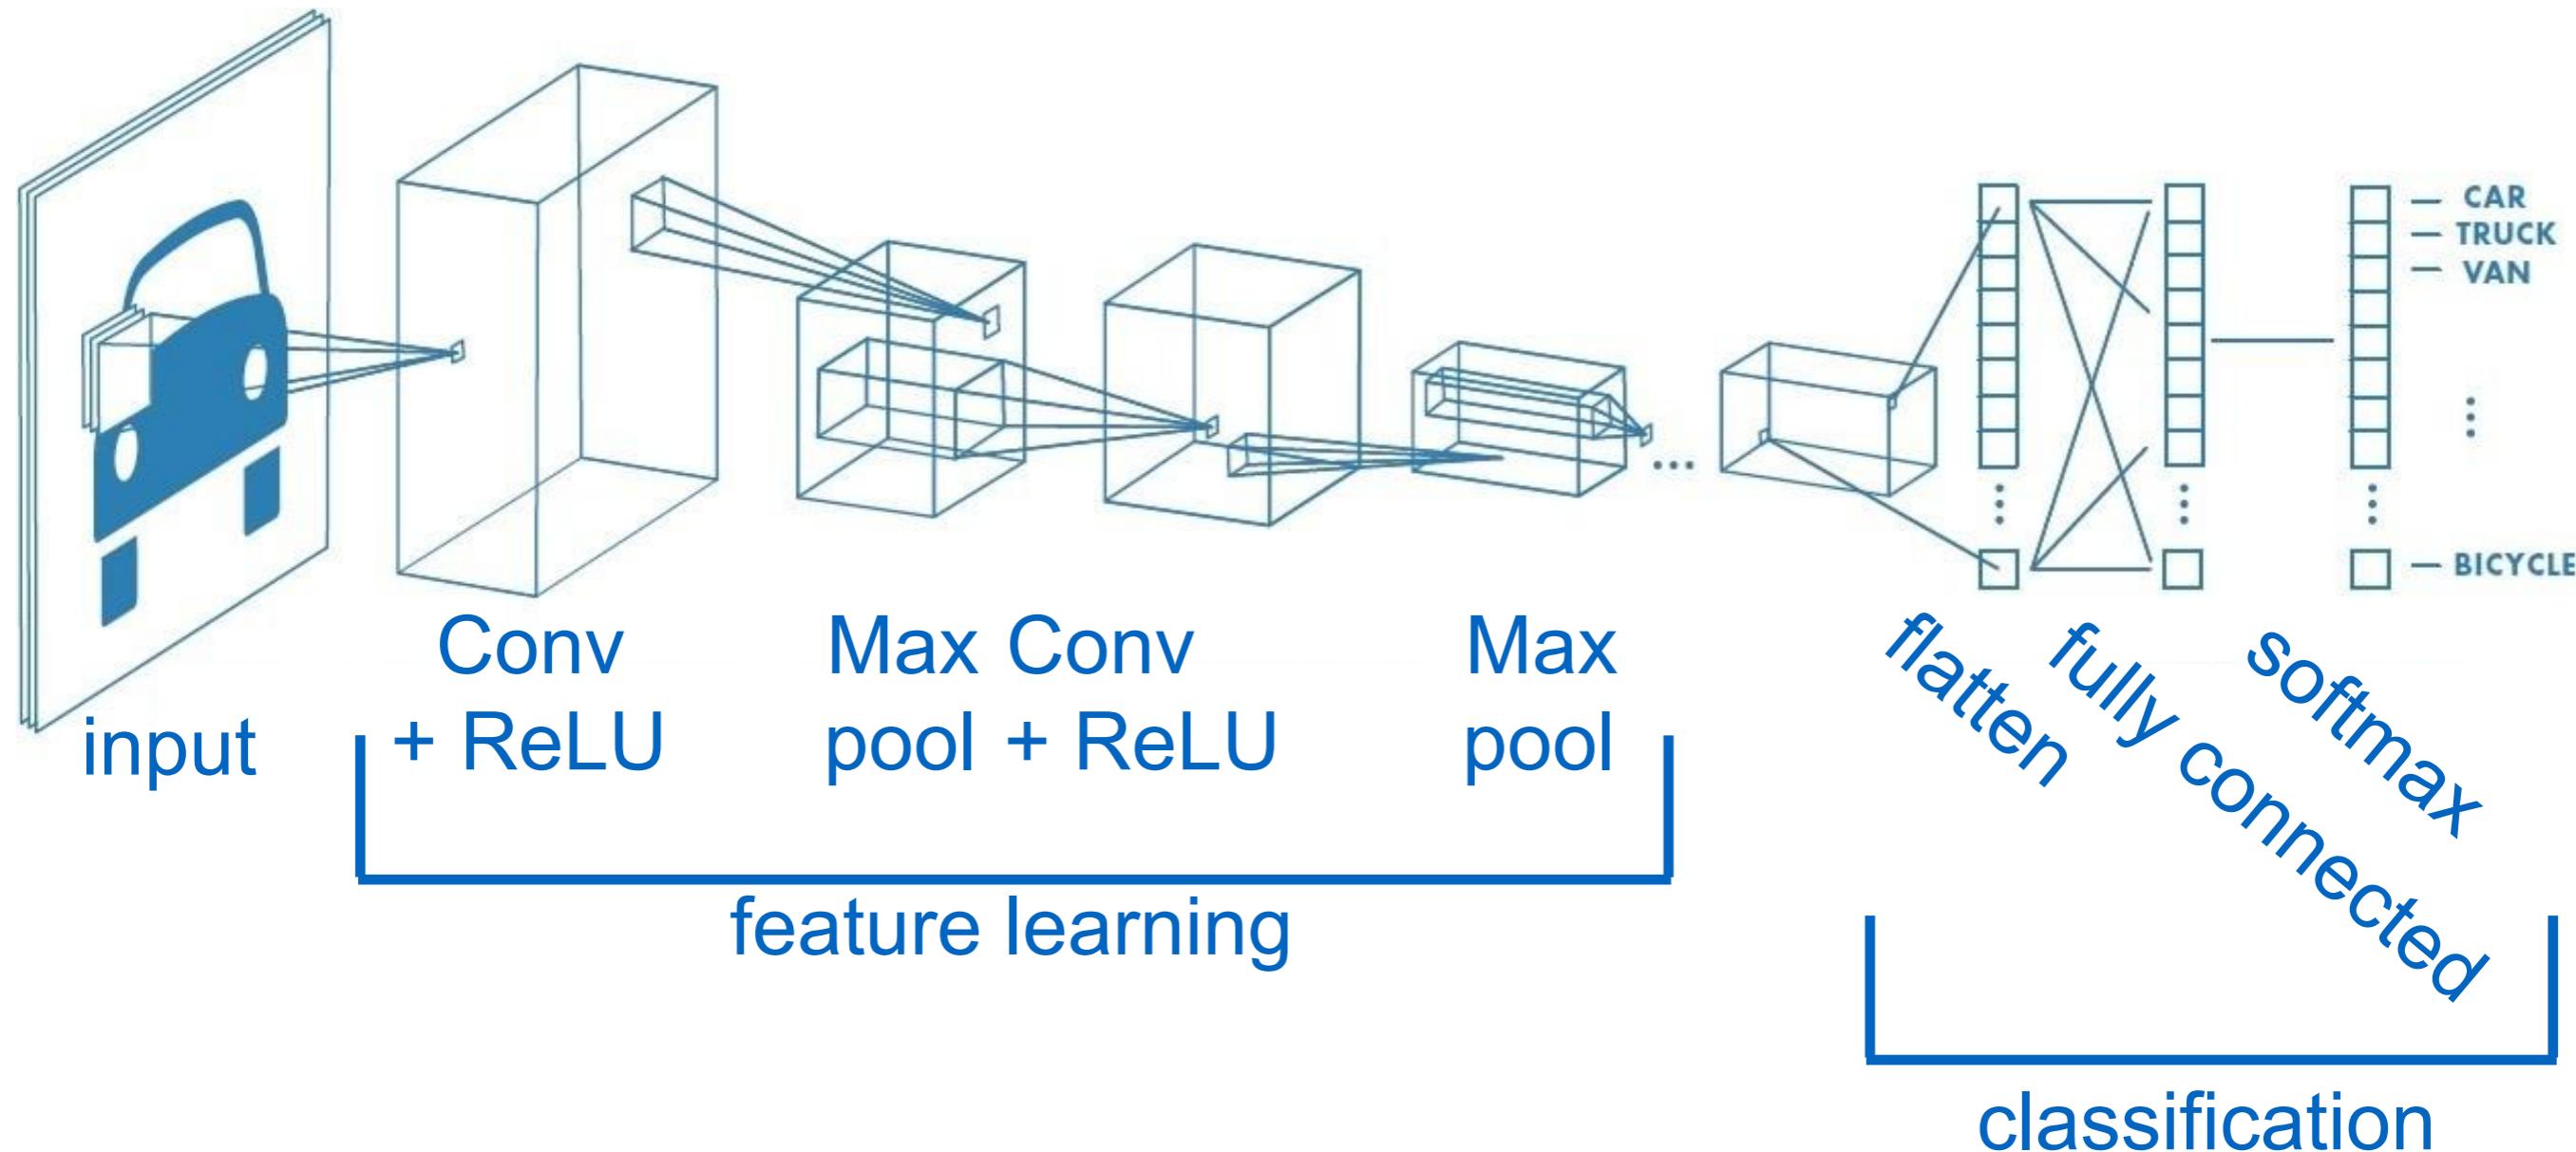
Common settings:

- $K = (\text{powers of 2, e.g. } 32, 64, 128, 512)$
- $F = 3, S = 1, P = 1$
 - $F = 5, S = 1, P = 2$
 - $F = 5, S = 2, P = ?$ (whatever fits)
 - $F = 1, S = 1, P = 0$

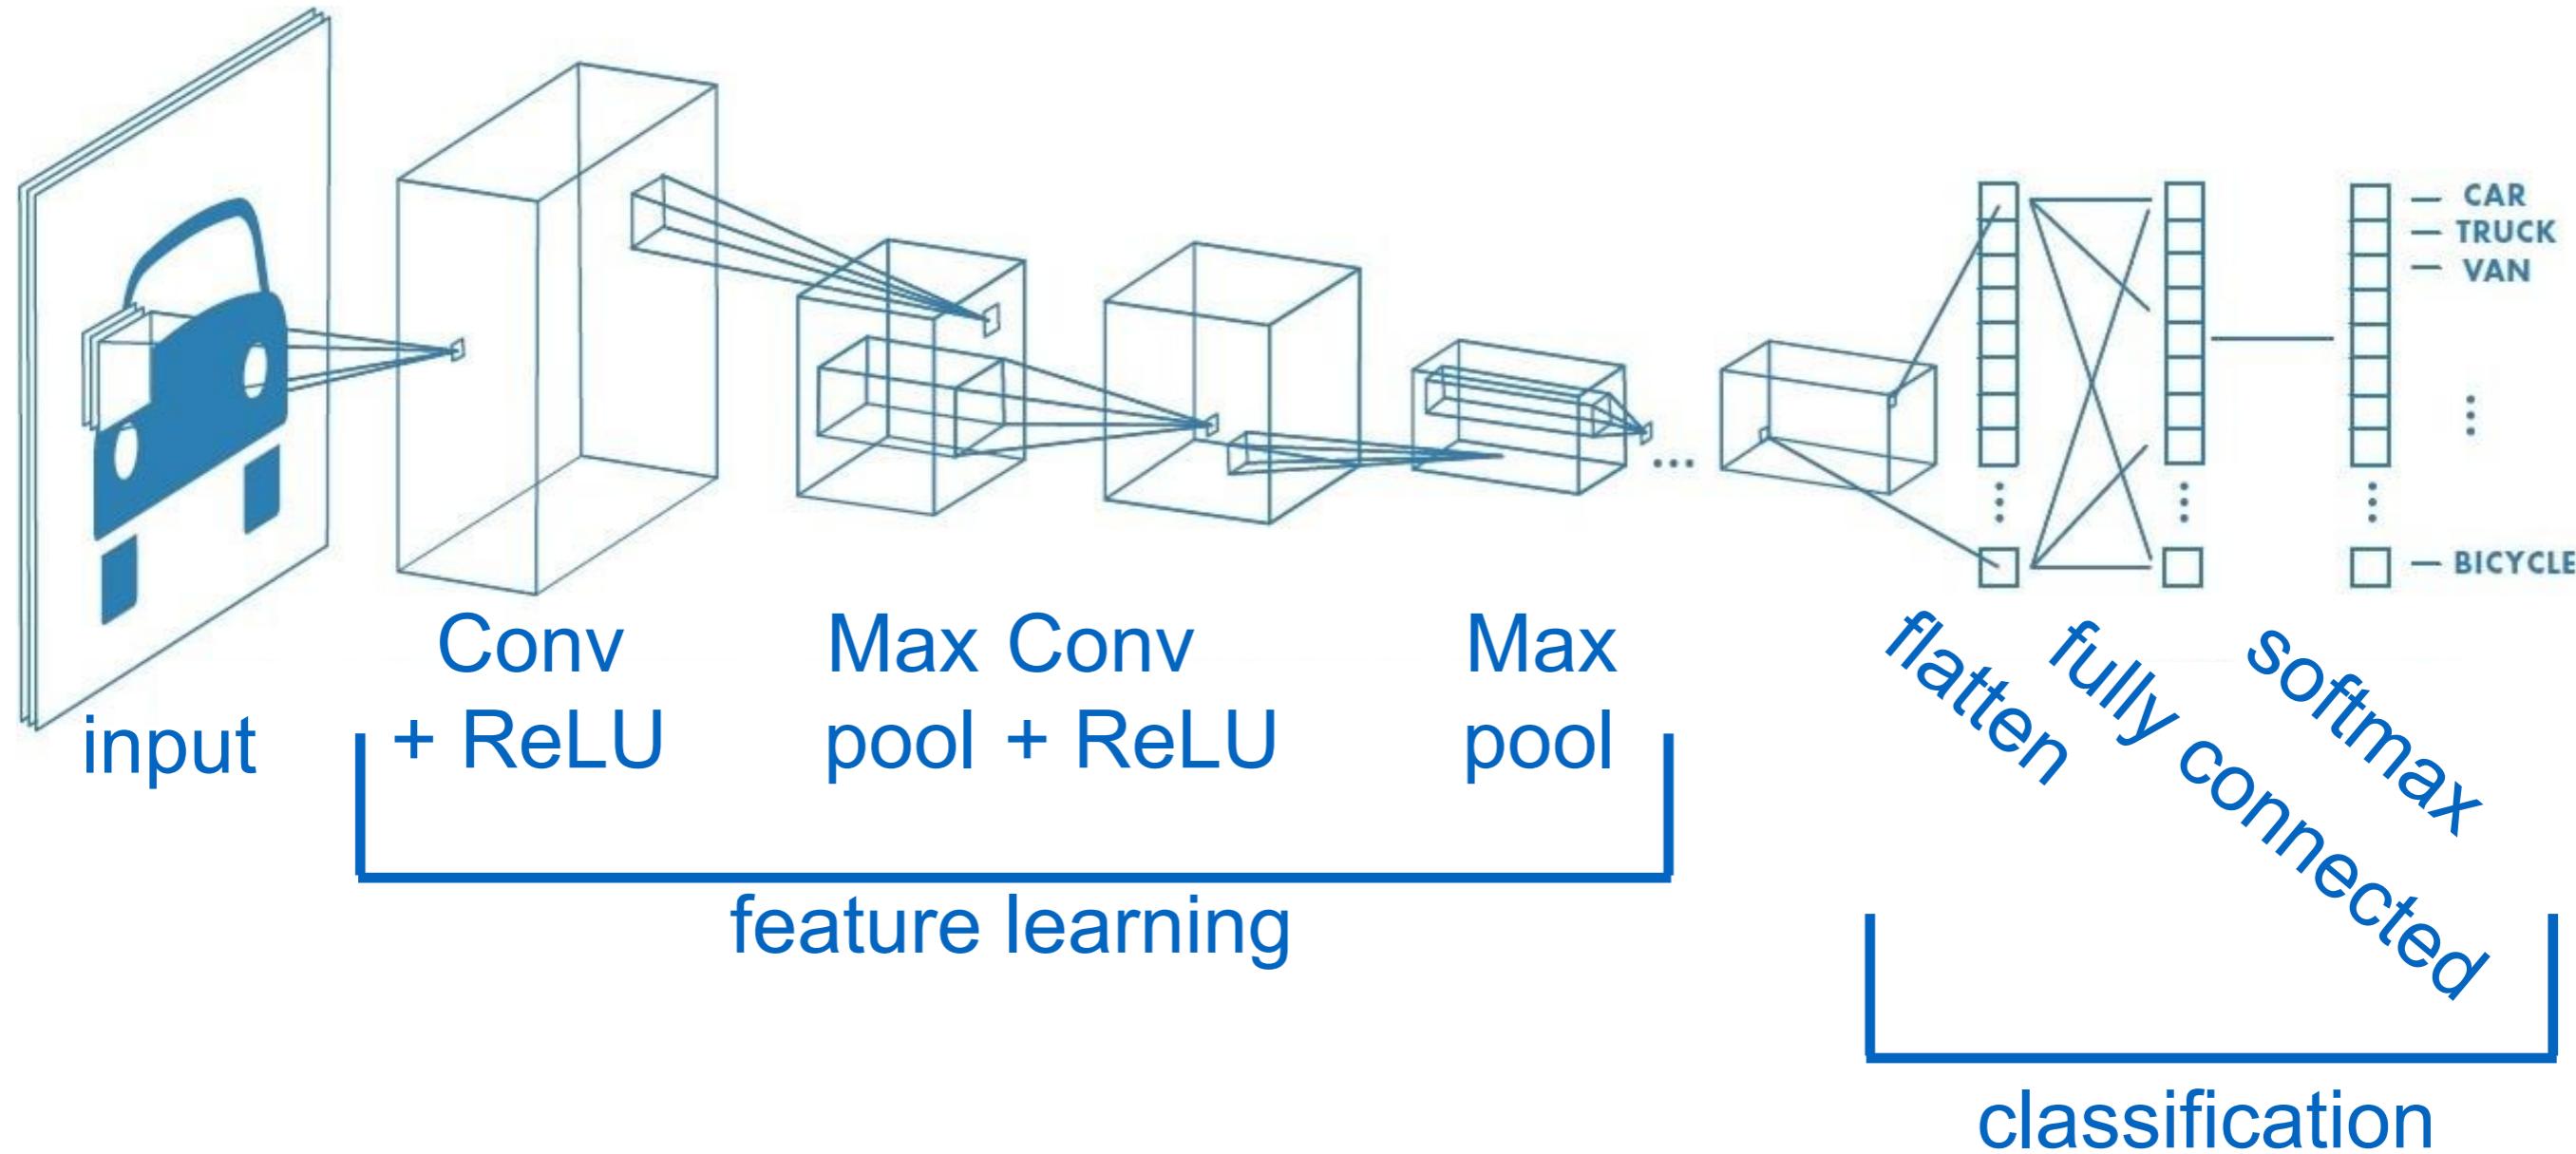
CNNs: typical architecture



CNNs: typical architecture



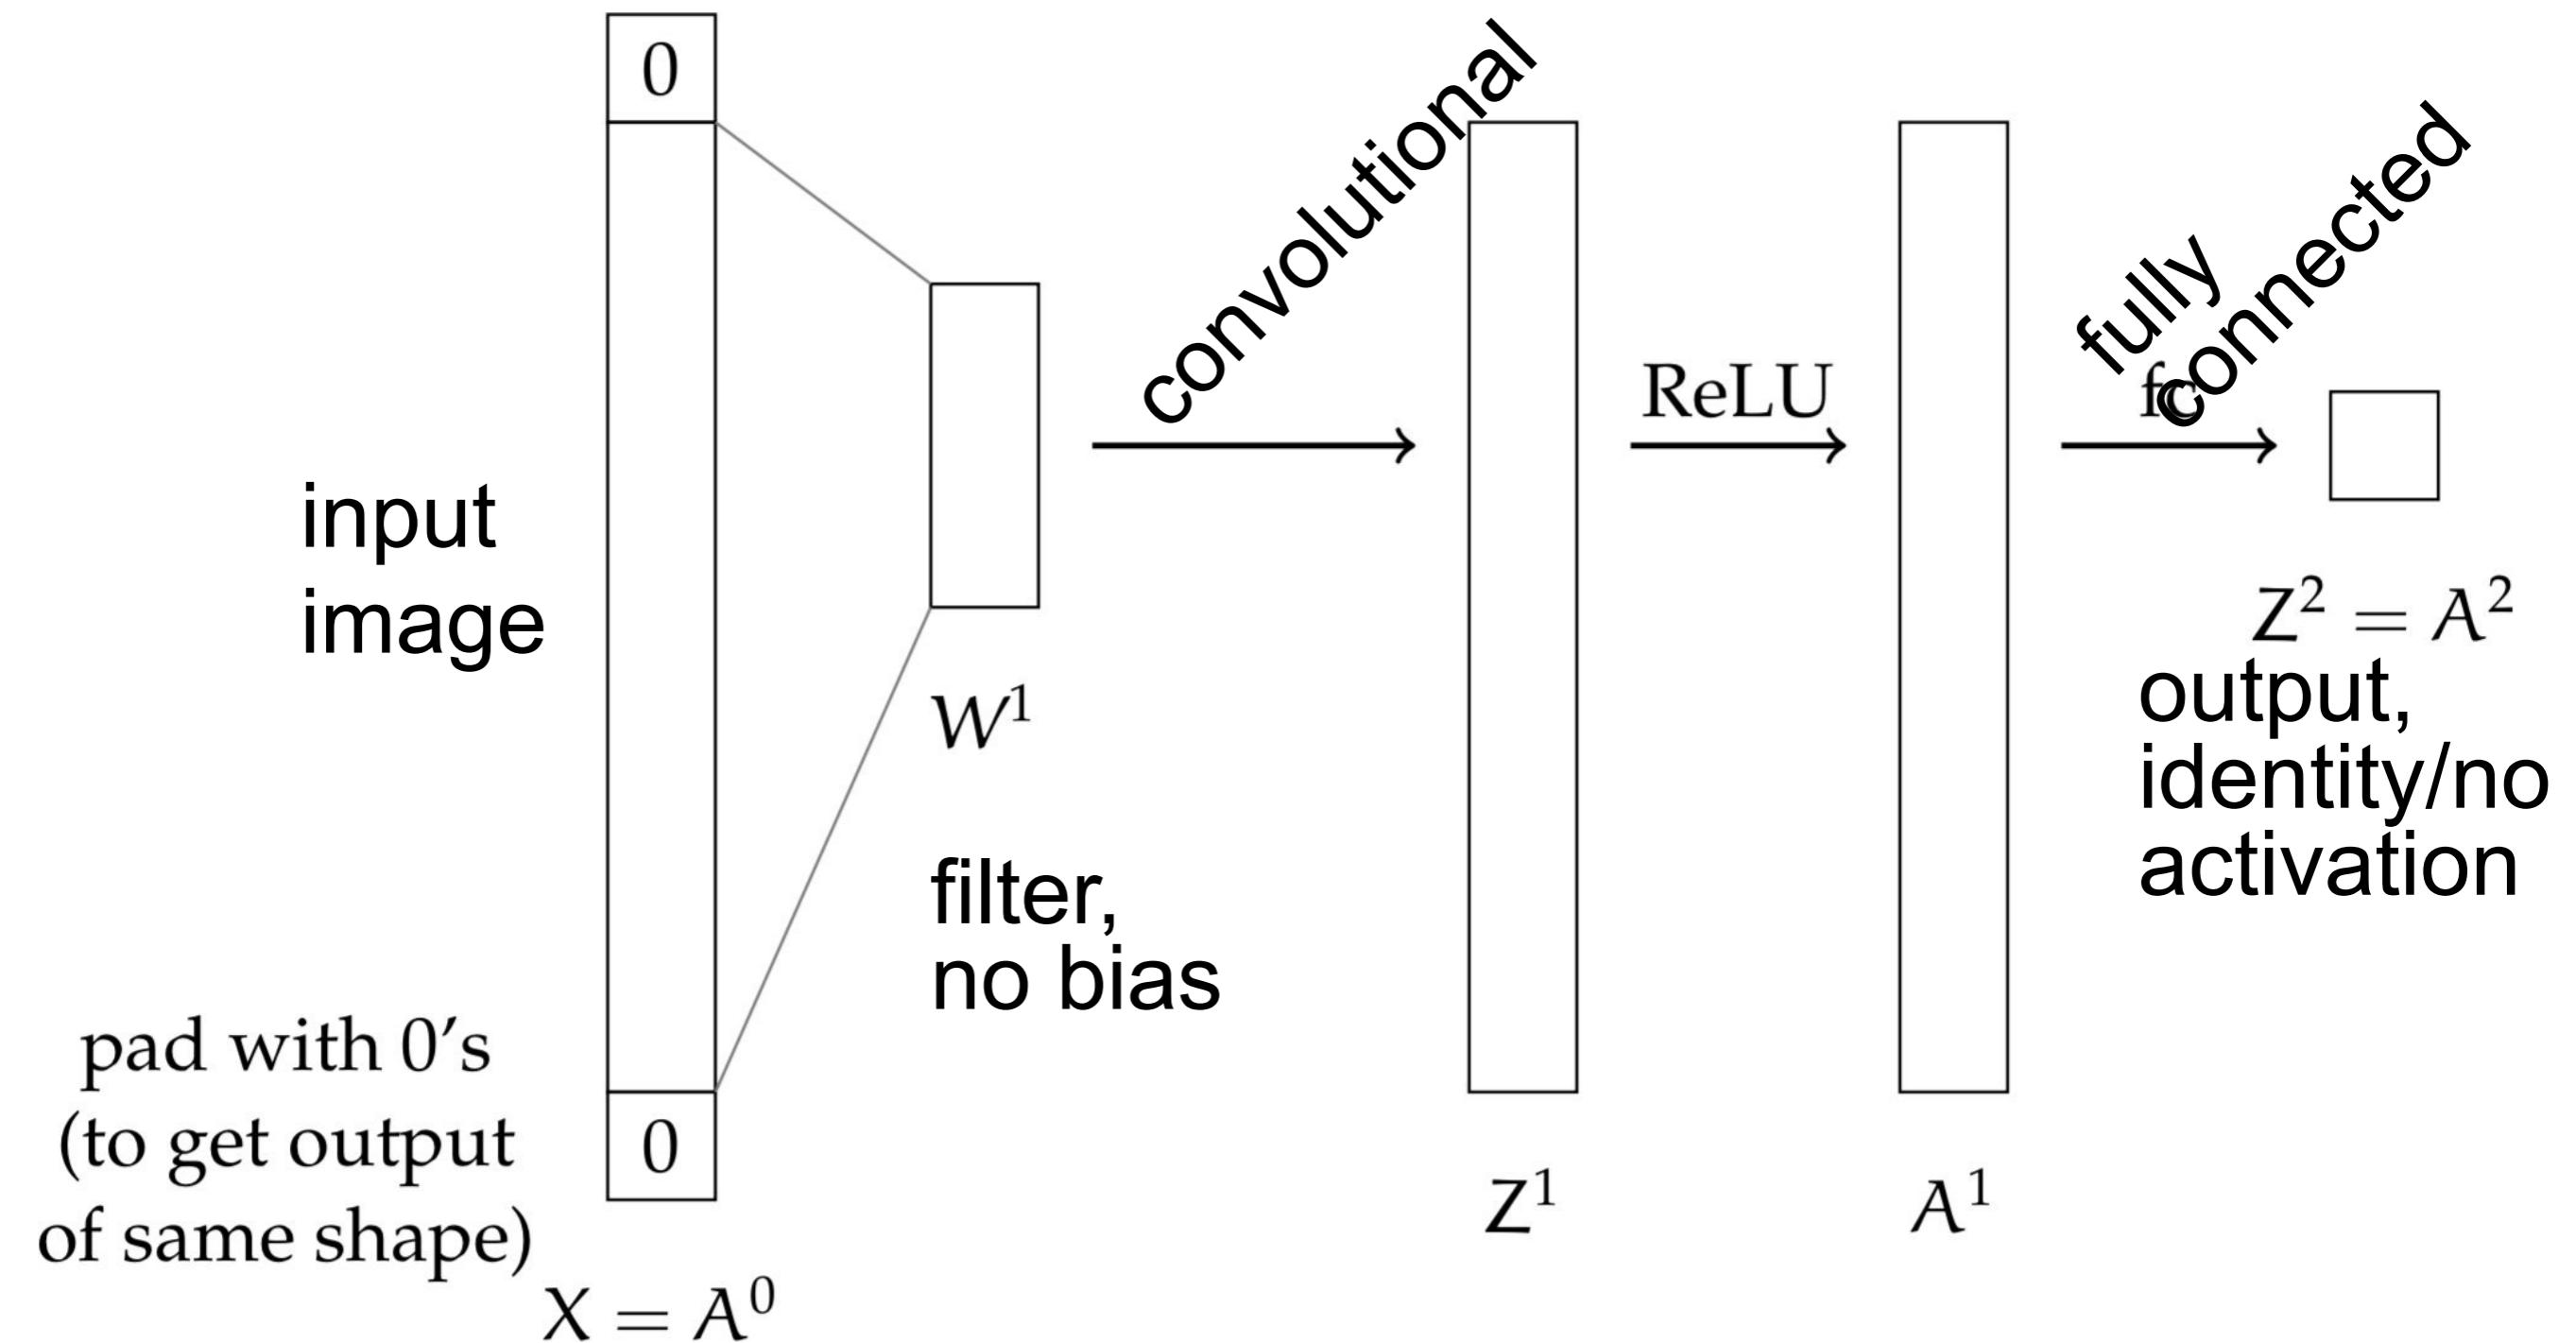
CNNs: typical architecture



Recall: we wanted to encode

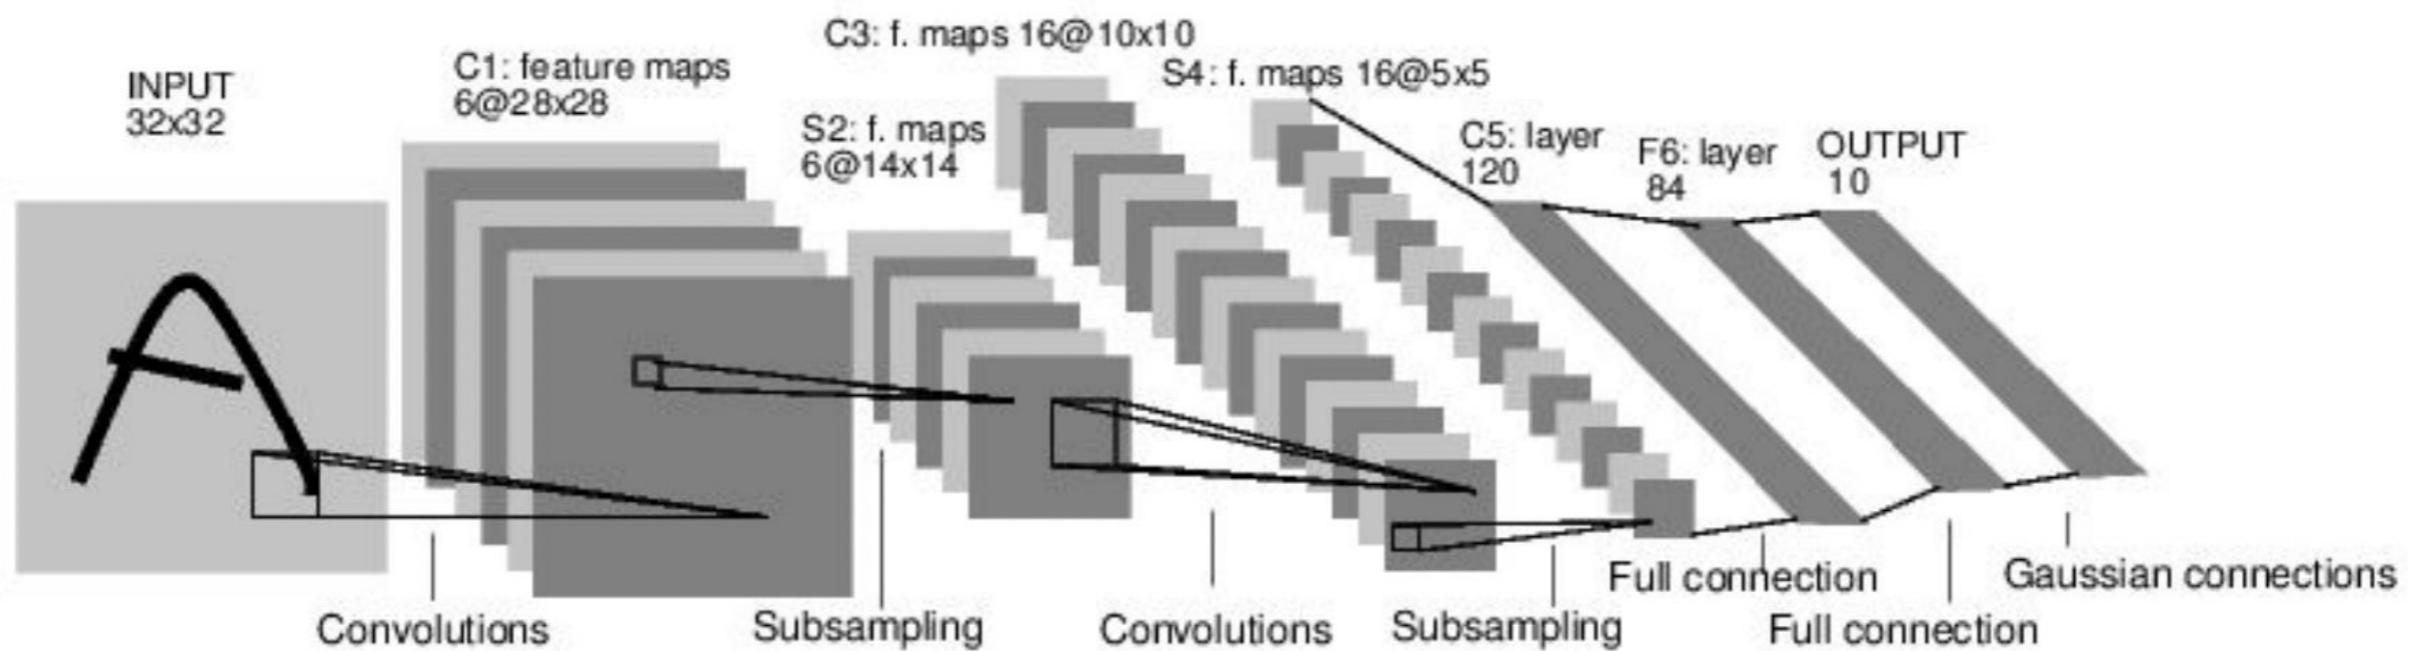
- Spatial locality
- Translation invariance

CNNs: a taste of backpropagation



Case Study: LeNet 5

[LeCun et al., 1998]



Conv filters were 5×5 , applied at stride 1

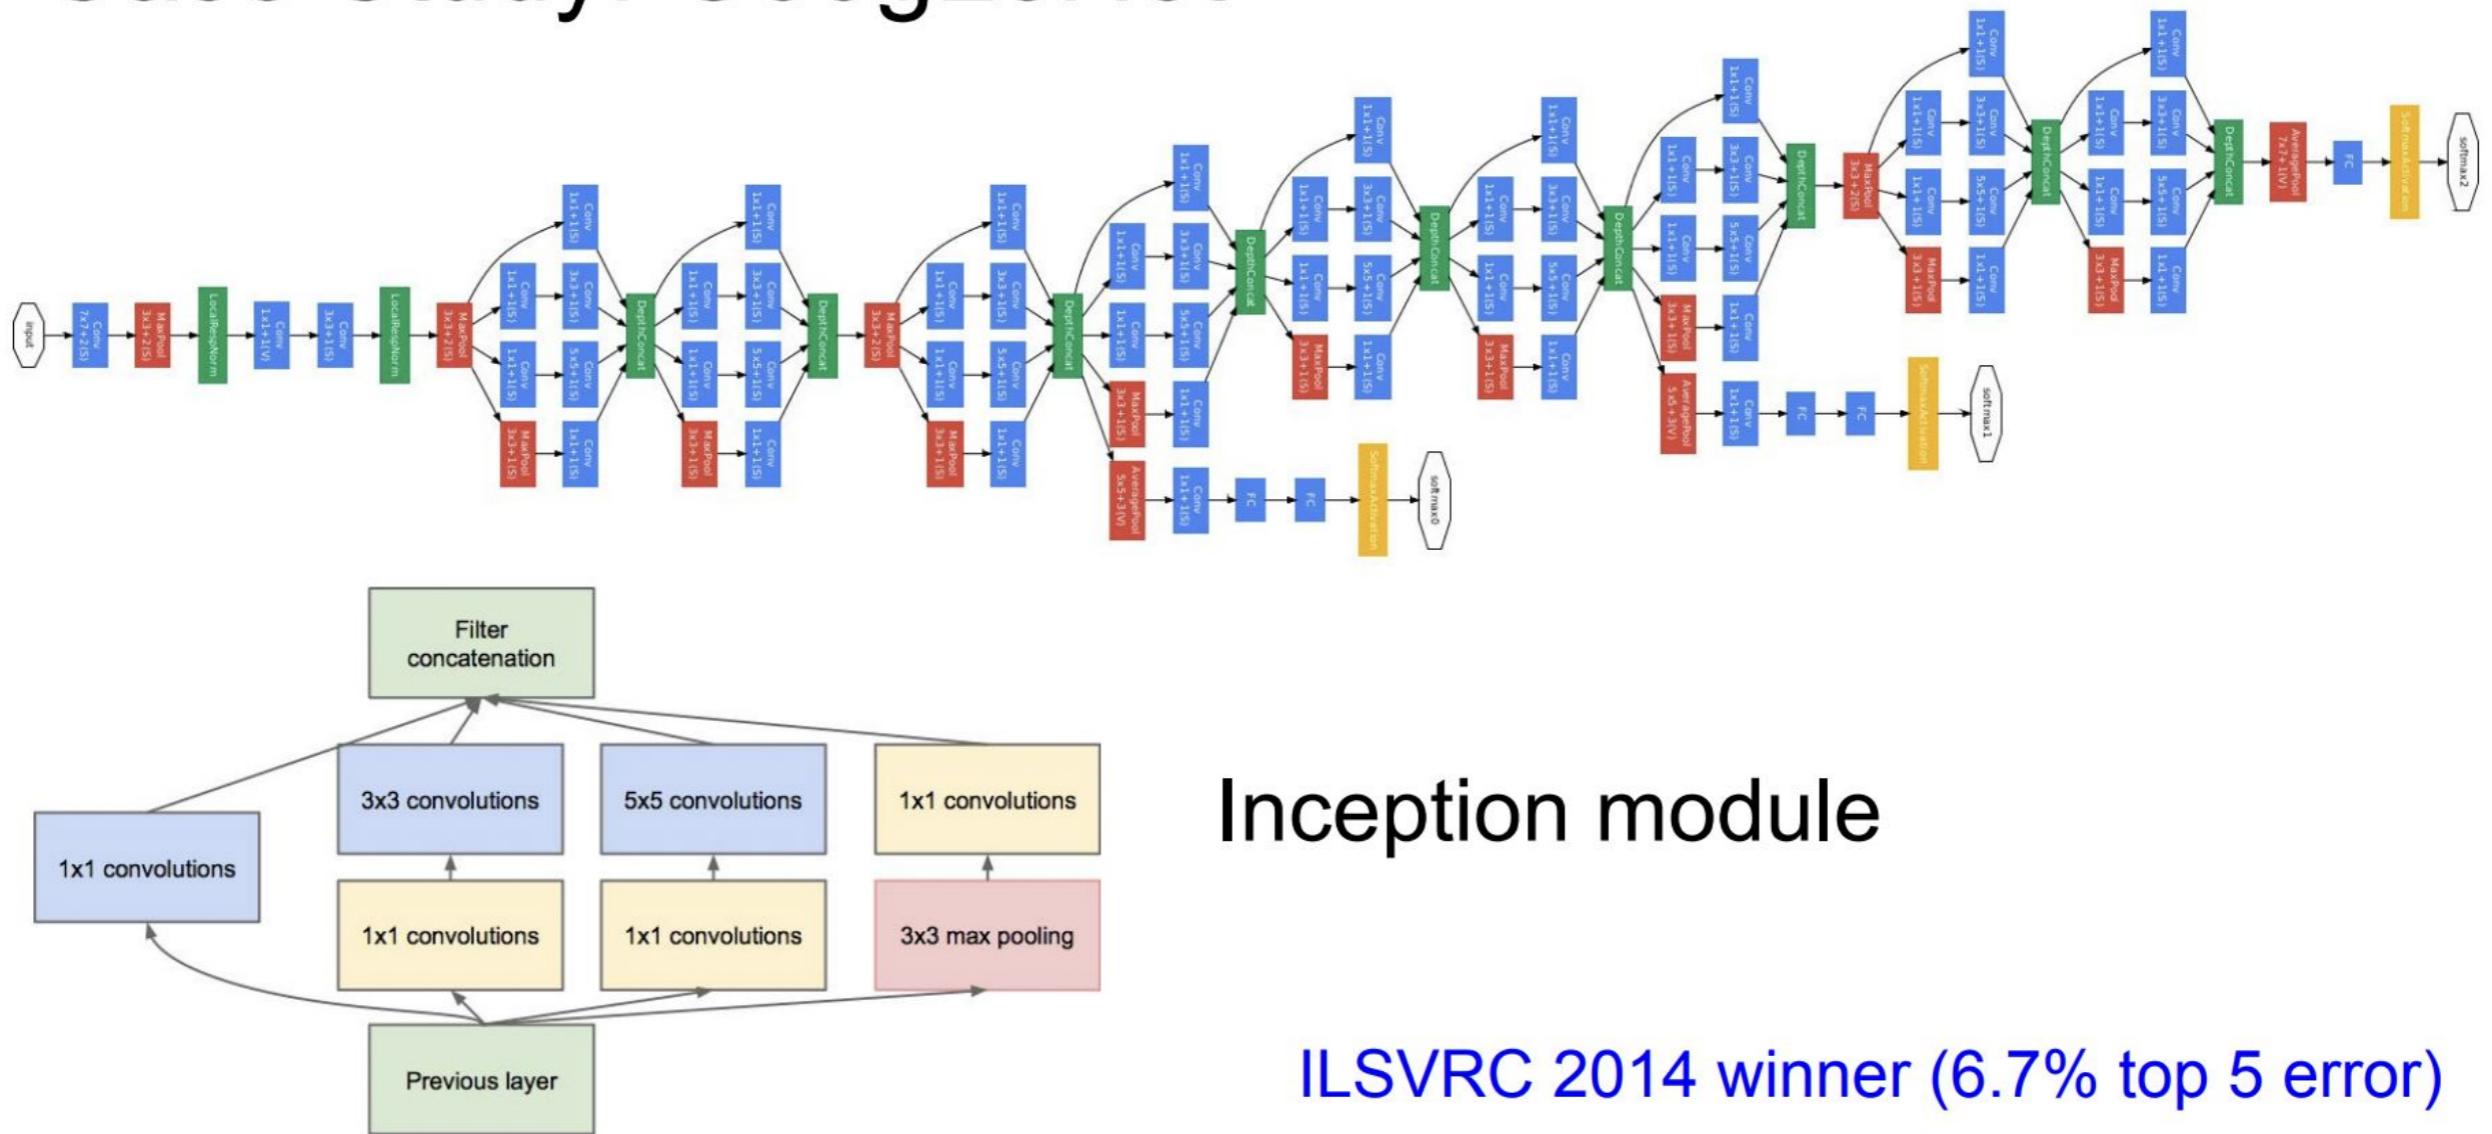
Subsampling (Pooling) layers were 2×2 applied at stride 2

i.e. architecture is [CONV-POOL-CONV-POOL-CONV-FC]

Case Study: GoogleNet

Case Study: GoogLeNet

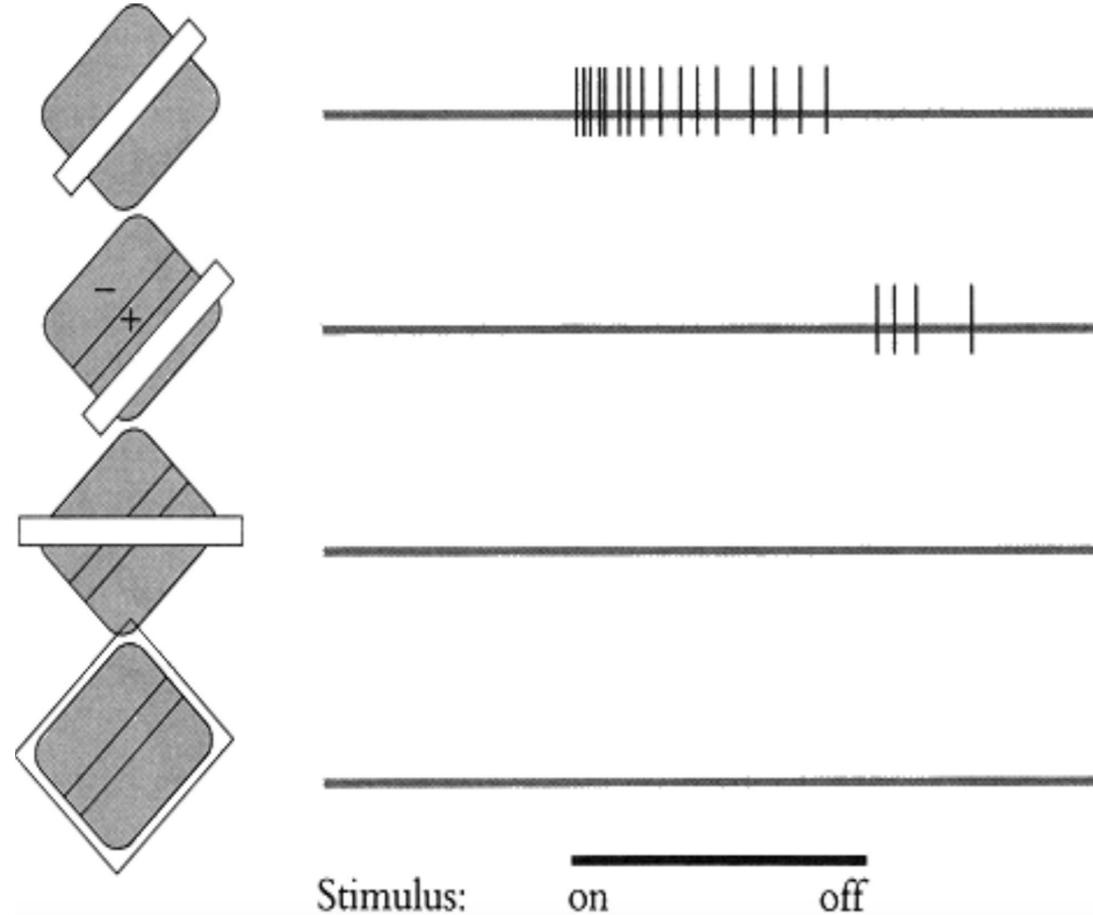
[Szegedy et al., 2014]



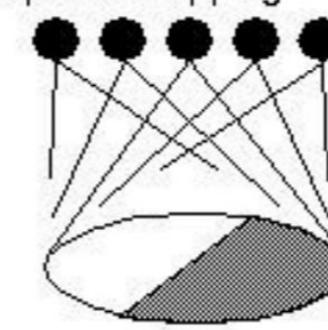
Cat neurons [Hubel, Weisel 1959, 1962]

(Be careful with biology analogies)

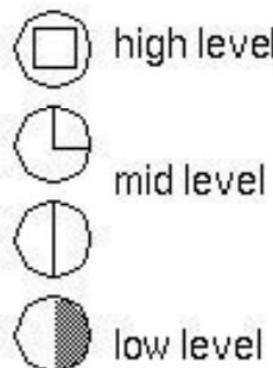
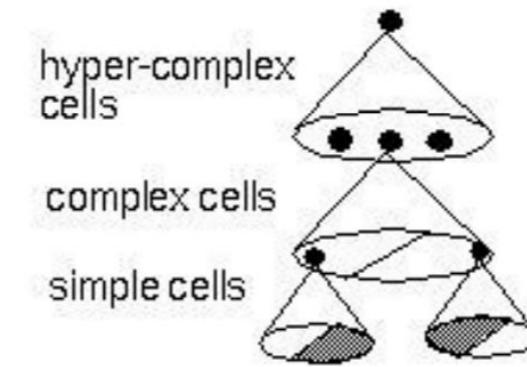
receptive field



Hubel & Weisel
topographical mapping



featural hierarchy



- *simple cells*
- *complex cells*

[<http://fourier.eng.hmc.edu/e180/lectures/v1/node7.html>]

Summary

- ConvNets stack CONV,POOL,FC layers
- Trend towards smaller filters and deeper architectures
- Trend towards getting rid of POOL/FC layers (just CONV)
- Typical architectures look like
 $[(CONV-RELU)^*N-POOL?]^*M-(FC-RELU)^*K, SOFTM$
AX
- where N is usually up to ~ 5 , M is large, $0 \leq K \leq 2$.
- but recent advances such as ResNet/GoogLeNet challenge this paradigm